

# Object Detection @ Google



Jonathan Huang ([jonathanhuang@google.com](mailto:jonathanhuang@google.com))

SIDAS 2017 Industry Keynote, 15 Sept 2017  
Presenting the work of **many** people @Google

**Google Research and Machine Intelligence**

# Imagenet Challenge

Given an image,  
predict one of 1000  
different classes

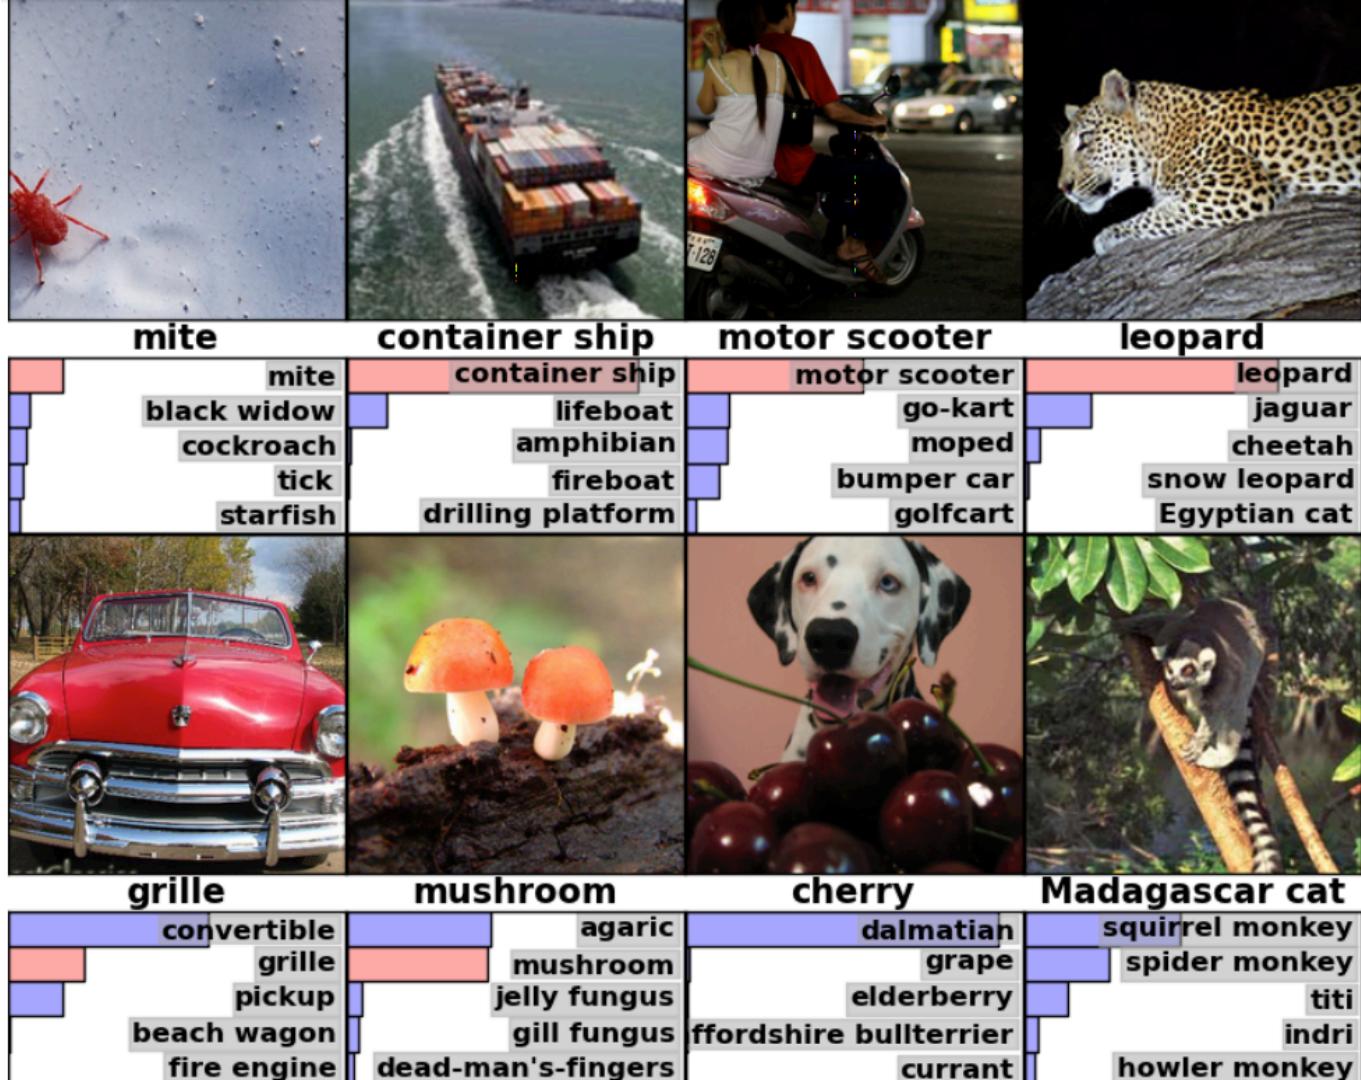
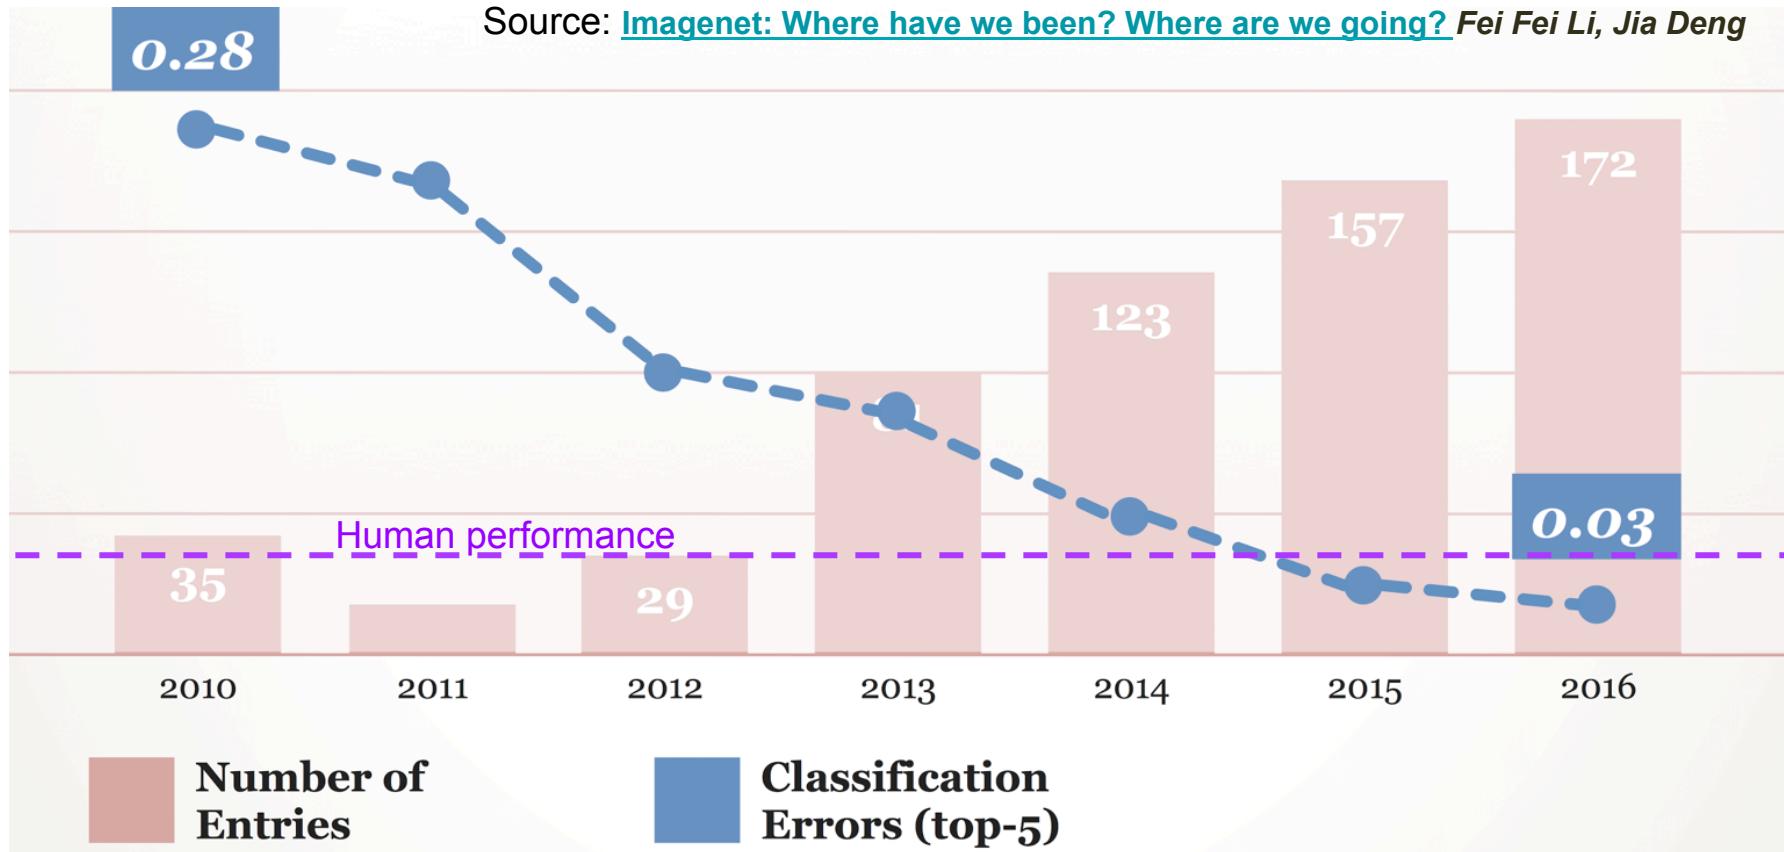


Image credit:

[www.cs.toronto.edu/~fritz/absps/imagenet.pdf](http://www.cs.toronto.edu/~fritz/absps/imagenet.pdf)

# Imagenet Progress Over the Years

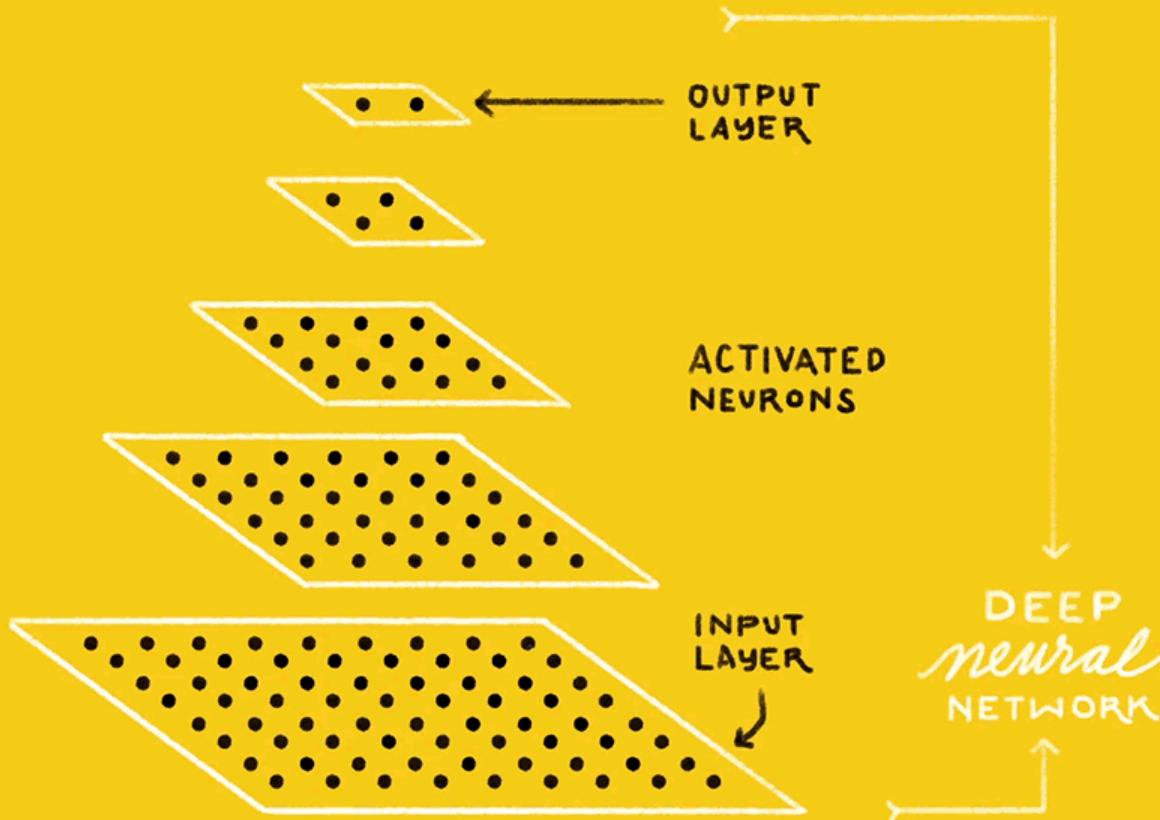
Source: [Imagenet: Where have we been? Where are we going? Fei Fei Li, Jia Deng](#)



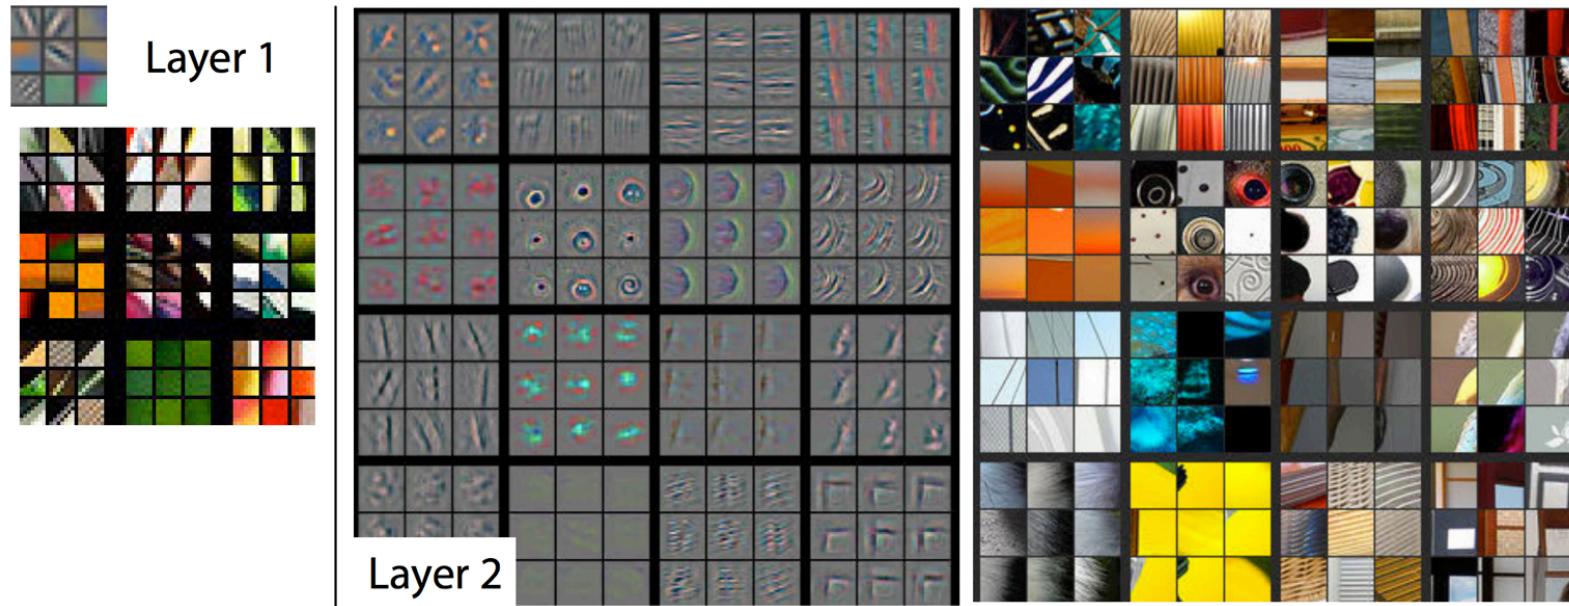
IS THIS A  
**CAT or DOG?**



CAT   DOG

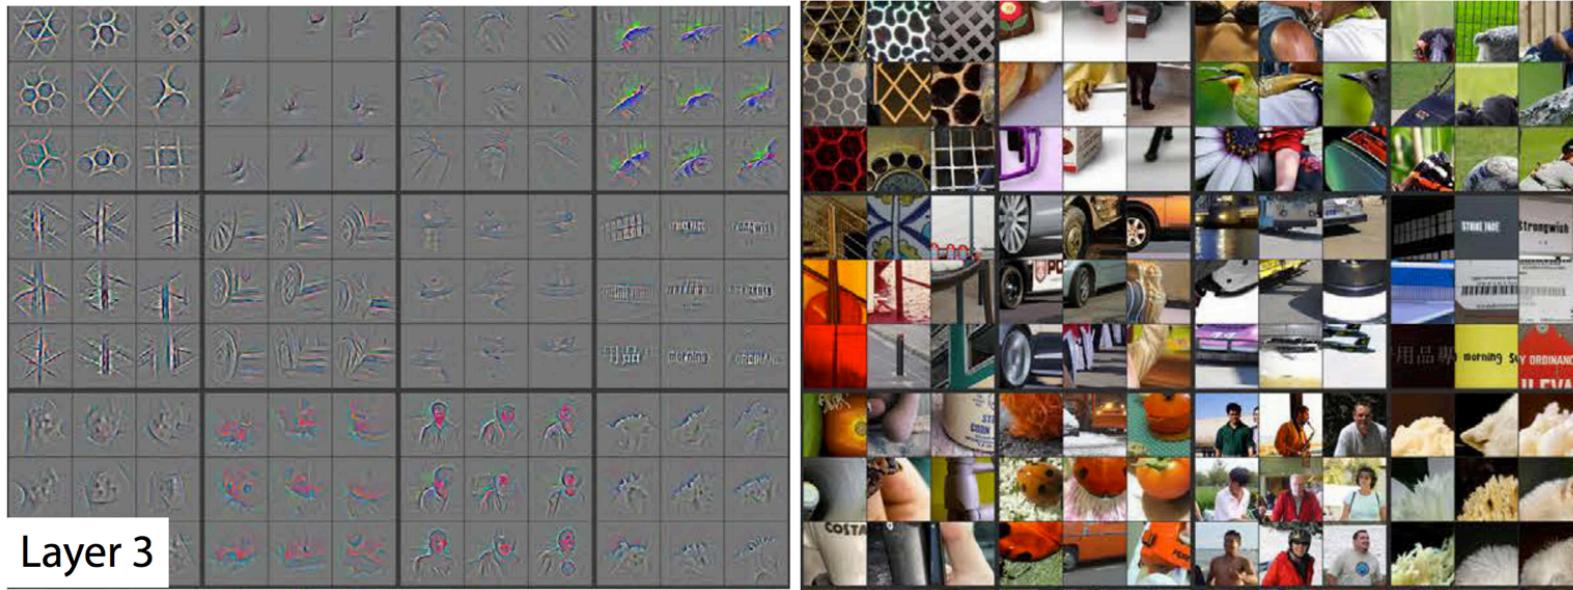


# Convolutions, Sliding Windows, Patterns!



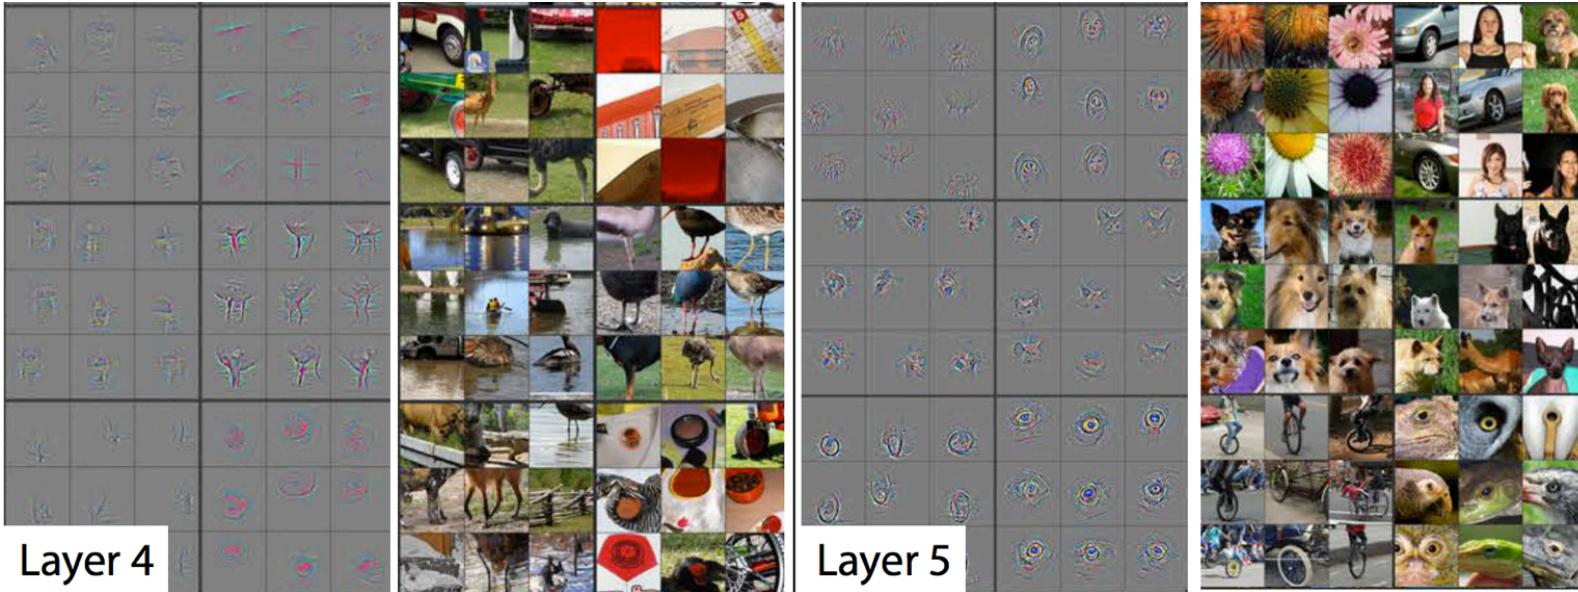
Source: [Visualizing and Understanding Convolutional Networks](#) M.D. Zeiler, R. Fergus

# Convolutions, Sliding Windows, Patterns!



Source: [Visualizing and Understanding Convolutional Networks](#) M.D. Zeiler, R. Fergus

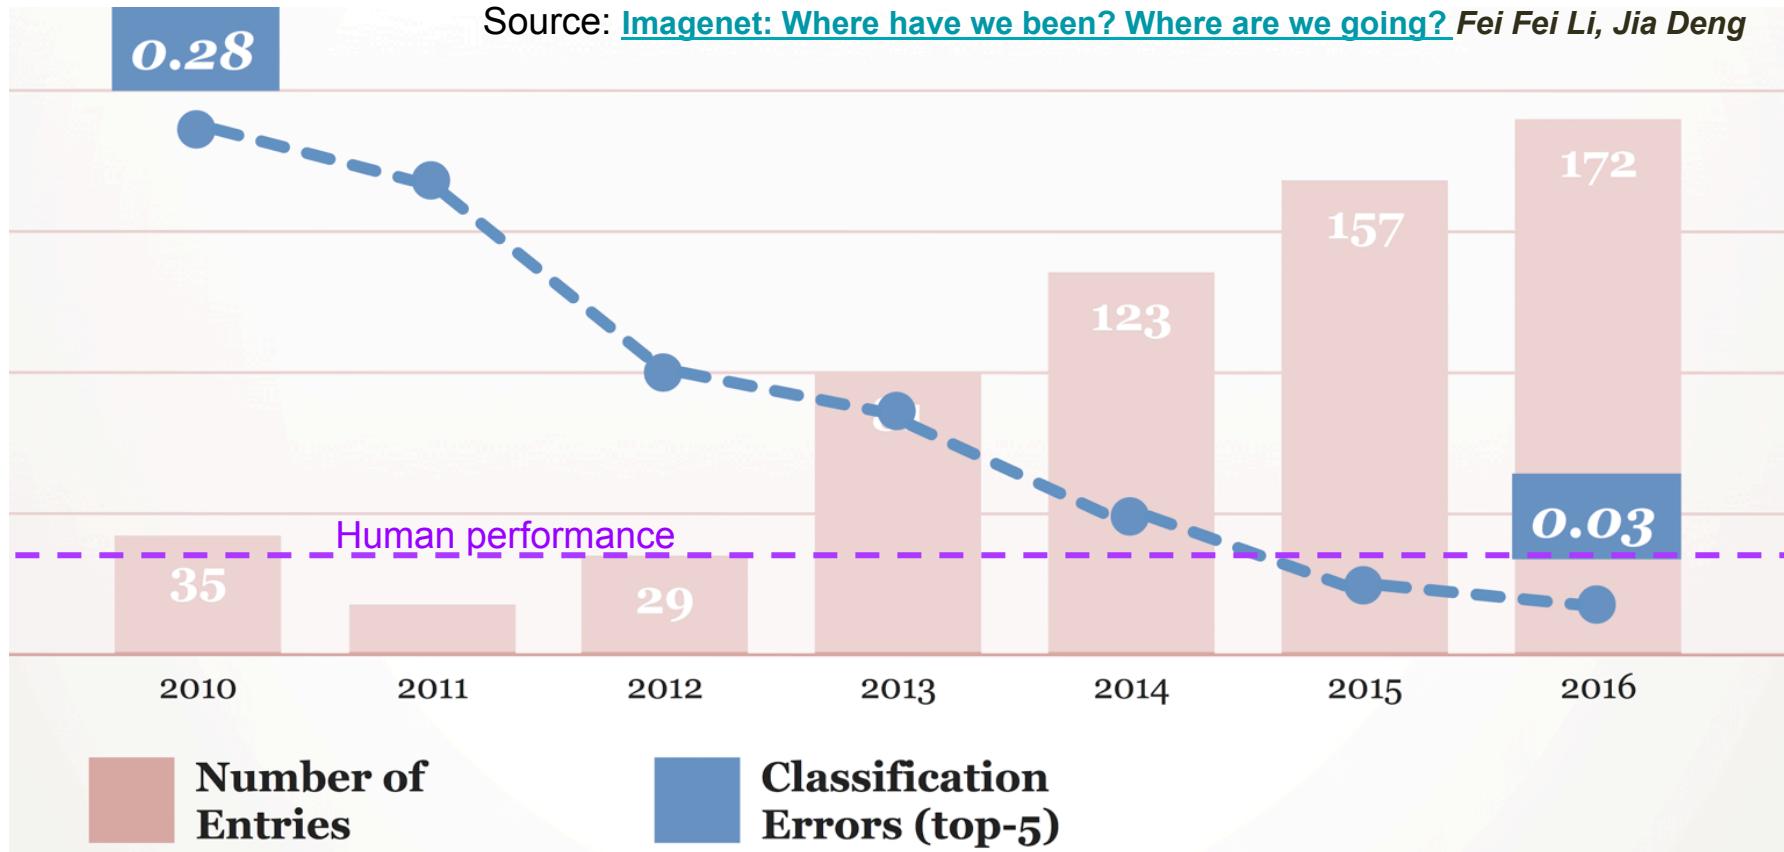
# Convolutions, Sliding Windows, Patterns!



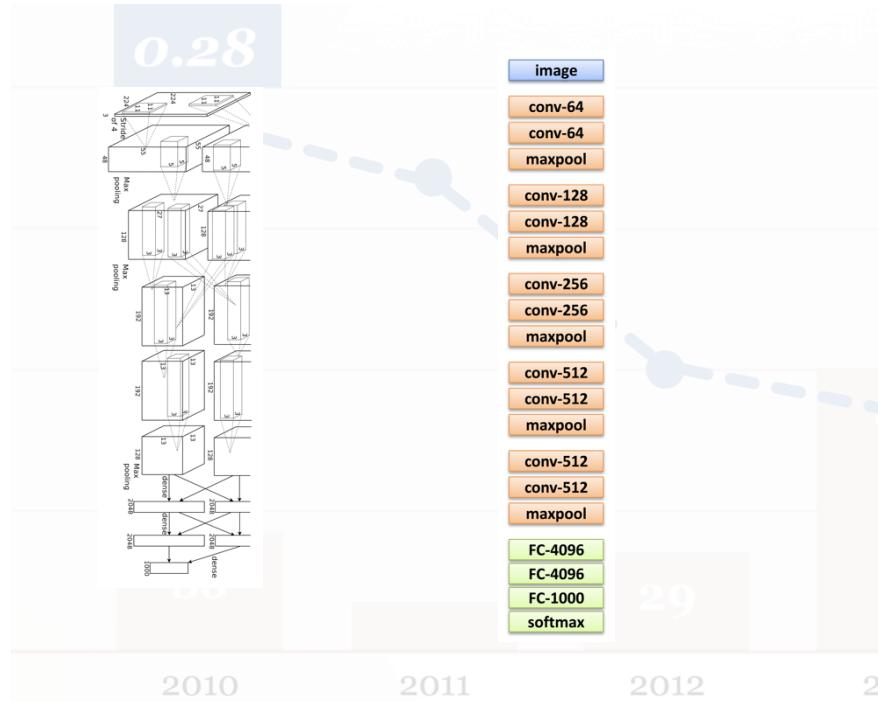
Source: [Visualizing and Understanding Convolutional Networks](#) M.D. Zeiler, R. Fergus

# Imagenet Progress Over the Years

Source: [Imagenet: Where have we been? Where are we going? Fei Fei Li, Jia Deng](#)



# “Famous Architectures”



AlexNet  
[Krizhevsky et al, 2012]

VGG

[Simonyan and Zisserman, 2014]

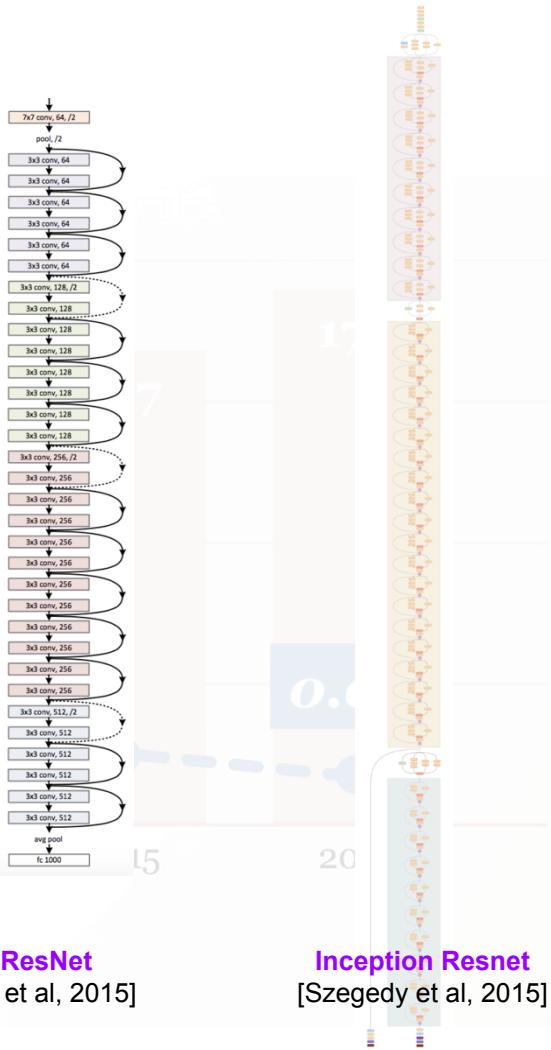
Inception  
Classification Error (top 5)

Inception

[Szegedy et al, 2014]

ResNet  
[He et al, 2015]

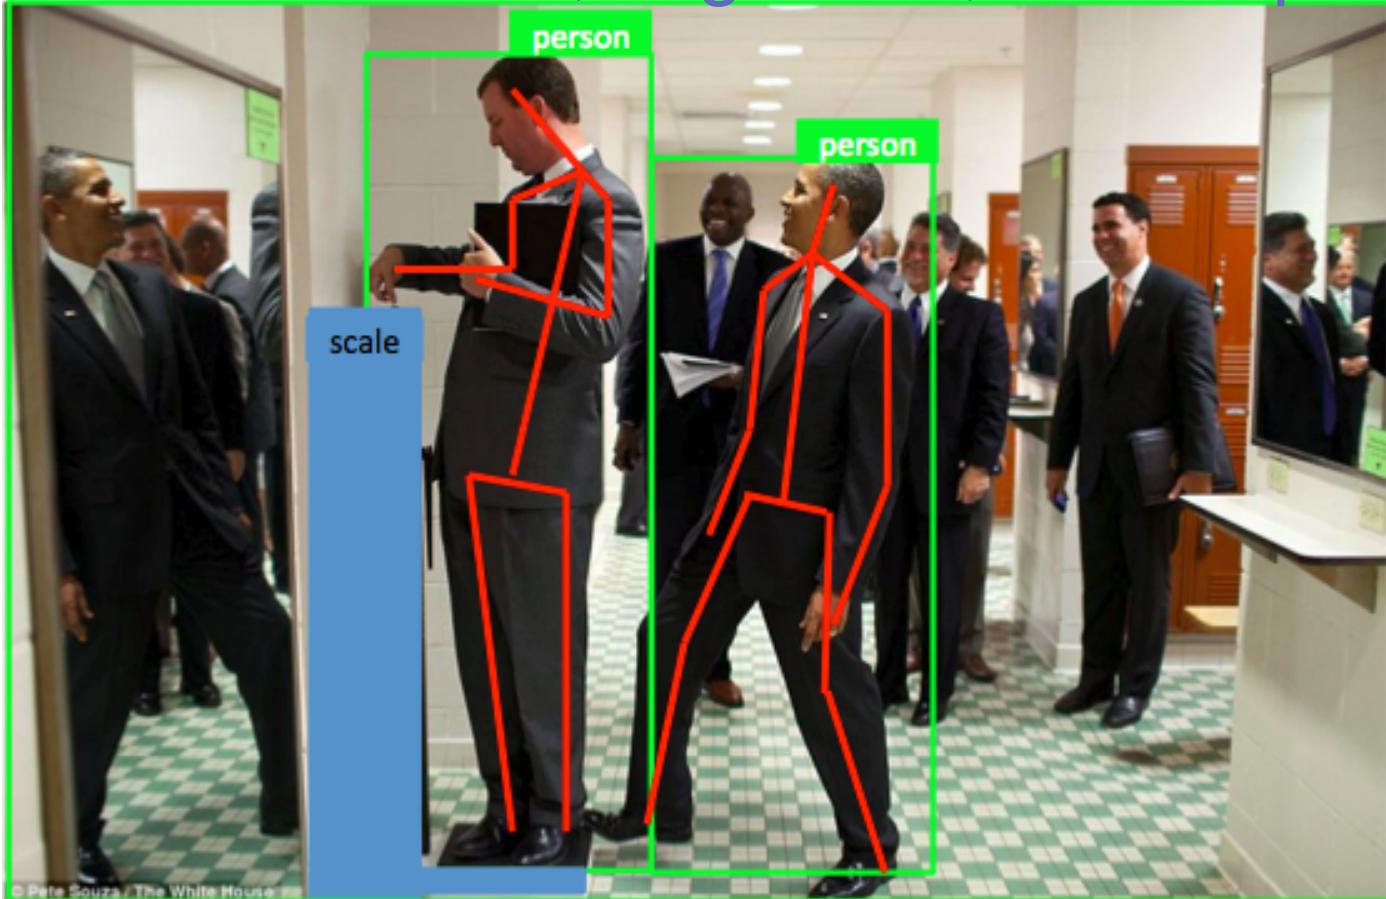
Inception Resnet  
[Szegedy et al, 2015]



Today's image taggers just return a bag of words...

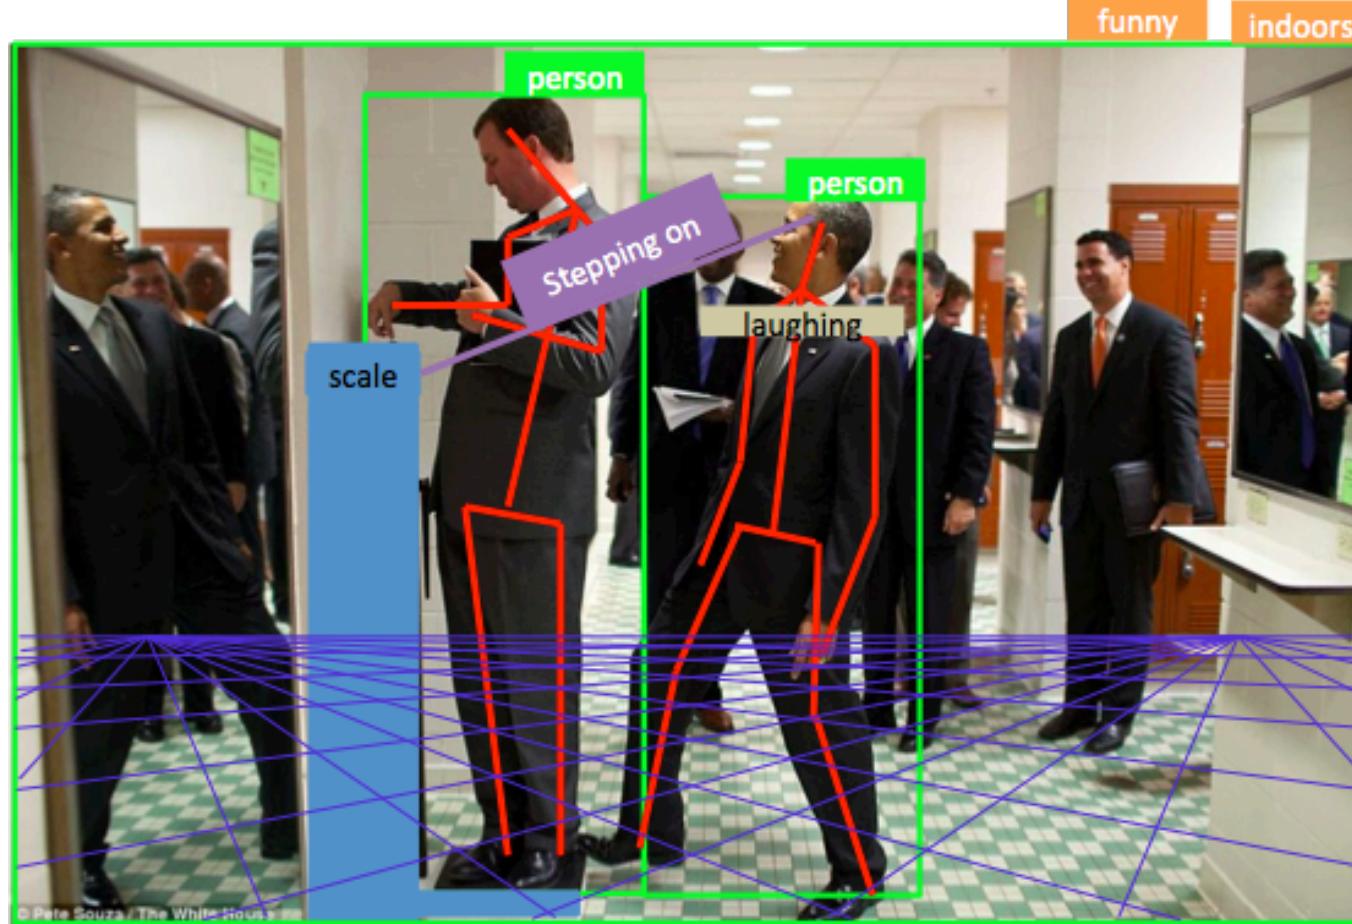


# What we need: boxes, segments, human pose...



Based on a figure from Jia Deng

# And finally: attributes, relations, 3-d!



# From Classification to Detection

## Challenges

- Variable # outputs
- Localization
- Running time
- Even with perfect localization, need to classify based on much fewer pixels than in Imagenet setting, requires context!

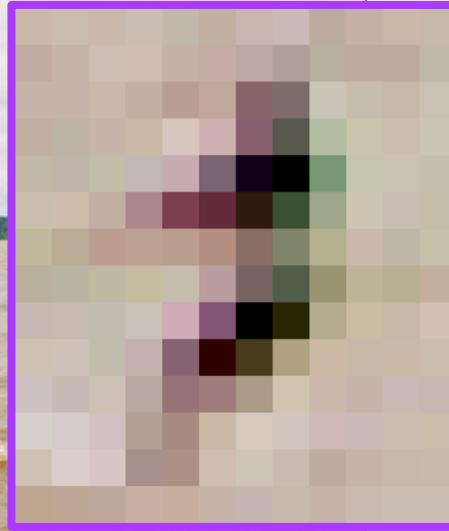
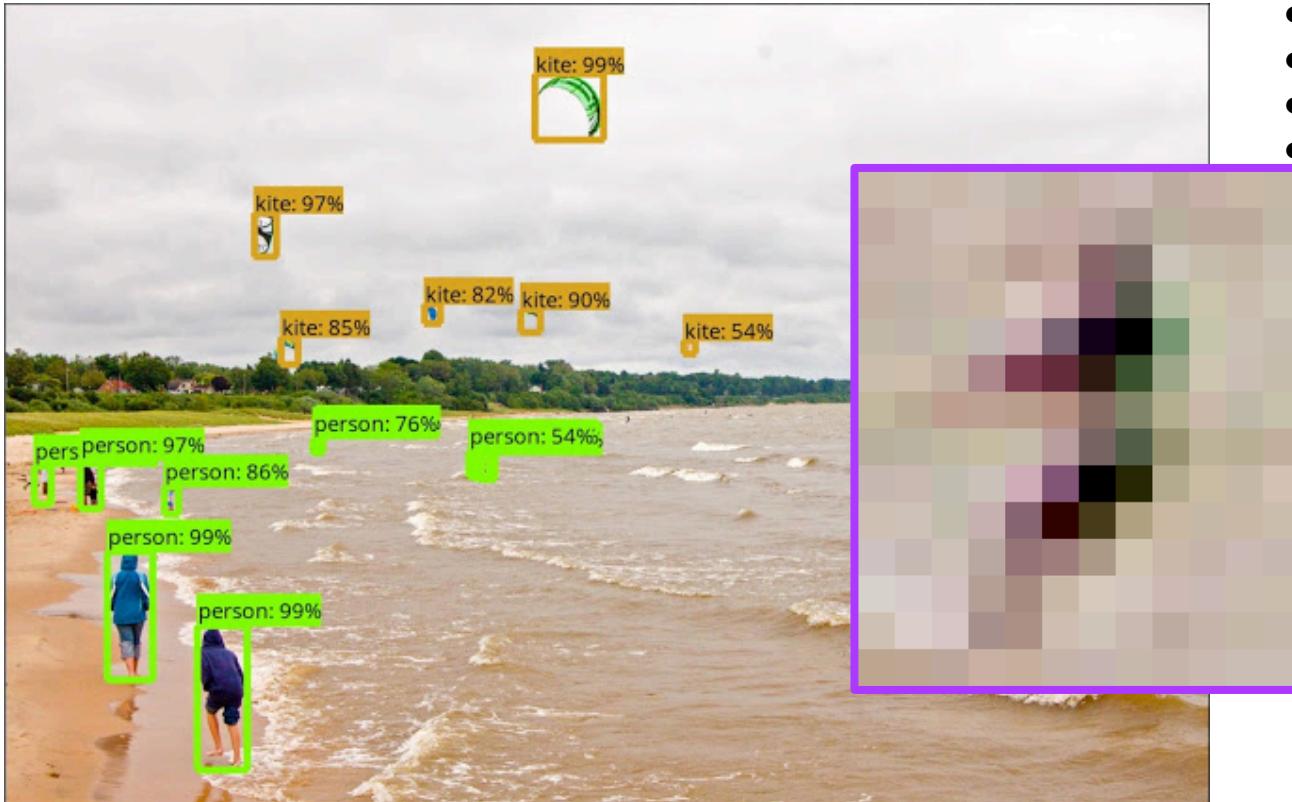
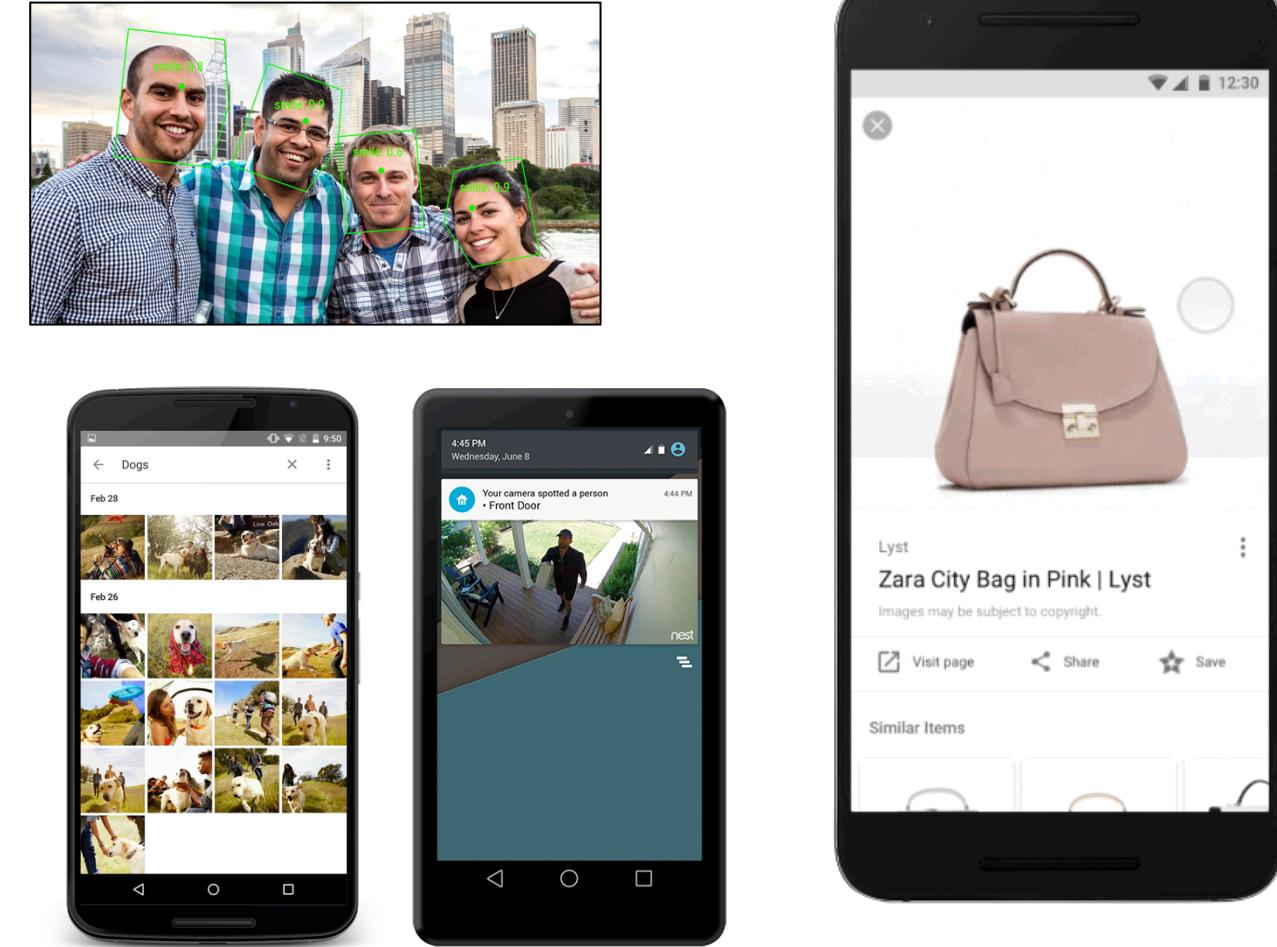
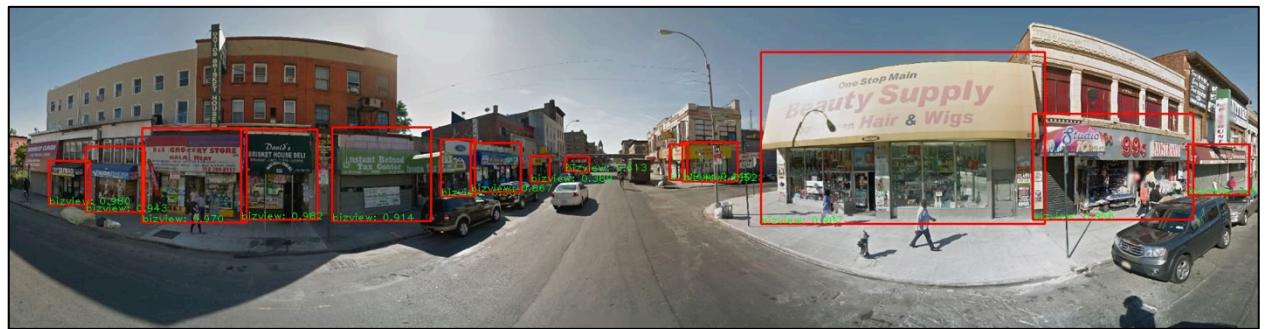


Photo credit: Michael Mina

# Object Detection @ Google



# Object Detection @ Google



# This Talk

- **The Tensorflow Object Detection API**
- Pushing the envelope on Accuracy
  - Better Models
  - Better Data
- Pushing the envelope on Speed
  - Lightweight Models
  - Adaptive Inference

# Object Detection in Tensorflow



## TensorFlow

32,986 stars, 14,327 forks on Github  
(as of Sept 29, 2016)



Your laptop



Datacenters



Mobile



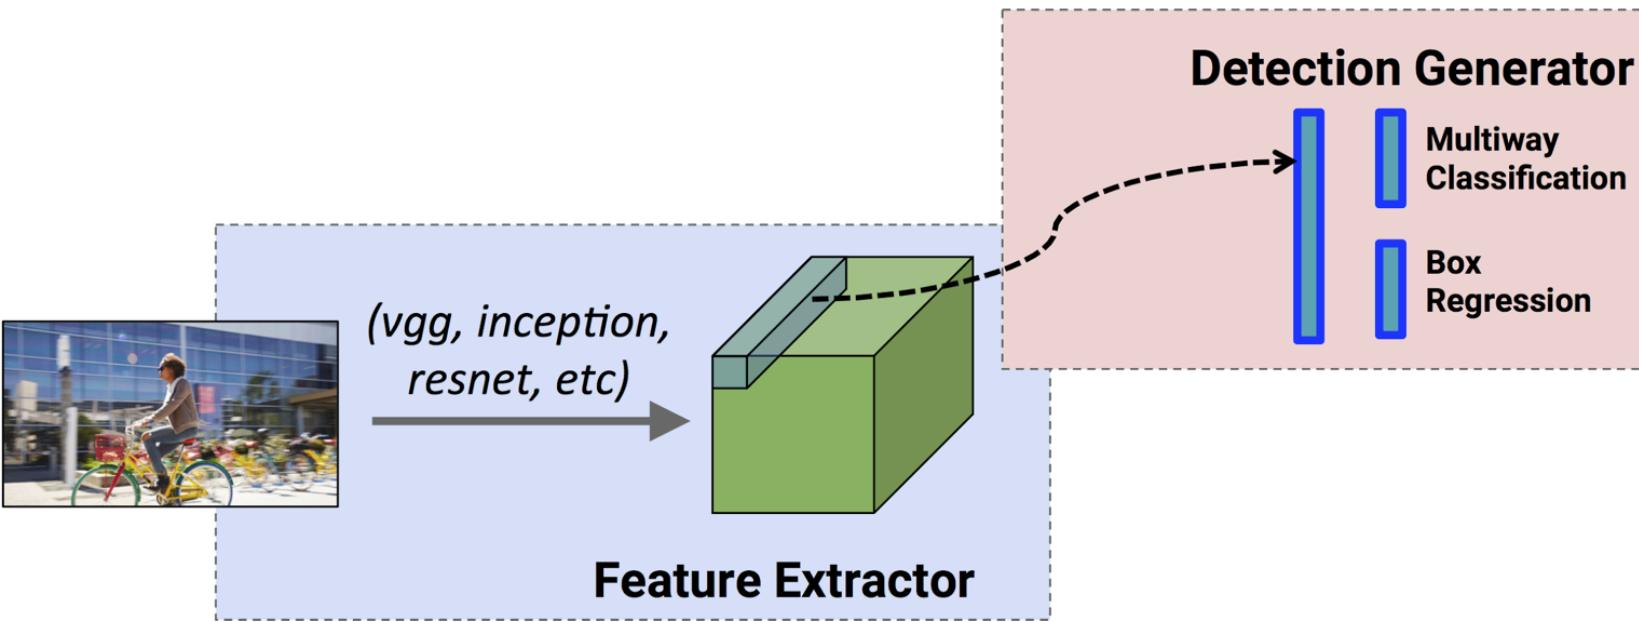
Raspberry Pi



Tensor Processing Unit

- Deploy models anywhere
- Scalable
- For research **and** production
- State-of-the-art performance
- Support leading methods  
Multibox/SSD, Fast/Faster  
RCNN, etc...

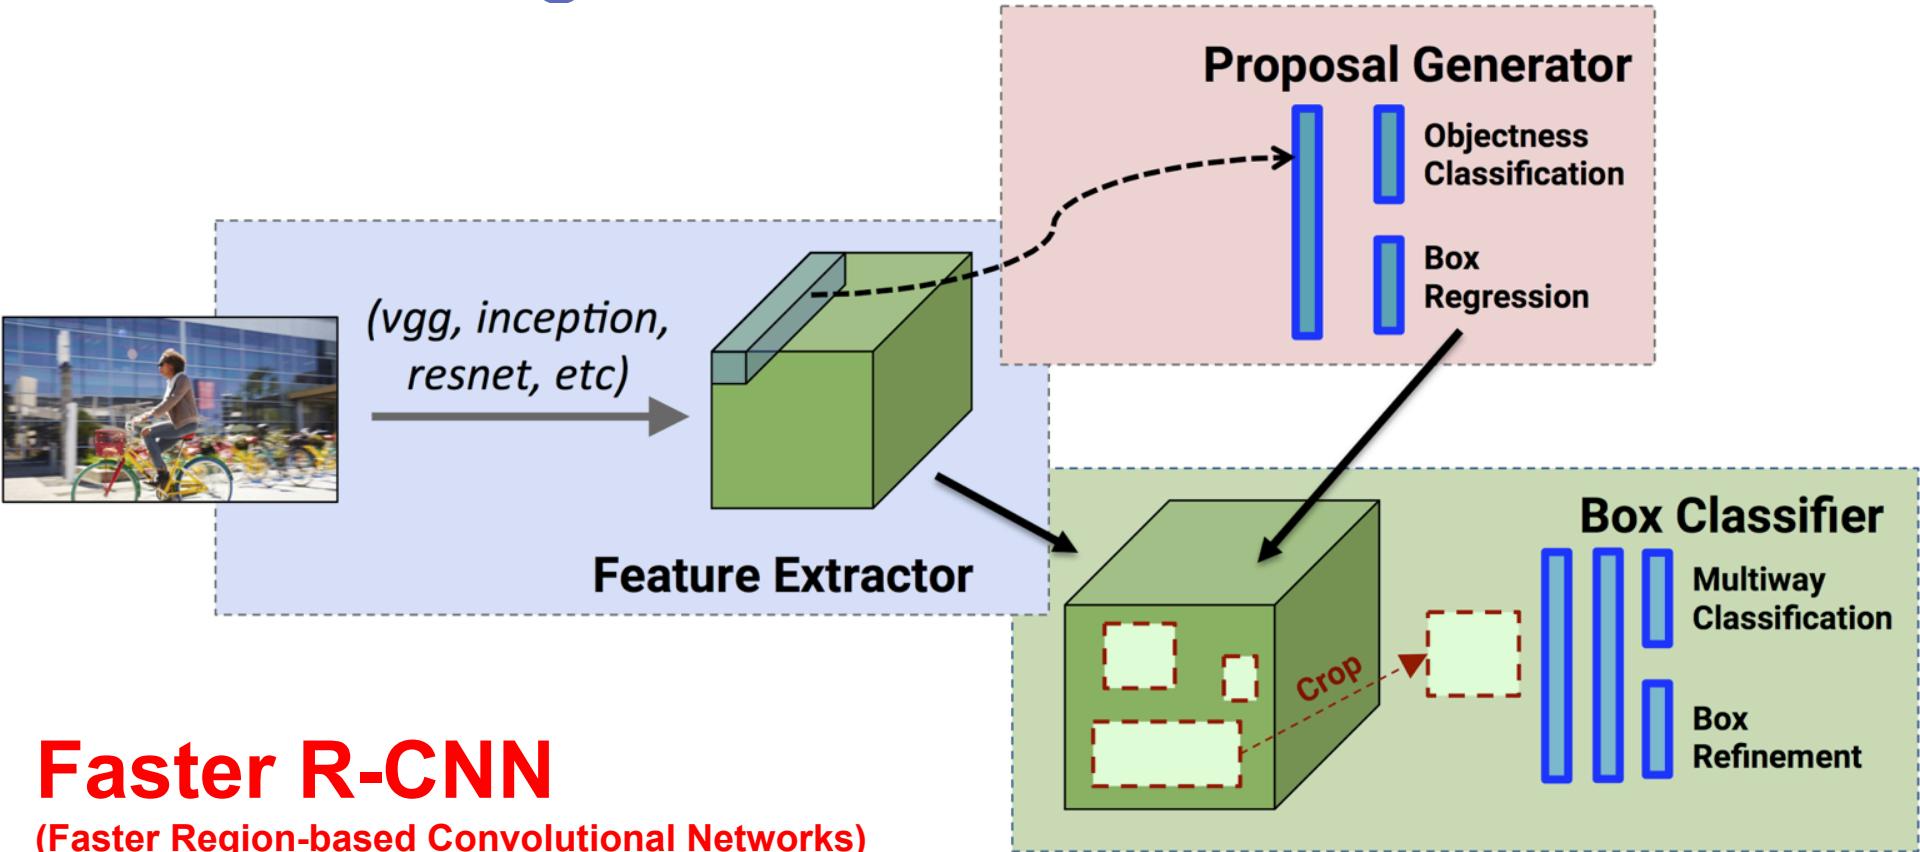
# A convergence of architectures



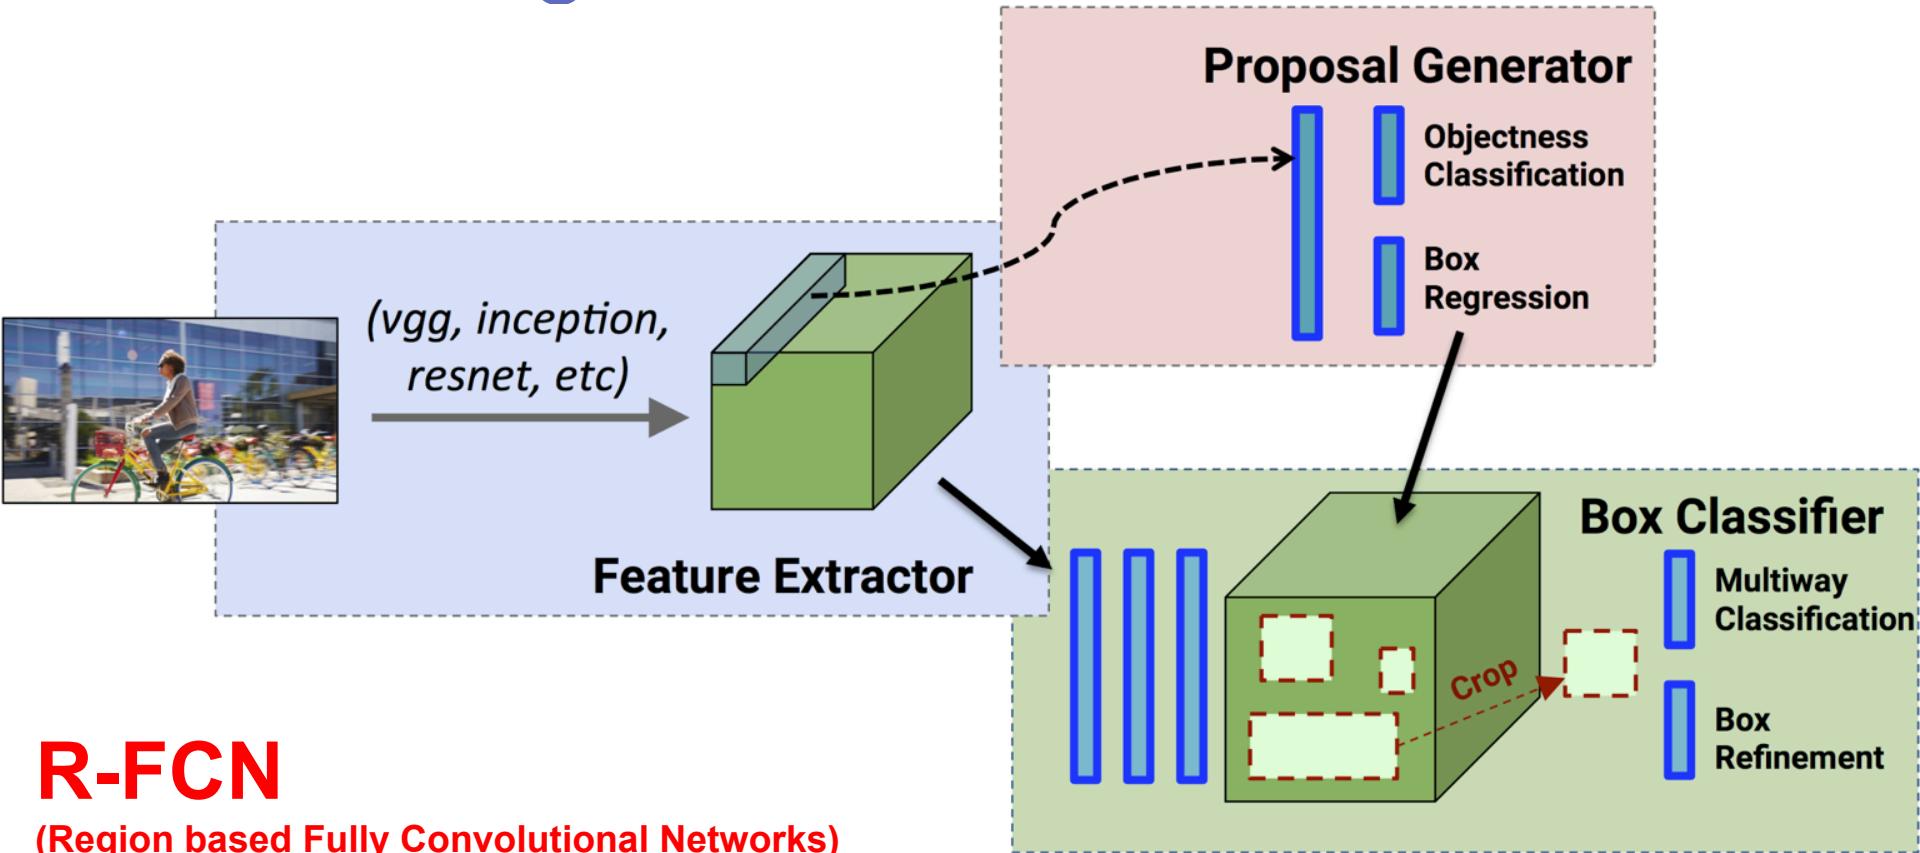
**SSD**

(Single Shot Detector --- encapsulates Multibox, YOLO, YOLO v2)

# A convergence of architectures



# A convergence of architectures



# Recent works in the detection literature

Paper	Meta-architecture	Feature Extractor	Matching	Box Encoding $\phi(b_a, a)$	Location Loss functions
Szegedy et al. [40]	SSD	InceptionV3	Bipartite	$[x_0, y_0, x_1, y_1]$	$L_2$
Redmon et al. [29]	SSD	Custom (GoogLeNet inspired)	Box Center	$[x_c, y_c, \sqrt{w}, \sqrt{h}]$	$L_2$
Ren et al. [31]	Faster R-CNN	VGG	Argmax	$[\frac{x_c}{w_a}, \frac{y_c}{h_a}, \log w, \log h]$	$SmoothL_1$
He et al. [13]	Faster R-CNN	ResNet-101	Argmax	$[\frac{x_c}{w_a}, \frac{y_c}{h_a}, \log w, \log h]$	$SmoothL_1$
Liu et al. [26] (v1)	SSD	InceptionV3	Argmax	$[x_0, y_0, x_1, y_1]$	$L_2$
Liu et al. [26] (v2, v3)	SSD	VGG	Argmax	$[\frac{x_c}{w_a}, \frac{y_c}{h_a}, \log w, \log h]$	$SmoothL_1$
Dai et al [6]	R-FCN	ResNet-101	Argmax	$[\frac{x_c}{w_a}, \frac{y_c}{h_a}, \log w, \log h]$	$SmoothL_1$

But which is the best?

# TensorFlow Object Detection API

Implement all the detectors



<b>Meta-architecture</b>	SSD, Faster R-CNN, R-FCN
<b>Feature Extractor</b>	Mobilenet, VGG, Inception V2, Inception V3, Resnet-50, Resnet-101, Resnet-152, Inception Resnet v2
<b>Learning schedule</b>	Manually Stepped, Exponential Decay, etc
<b>Location Loss function</b>	L2, L1, Huber, IOU
<b>Classification Loss function</b>	SigmoidCrossEntropy, SoftmaxCrossEntropy

CONFIG\_PATH="tensorflow/models/samples/configs/faster\_rcnn\_resnet101\_pets.config"

## Model parameters

```
model {  
    faster_rcnn {  
        num_classes: 37  
        image_resizer {  
            keep_aspect_ratio_resizer {  
                min_dimension: 600  
                max_dimension: 1024  
            }  
        }  
        feature_extractor {  
            type: 'faster_rcnn_resnet101'  
            first_stage_features_stride: 16  
        }  
        first_stage_anchor_generator {  
            grid_anchor_generator {  
                scales: [0.25, 0.5, 1.0, 2.0]  
                aspect_ratios: [0.5, 1.0, 2.0]  
                height_stride: 16  
                width_stride: 16  
            }  
        }  
        first_stage_box_predictor_conv_hyperparams {  
            op: CONV  
            regularizer {  
                l2_regularizer {  
                    weight: 0.0  
                }  
            }  
            initializer {  
                truncated_normal_initializer {  
                    stddev: 0.01  
                }  
            }  
        }  
        first_stage_nms_score_threshold: 0.0  
        first_stage_nms_iou_threshold: 0.7  
        first_stage_max_proposals: 300  
        first_stage_localization_loss_weight: 2.0  
        first_stage_objectness_loss_weight: 1.0  
    }  
}
```

37 dog/cat breeds

high resolution input images

Faster R-CNN, Resnet 101

## Train/Eval parameters

```
train_config: {  
    batch_size: 1  
    optimizer {  
        momentum_optimizer: {  
            learning_rate: {  
                manual_step_learning_rate {  
                    initial_learning_rate: 0.0003  
                }  
                schedule {  
                    step: 0  
                    learning_rate: .0003  
                }  
                schedule {  
                    step: 900000  
                    learning_rate: .00003  
                }  
                schedule {  
                    step: 1200000  
                    learning_rate: .000003  
                }  
            }  
            momentum_optimizer_value: 0.9  
        }  
        use_moving_average: false  
    }  
    gradient_clipping_by_norm: 10.0  
    fine_tune_checkpoint: "PATH_TO_BE_CONFIGURED/model.ckpt"  
    from_detection_checkpoint: true  
    data_augmentation_options {  
        random_horizontal_flip {  
        }  
    }  
    eval_config: {  
        num_examples: 2000  
    }  
}
```

manually stepped learning rate schedule

Detection checkpoint pretrained on COCO

## Data sources

```
train_input_reader: {  
    tf_record_input_reader {  
        input_path: "PATH_TO_BE_CONFIGURED/pet_train.record"  
    }  
    label_map_path: "PATH_TO_BE_CONFIGURED/pet_label_map.pbtxt"  
}  
  
eval_input_reader: {  
    tf_record_input_reader {  
        input_path: "PATH_TO_BE_CONFIGURED/pet_val.record"  
    }  
    label_map_path: "PATH_TO_BE_CONFIGURED/pet_label_map.pbtxt"  
}
```

CONFIG\_PATH="tensorflow/models/samples/configs/faster\_rcnn\_resnet101\_pascal.config"

## Model parameters

```
model {  
  faster_rcnn {  
    num_classes: 20  
  }  
  image_resizer {  
    keep_aspect_ratio_resizer {  
      min_dimension: 600  
      max_dimension: 1024  
    }  
  }  
  feature_extractor {  
    type: 'faster_rcnn_resnet101'  
    first_stage_features_stride: 16  
  }  
  first_stage_anchor_generator {  
    grid_anchor_generator {  
      scales: [0.25, 0.5, 1.0, 2.0]  
      aspect_ratios: [0.5, 1.0, 2.0]  
      height_stride: 16  
      width_stride: 16  
    }  
  }  
  first_stage_box_predictor_conv_hyperparams {  
    op: CONV  
    regularizer {  
      l2_regularizer {  
        weight: 0.0  
      }  
    }  
    initializer {  
      truncated_normal_initializer {  
        stddev: 0.01  
      }  
    }  
  }  
  first_stage_nms_score_threshold: 0.0  
  first_stage_nms_iou_threshold: 0.7  
  first_stage_max_proposals: 300  
  first_stage_localization_loss_weight: 2.0  
  first_stage_objectness_loss_weight: 1.0
```

## Train/Eval parameters

```
train_config: {  
  batch_size: 1  
  optimizer {  
    momentum_optimizer: {  
      learning_rate: {  
        manual_step_learning_rate {  
          initial_learning_rate: 0.0003  
        }  
        schedule {  
          step: 0  
          learning_rate: .0003  
        }  
        schedule {  
          step: 900000  
          learning_rate: .00003  
        }  
        schedule {  
          step: 1200000  
          learning_rate: .000003  
        }  
      }  
      momentum_optimizer_value: 0.9  
    }  
    use_moving_average: false  
  }  
  gradient_clipping_by_norm: 10.0  
  fine_tune_checkpoint: "PATH_TO_BE_CONFIGURED/model.ckpt"  
  from_detection_checkpoint: true  
  data_augmentation_options {  
    random_horizontal_flip {  
    }  
  }  
}  
eval_config: {  
  num_examples: 2000  
}
```

## Data sources

```
train_input_reader: {  
  tf_record_input_reader {  
    input_path: "PATH_TO_BE_CONFIGURED/pascal_train.record"  
  }  
  label_map_path: "PATH_TO_BE_CONFIGURED/  
pascal_label_map.pbtxt"  
}  
  
eval_input_reader: {  
  tf_record_input_reader {  
    input_path: "PATH_TO_BE_CONFIGURED/pascal_val.record"  
  }  
  label_map_path: "PATH_TO_BE_CONFIGURED/  
pascal_label_map.pbtxt"  
}
```

This repository Search

Pull requests Issues Marketplace Gist

Unwatch 1,311 Unstar 17,218 Fork 6,704

Code Issues 196 Pull requests 29 Projects 0 Wiki Insights

Branch: master models / object\_detection / Create new file Upload files Find file History

derekjchow committed with sguada Make Record scripts python3 compatible. (#1614) Latest commit 057203e 2 hours ago

..

**anchor\_generators** Add Tensorflow Object Detection API. (#1561) 6 days ago

**box\_coders** Add Tensorflow Object Detection API. (#1561) 6 days ago

**builders** Fix compatibility for model\_builder\_test.py (#1571) 4 days ago

**core** Add Tensorflow Object Detection API. (#1561) 6 days ago

**data** Add Tensorflow Object Detection API. (#1561) 6 days ago

**data\_decoders** Add Tensorflow Object Detection API. (#1561) 6 days ago

**g3doc** Fix ML Engine Dashboard link (#1599) a day ago

**matchers** Add Tensorflow Object Detection API. (#1561) 6 days ago

**meta\_architectures** Add Tensorflow Object Detection API. (#1561) 6 days ago

**models** Use spatial\_squeeze=False for ResNet feature extractors. (#1586) 4 days ago

**protos** Add Tensorflow Object Detection API. (#1561) 6 days ago

**samples** Reduce batchsize from 32->24 for SSD configs. 5 days ago

**test\_images** Add Tensorflow Object Detection API. (#1561) 6 days ago

**utils** Change visualizer font and jupyter notebook line thickness (#1589) 4 days ago

**BUILD** Add Tensorflow Object Detection API. (#1561) 6 days ago

**CONTRIBUTING.md** Add Tensorflow Object Detection API. (#1561) 6 days ago

**README.md** Clean up documentation. (#1563) 5 days ago

**\_init\_\_.py** Add Tensorflow Object Detection API. (#1561) 6 days ago

**create\_pascal\_tf\_record.py** Make Record scripts python3 compatible. (#1614) 2 hours ago

**create\_pascal\_tf\_record\_test.py** Add Tensorflow Object Detection API. (#1561) 6 days ago

**create\_pet\_tf\_record.py** Make Record scripts python3 compatible. (#1614) 2 hours ago

**eval.py** Add Tensorflow Object Detection API. (#1561) 6 days ago

**eval\_util.py** Add Tensorflow Object Detection API. (#1561) 6 days ago

**evaluator.py** Add Tensorflow Object Detection API. (#1561) 6 days ago

**export\_inference\_graph.py** Add Tensorflow Object Detection API. (#1561) 6 days ago

README.md

## Tensorflow Object Detection API

Creating accurate machine learning models capable of localizing and identifying multiple objects in a single image remains a core challenge in computer vision. The TensorFlow Object Detection API is an open source framework built on top of TensorFlow that makes it easy to construct, train and deploy object detection models. At Google we've certainly found this codebase to be useful for our computer vision needs, and we hope that you will as well.



Contributions to the codebase are welcome and we would love to hear back from you if you find this API useful. Finally if you use the Tensorflow Object Detection API for a research publication, please consider citing:

"Speed/accuracy trade-offs for modern convolutional object detectors."  
Huang J, Rathod V, Sun C, Zhu M, Korattikara A, Fathi A, Fischer I, Wojna Z,  
Song Y, Guadarrama S, Murphy K, CVPR 2017

[link][bibtext]

### Maintainers

- Jonathan Huang, github: [jch1](#)
- Vivek Rathod, github: [tombstone](#)
- Derek Chow, github: [derekjchow](#)  
Chen Sun, github: [chen\\_sun](#)

## Tensorflow detection model zoo

We provide a collection of detection models pre-trained on the [COCO dataset](#). These models can be useful for out-of-the-box inference if you are interested in categories already in COCO (e.g., humans, cars, etc). They are also useful for initializing your models when training on novel datasets.

In the table below, we list each such pre-trained model including:

- a model name that corresponds to a config file that was used to train this model in the `samples/configs` directory,
- a download link to a tar.gz file containing the pre-trained model,
- model speed (one of {slow, medium, fast}),
- detector performance on COCO data as measured by the COCO mAP measure. Here, higher is better, and we only report bounding box mAP rounded to the nearest integer.
- Output types (currently only `Boxes`)

You can un-tar each tar.gz file via, e.g.:

```
tar -xzvf ssd_mobilenet_v1_coco.tar.gz
```

Inside the un-tar'ed directory, you will find:

- a graph proto (`graph.pbtxt`)
- a checkpoint (`model.ckpt.data-00000-of-00001`, `model.ckpt.index`, `model.ckpt.meta`)
- a frozen graph proto with weights baked into the graph as constants (`frozen_inference_graph.pb`) to be used for out of the box inference (try this out in the Jupyter notebook!)

Model name	Speed	COCO mAP	Outputs
<a href="#">ssd_mobilenet_v1_coco</a>	fast	21	Boxes
<a href="#">ssd_inception_v2_coco</a>	fast	24	Boxes
<a href="#">rfcn_resnet101_coco</a>	medium	30	Boxes
<a href="#">faster_rcnn_resnet101_coco</a>	medium	32	Boxes
<a href="#">faster_rcnn_inception_resnet_v2_atrous_coco</a>	slow	37	Boxes

# API Organization

## Front Ends

Train, Eval, Export for serving,  
Demos (Jupyter Notebook,  
Android)

## High Level API

Meta-architectures (SSD, Faster R-CNN, ...)  
Models (SSD w/Mobilenet, Faster R-CNN  
w/Resnet, ...)

## Low Level API

Box operations (e.g., IOU, Clipping, Non max suppression), Box Representations, Target Assignment, Visualization ops,

*API can be used at any of these levels*

# Ongoing Community Engagement...

Is the IOU in Tensorflow Object Detection API wrong?

 I just dugged a bit through the Tensorflow Object Detection API code especially the eval\_util part, as I wanted to implement the COCO metrics.

 0  
 But I noticed that the metrics are solely calculated using the bounding boxes which have normalized coordinates between [0, 1]. There are no aspect ratios or absolute coordinates used.

 So, doesn't this mean that the intersection over unions calculated on these results are incorrect? Let's take an 200x100 image pixel as an example. If the box would be off by 20px to the left, that's 0.1 in normalized coordinates. But if it would be off by 20px to the top, that would be 0.2 in normalized coordinates.

Doesn't that mean, being off to the top is harder penalizing the score than being off to the side?

[tensorflow](#) [object-detection](#)

[share](#) [edit](#) [flag](#)

asked Jul 13 at 19:22  
 lapayo  
27 ● 6

[add a comment](#)

[start a bounty](#)

1 Answer

[active](#) [oldest](#) 

 I believe the predicted coordinates are resized to the absolute image coordinates in the eval binary.

 2  
 But the other thing I would say is that IOU is scale invariant in the sense that if you scale two boxes by some factor, they will still have the same IOU overlap. As an example if we scale by 2 in the x-direction and scale by 3 in the y direction: If A is  $(x_1, y_1, x_2, y_2)$  and B is  $(u_1, v_1, u_2, v_2)$ , then  $\text{IOU}((A, B)) = \text{IOU}((2^2x_1, 3^2y_1, 2^2x_2, 3^2y_2), (2^2u_1, 3^2v_1, 2^2u_2, 3^2v_2))$

 What this means is that evaluating in normalized coordinates should give the same result as evaluating in absolute coordinates.

[share](#) [edit](#) [delete](#) [flag](#)

answered Jul 13 at 22:09  
 Jonathan Huang  
646 ● 1 ● 8

Thanks for your answer, I was wrong, you are right about scale invariance. :-)- [lapayo](#) Jul 14 at 8:16

[add a comment](#)

*We provide regular support on StackOverflow (and Github)*

# Ongoing Community Engagement...

*Community members support each other too :)*

Tensorflow serving No versions of servable <MODEL> found under base path

[Ask Question](#)

I was following [this tutorial](#) to use tensorflow serving using my object detection model. I am using tensorflow object detection for generating the model. I have created a frozen model using [this](#) exporter (the generated frozen model [works](#) using python script).

The frozen graph directory has following contents ( nothing on variables directory)

variables/  
saved\_model.pb

Now when I try to serve the model using the following command,

```
tensorflow_model_server --port=9000 --model_name=ssd --model_base_path=/serving/ssd_froz
```

It always shows me

```
...  
tensorflow_serving/model_servers/server_core.cc:421] [Re-adding model: ssd 2017-08-07  
10:22:43.892834: W tensorflow_serving/sources/storage_path/file_system_storage_path_source.cc:262] No versions  
of servable ssd found under base path /serving/ssd_frozen/ 2017-08-07 10:22:44.892901: W tensorflow_serving/sources/storage_path/file_system_storage_path_source.cc:262] No versions  
of servable ssd found under base path /serving/ssd_frozen/  
...
```

Please help me. I have already spent a lot of time.

[tensorflow](#) [deep-learning](#) [object-detection](#) [tensorflow-serving](#)

[share](#) [edit](#) [flag](#)

edited 16 hours ago

asked 16 hours ago



Ultraviolet  
354 2 12

[add a comment](#)

1 Answer

[active](#) [oldest](#) [votes](#)

I had same problem, the reason is because object detection api does not assign version of your model when exporting your detection model. However, tensorflow serving requires you to assign a version number of your detection model, so that you could choose different versions of your models to serve. In your case, you should put your detection model(.pb file and variables folder) under folder: /serving/ssd\_frozen/1/. In this way, you will assign your model to version 1, and tensorflow serving will automatically load this version since you only have one version. By default tensorflow serving will automatically serve the latest version(e, the largest number of versions).

Note, after you created 1/ folder, the `model_base_path` is still need to be set to `--model_base_path=/serving/ssd_frozen/`.

[share](#) [edit](#) [flag](#)

edited 9 hours ago

answered 9 hours ago



Xinyao Wang  
38 4

Thank you very much for the help. Resolved the issue. :) – [Ultraviolet](#) 8 hours ago

[add a comment](#)

FEATURED ON META

- [Control the types of email you receive via our new Email Settings feature](#)
- [2017 Community Moderator Election RESULTS](#)
- [Sunsetting Documentation](#)
- [Help set Q&A \(TeamDAG\) product development priorities](#)

HOT META POSTS

- [Overlapping/conflicting review bar bug](#)

Microsoft

Build smarter apps with cognitive APIs and machine learning.

[Try Azure free](#)

Want a [python](#) job?

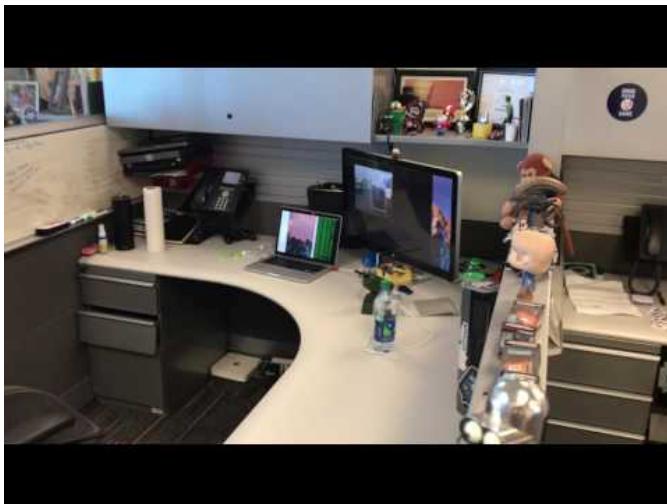
**Full-Stack Software Engineer**  
Charlie Finance Co. 9 San Francisco, CA  
**\$80K - \$130K** [RELOCATION](#)  
[python](#) [machine-learning](#)

**Software Engineer, Tools and Infrastructure**  
Google 9 San Francisco, CA  
[python](#) [shell](#)

Related

- [Tensorflow Serving - No versions of servable <model> found under base path](#)
- [Use keras with tensorflow serving](#)
- [How to create a tensorflow serving client](#)

# Community Creations!



Testing Custom Object Detector - TensorFlow Object Detection API  
4,524 views • 3 days ago



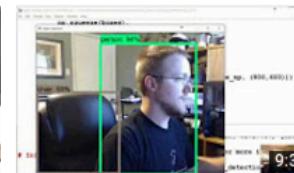
Training Custom Object Detector - TensorFlow Object Detection  
3,007 views • 3 days ago



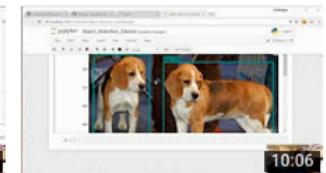
Creating TFRecords - TensorFlow Object Detection API Tutorial p.4  
3,145 views • 3 days ago



Tracking Custom Objects - TensorFlow Object Detection API  
4,600 views • 3 days ago



Adapting to video feed - TensorFlow Object Detection API  
15,157 views • 6 days ago



Intro - TensorFlow Object Detection API Tutorial p.1  
16,571 views • 1 week ago

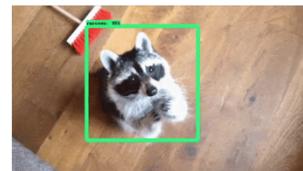


Dat Tran [Follow](#)  
Data Scientist at Pindar Labs, Berlin.  
Jul 26 • 8 min read

## How to train your own Object Detector with TensorFlow's Object Detector API

This is a follow-up post on "Building a Real-Time Object Recognition App with Tensorflow and OpenCV" where I focus on training my own classes. Specifically, I trained my own Raccoon detector on a dataset that I collected and labeled by myself. The full dataset is available on [my Github repo](#).

By the way, here is the Raccoon detector in action:



The Raccoon detector.



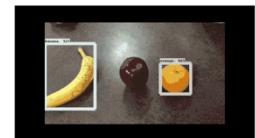
Priya Daveli [Follow](#)  
Priya Daveli is a machine learning, deep learning and AI

Jul 11 • 4 min read

## Is Google Tensorflow Object Detection API the easiest way to implement image recognition?

*Doing cool things with data!*

There are many different ways to do image recognition. Google recently released a new Tensorflow Object Detection API to give computer vision everywhere a boost. Any offering from Google is not to be taken lightly, and so I decided to try my hands on this new API and use it on videos from youtube :) See the result below:



Object Detection from Tensorflow API

You can find the full code on my [Github repo](#)

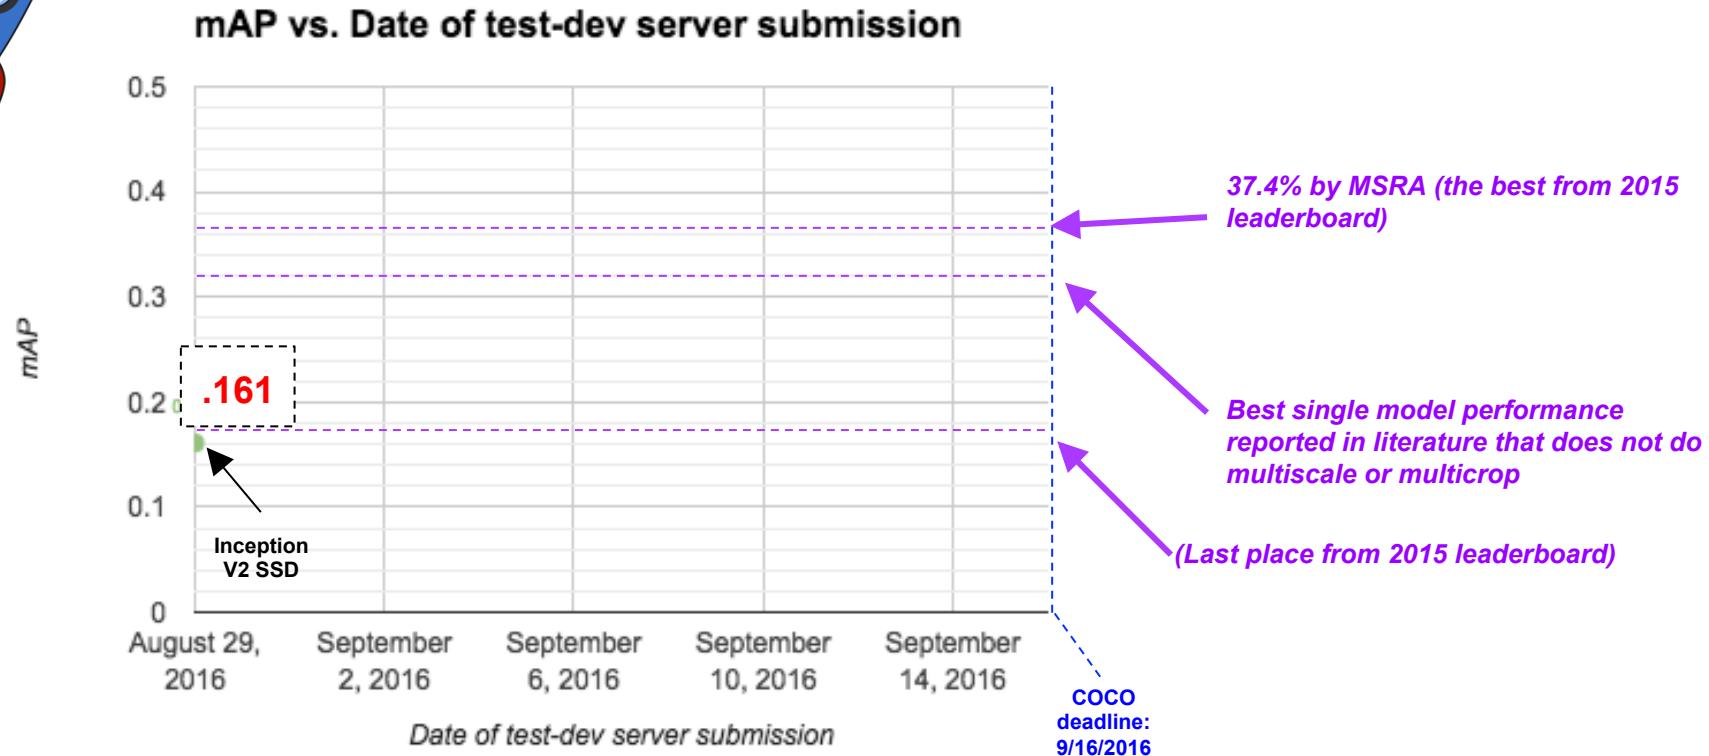
# This Talk

- The Tensorflow Object Detection API
- **Pushing the envelope on Accuracy**
  - **Better Models**
  - Better Data
- Pushing the envelope on Speed
  - Lightweight Models
  - Adaptive Inference

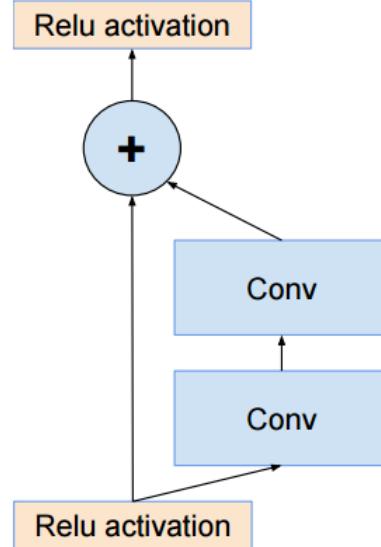
# Race to the Top

First submission to test-dev: Tensorflow implementation of SSD(ish) model with 224x224 input images

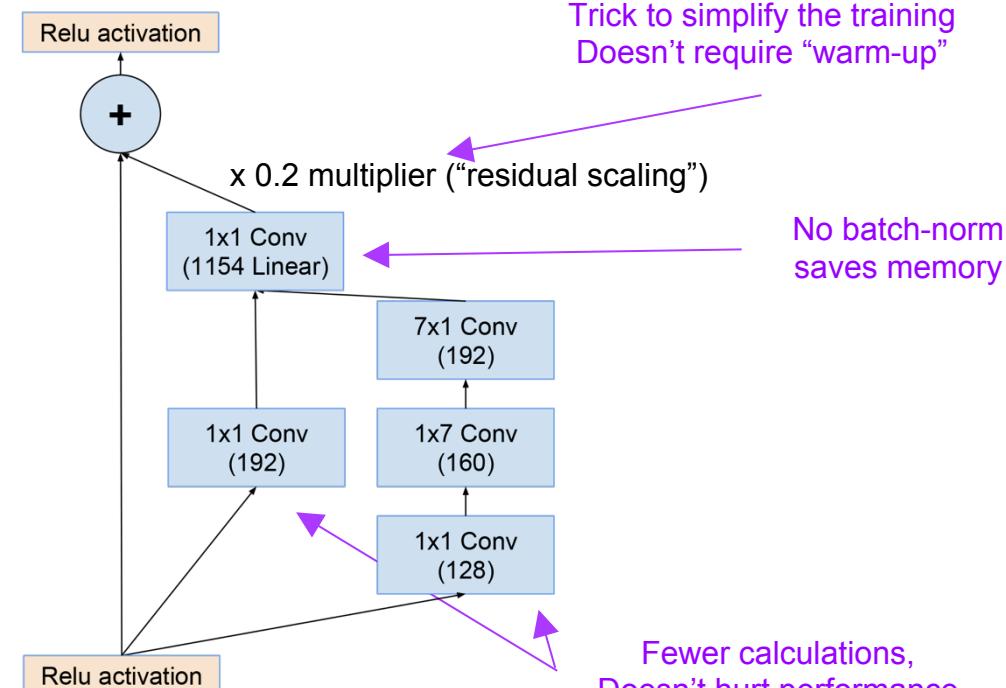
Liu, Wei, et al. "SSD: Single Shot MultiBox Detector." *arXiv preprint arXiv:1512.02325* (2015).



# Residual Blocks vs. Inception Resnet Blocks



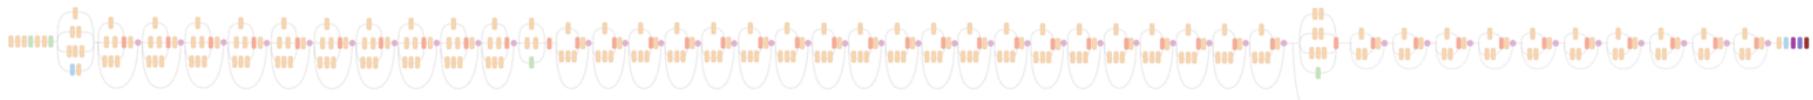
Residual Block



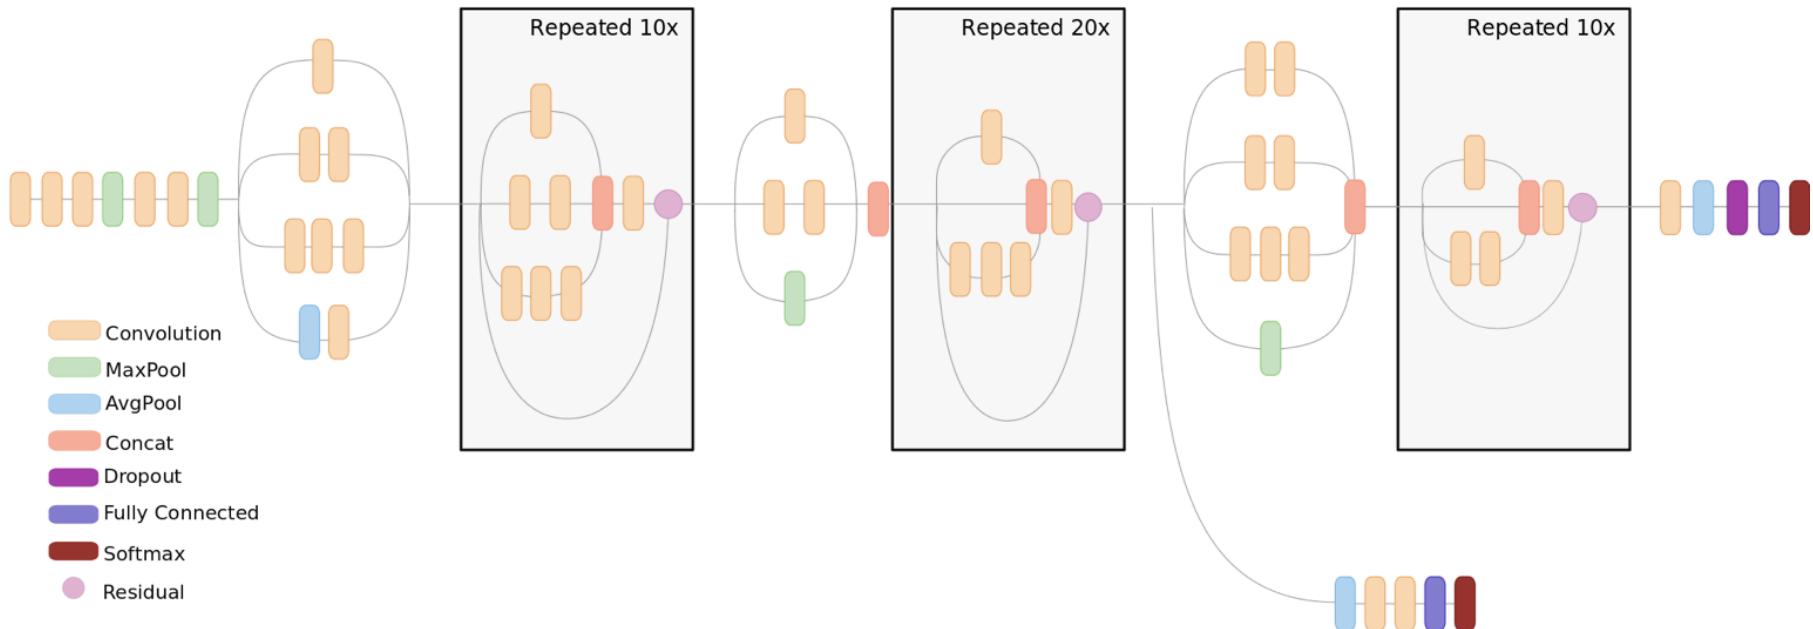
Inception Resnet Block

# Inception Resnet (v2) Feature Extractor

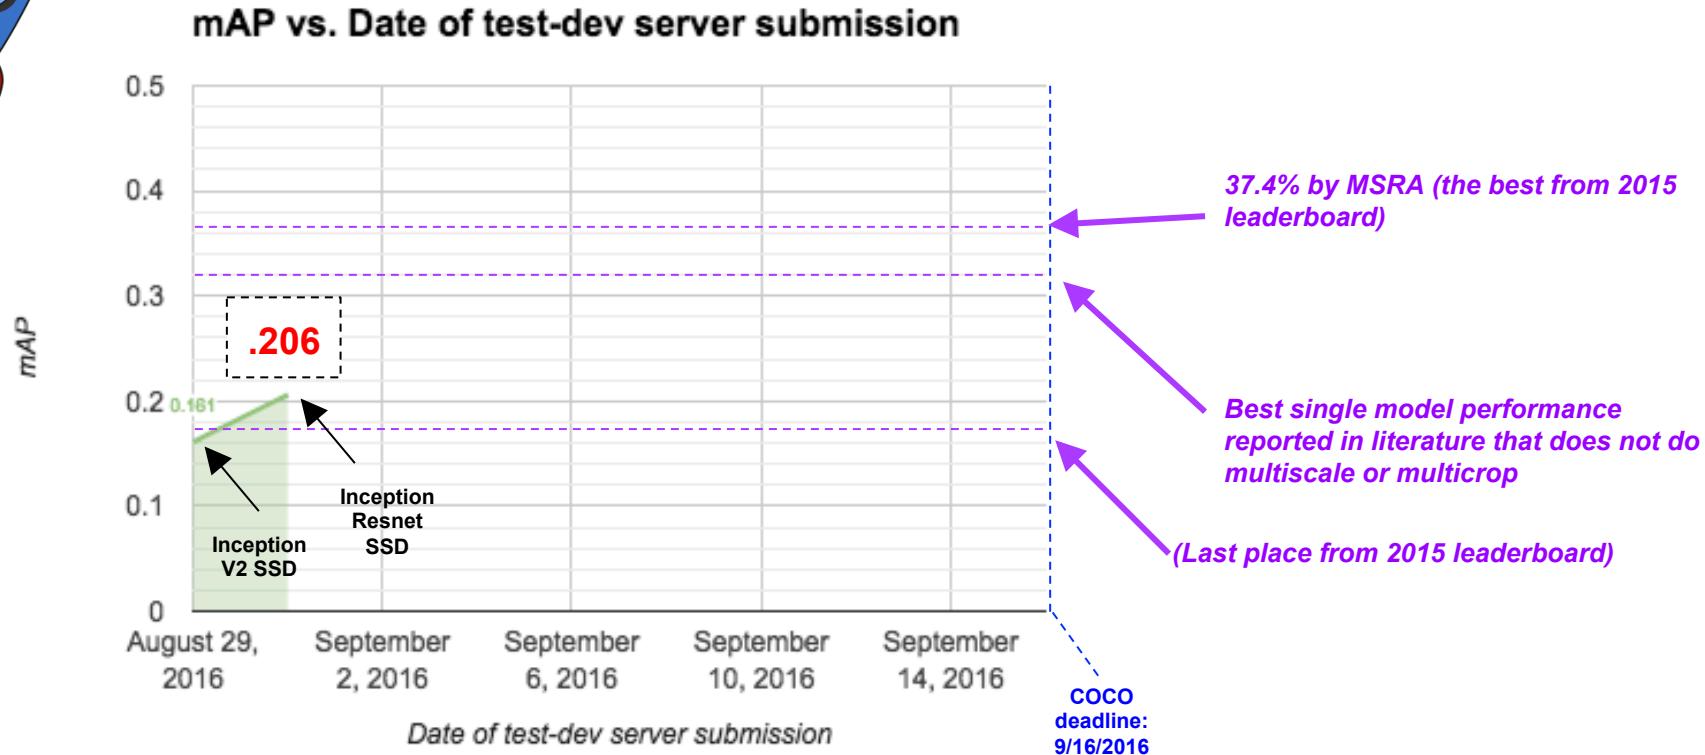
Full Inception Resnet V2 Network



Detailed view with compressed residual blocks



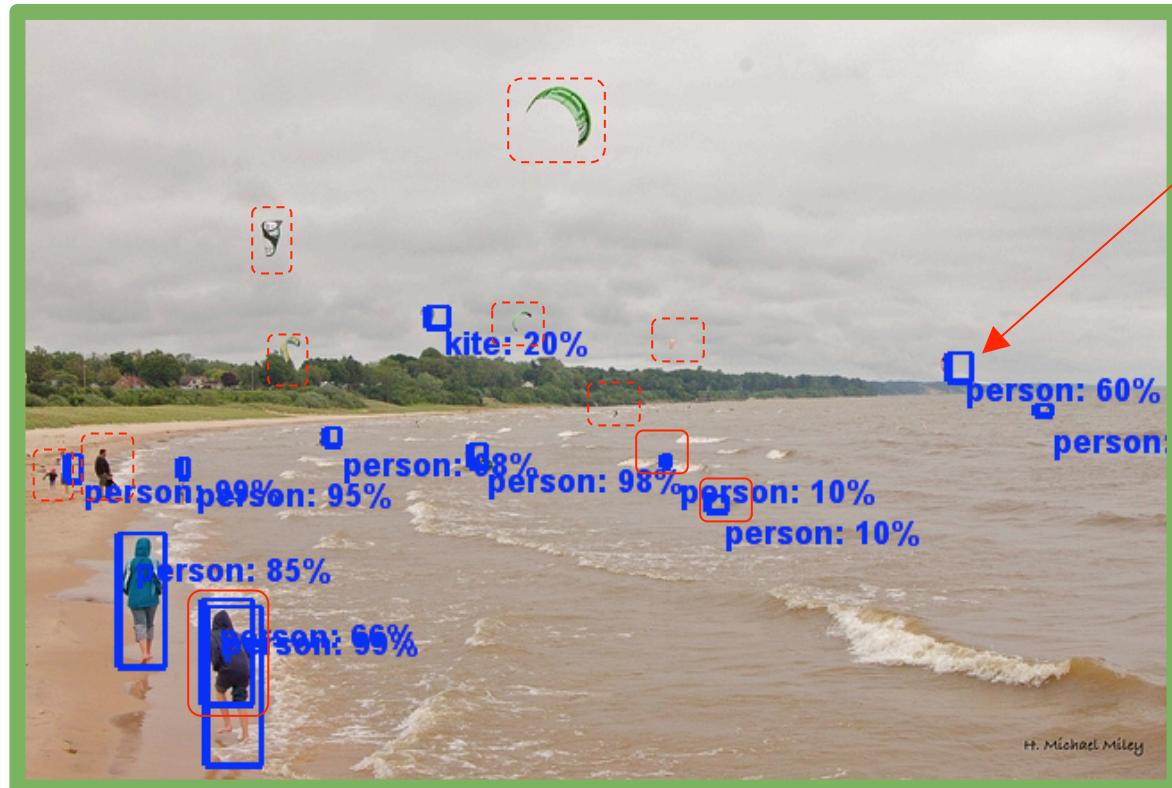
# Race to the Top





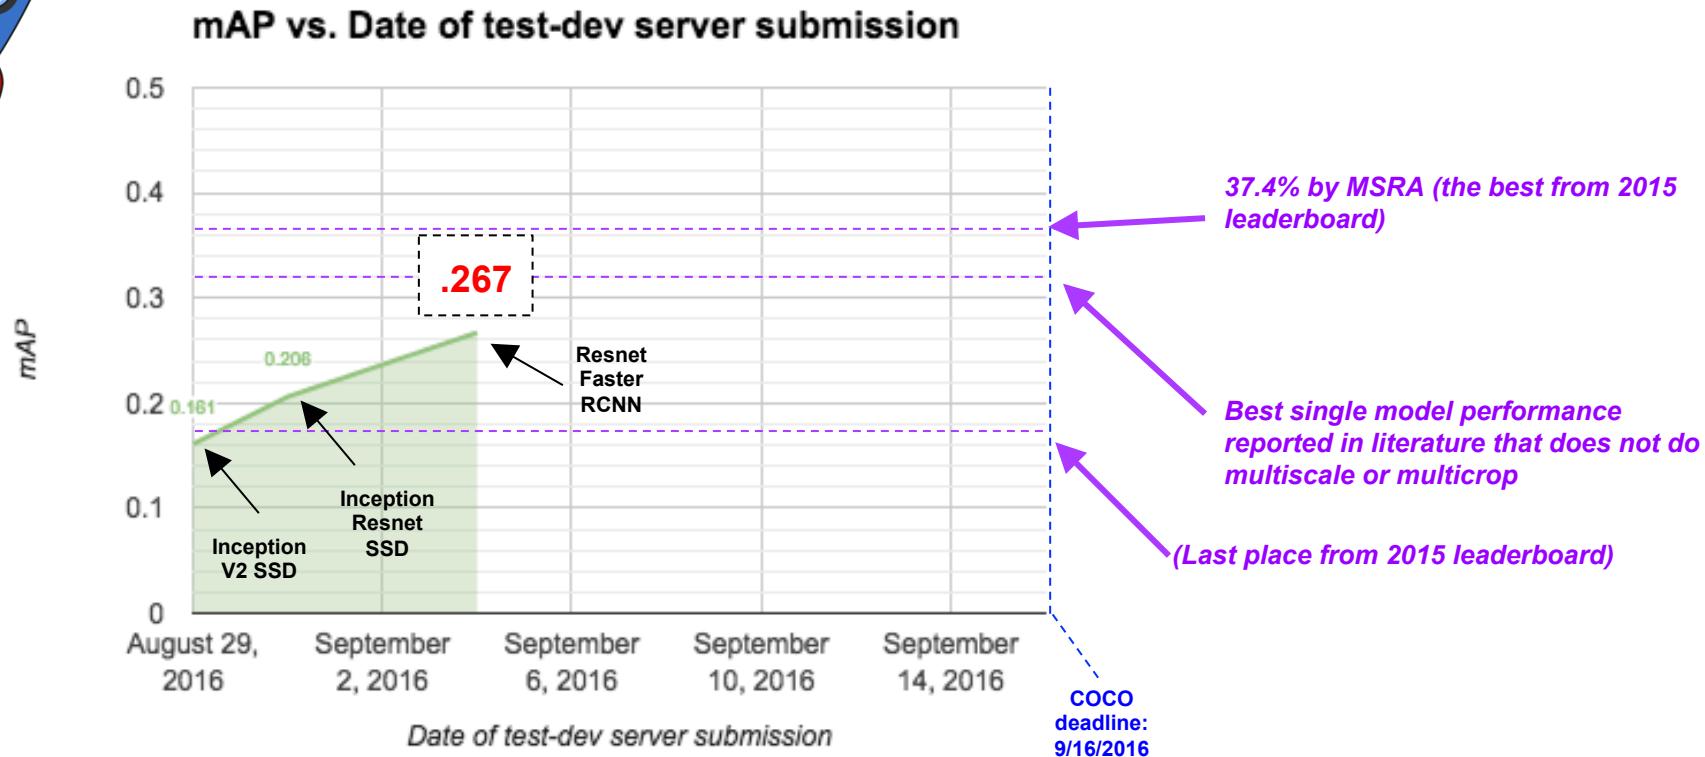
H. Michael Miley

### Inception Resnet SSD

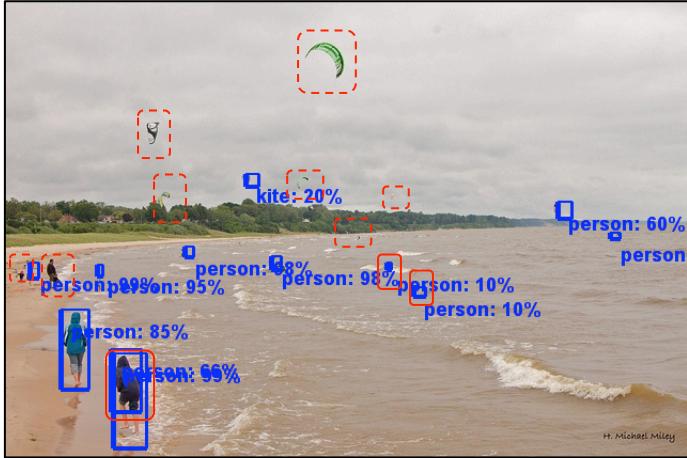


inaccurate localization  
for small objects

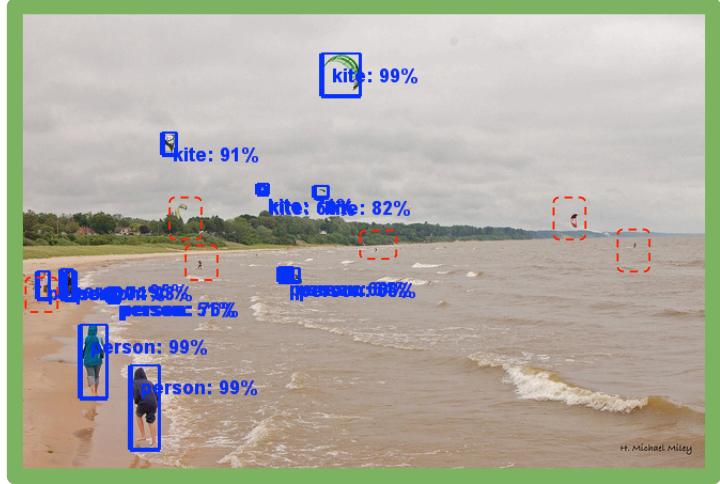
# Race to the Top



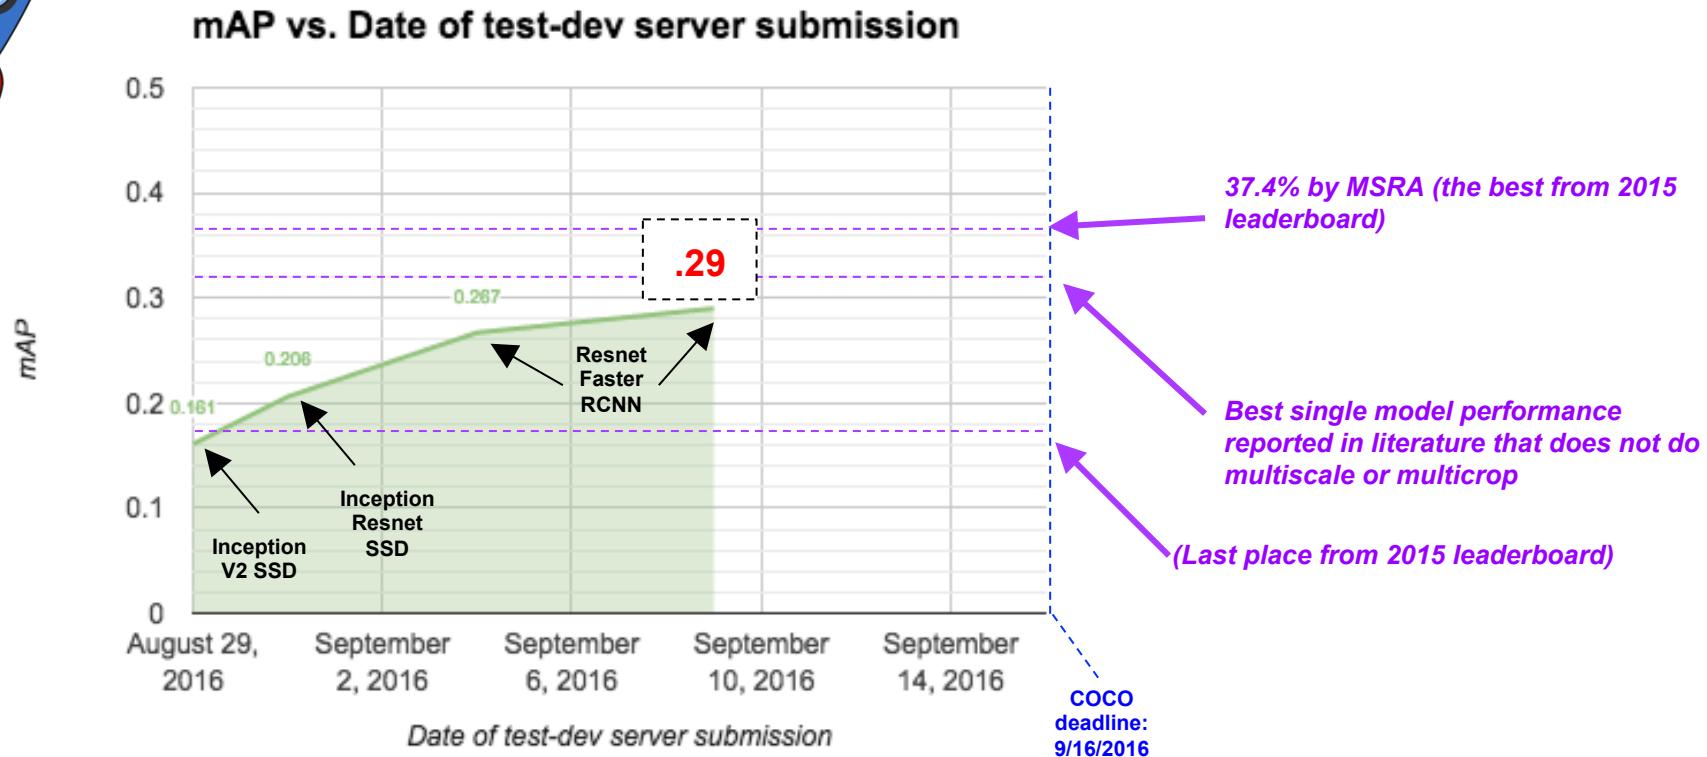
### Inception Resnet SSD



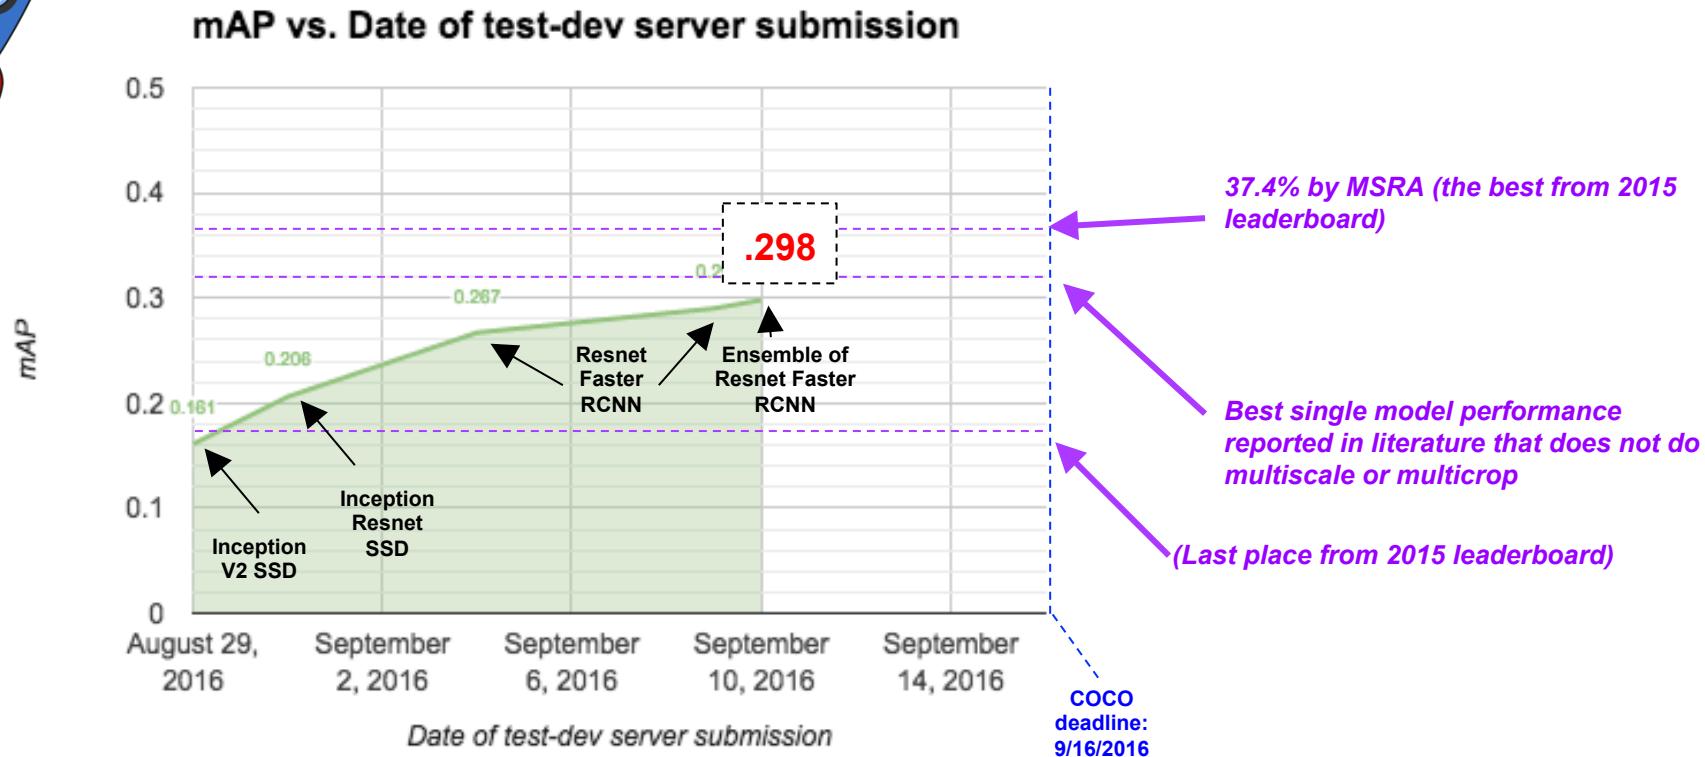
### Resnet Faster RCNN



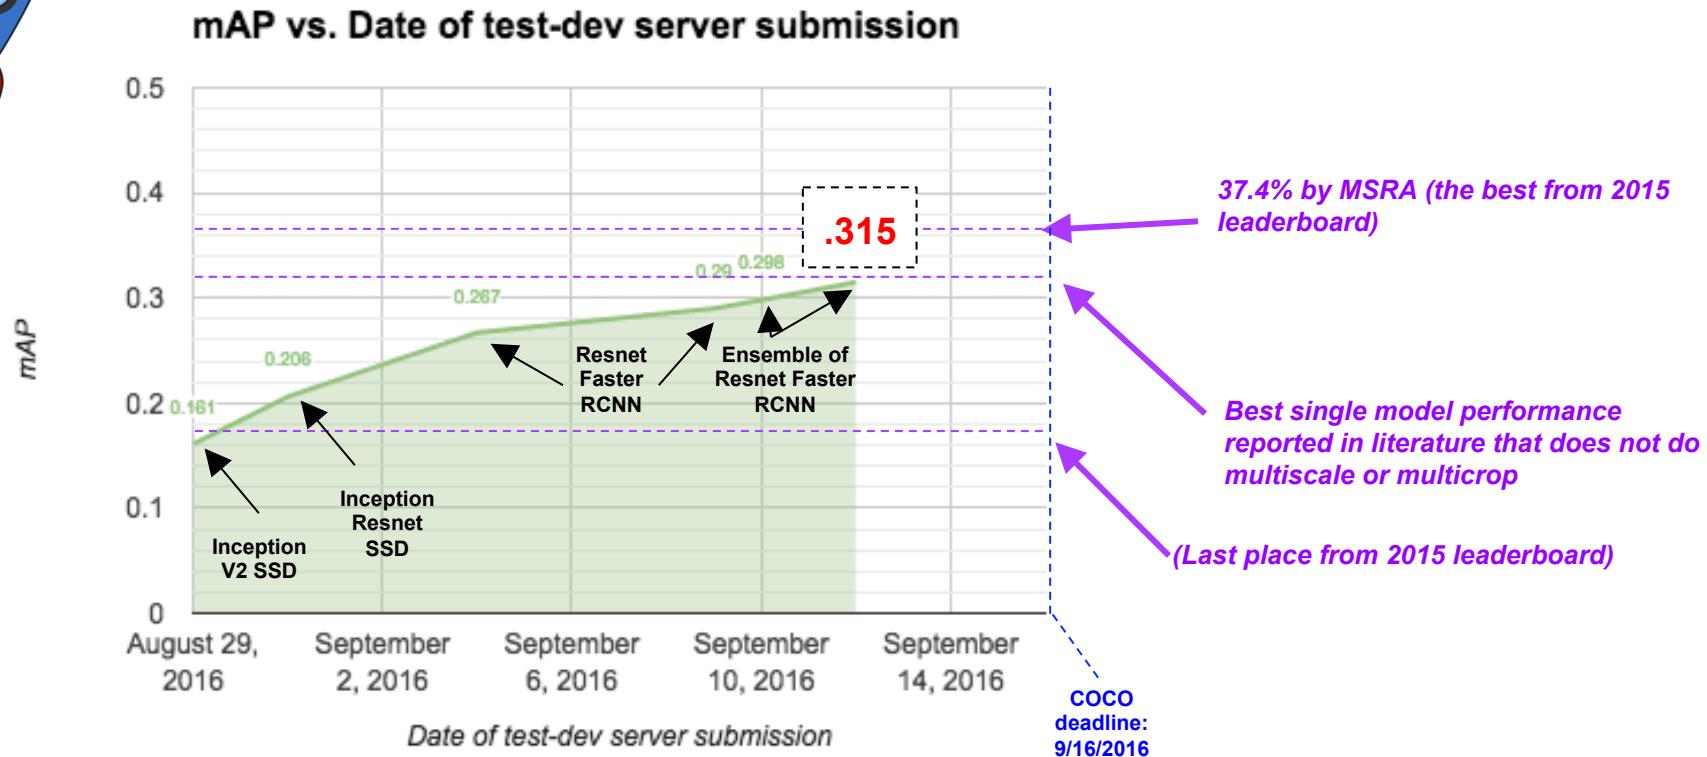
# Race to the Top



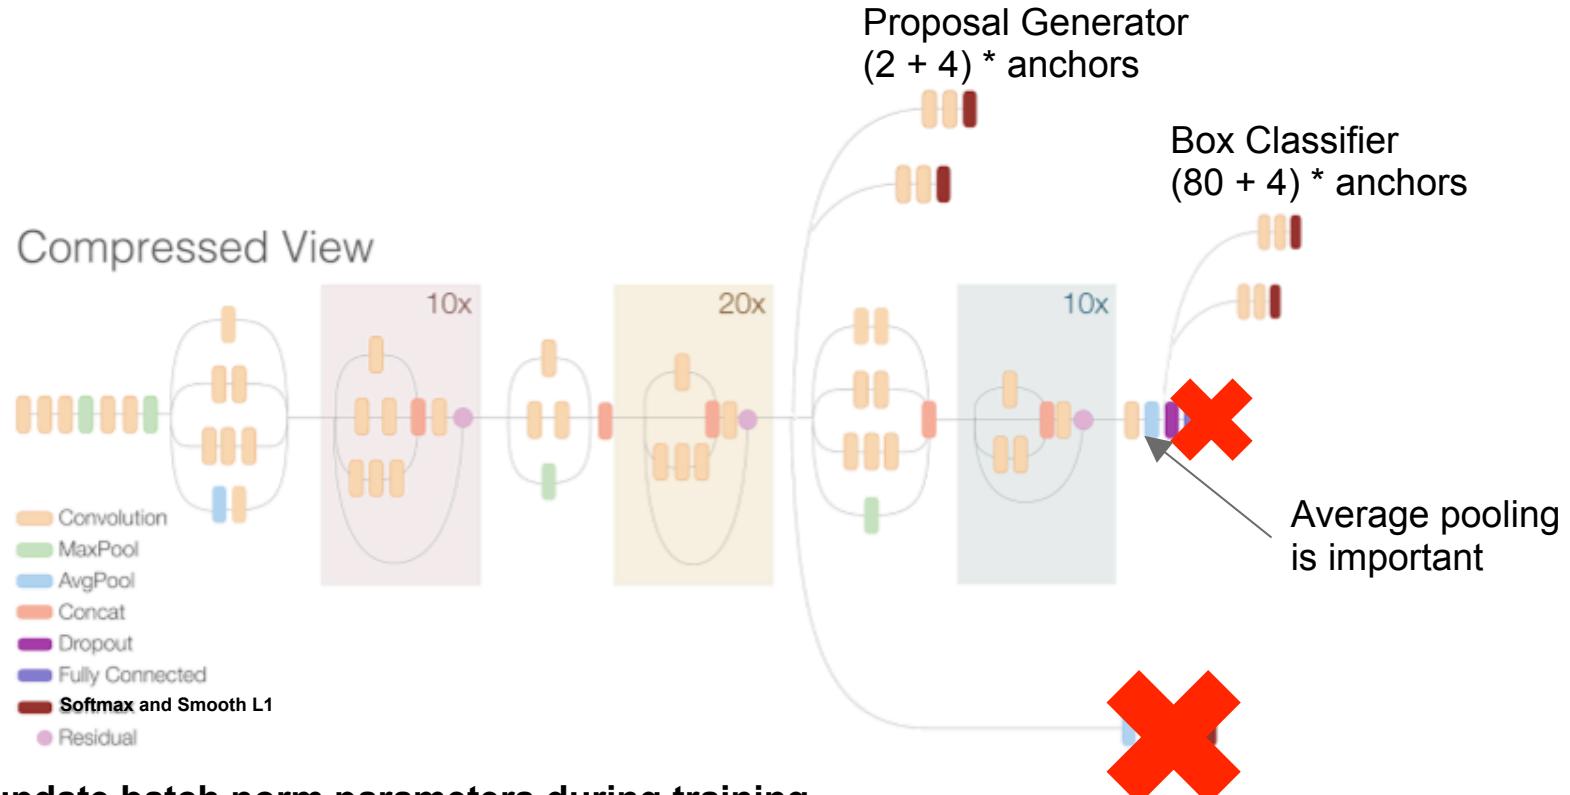
# Race to the Top



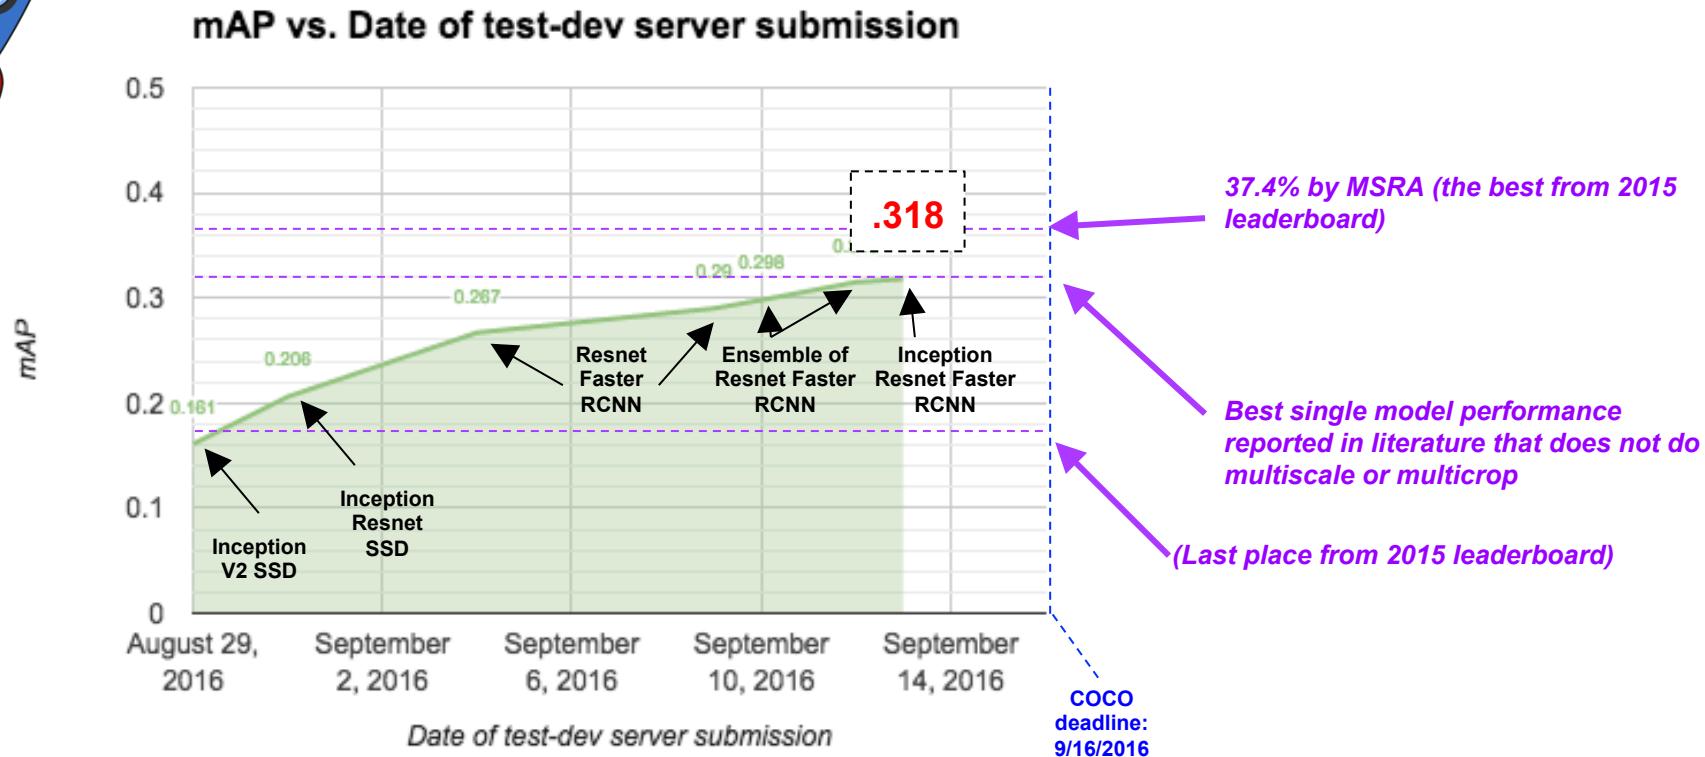
# Race to the Top



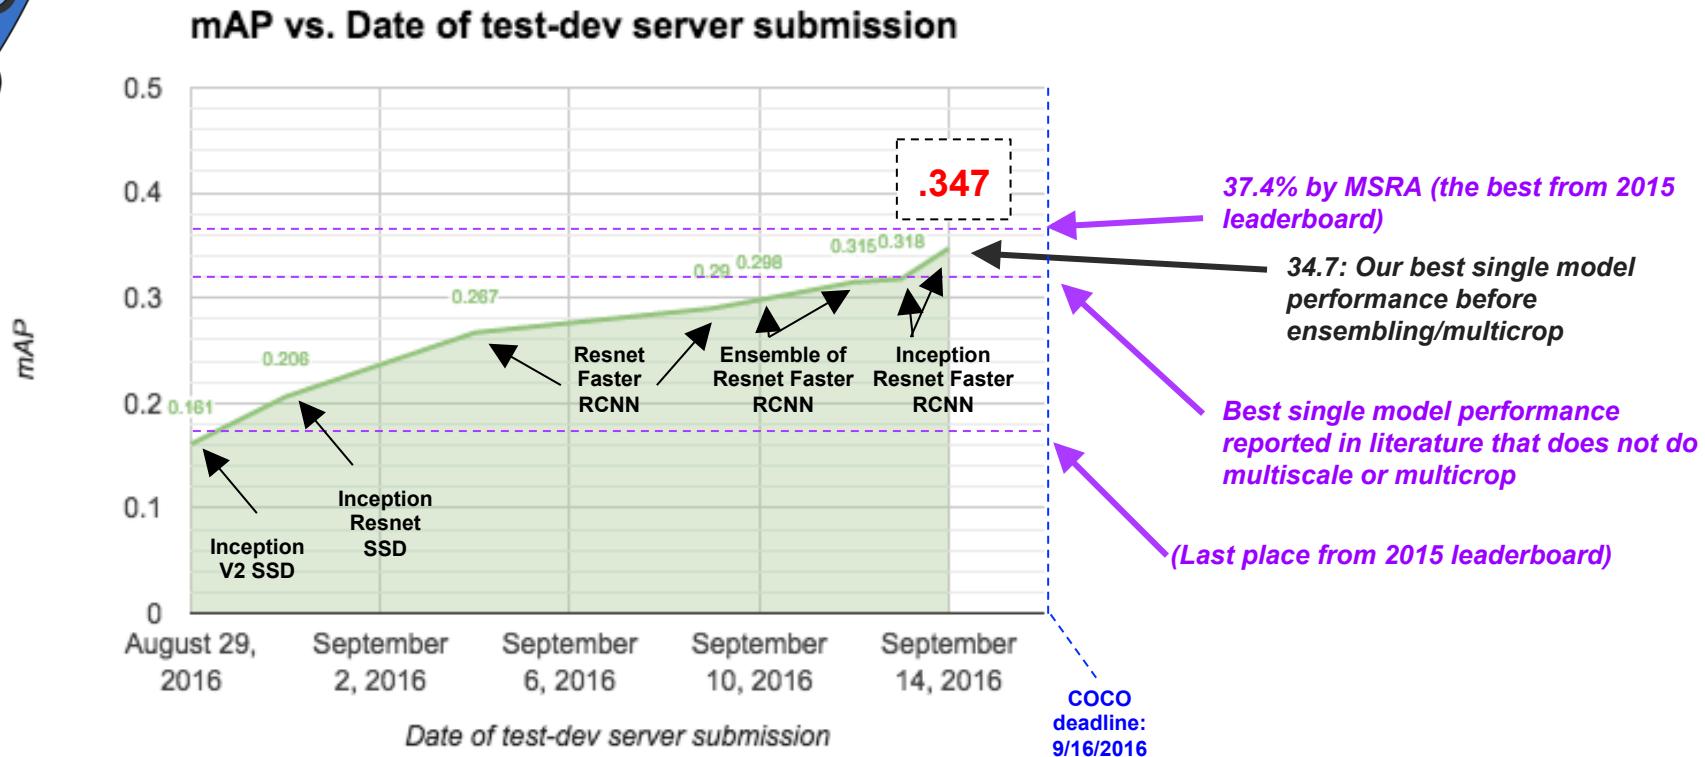
# Faster RCNN w/Inception Resnet (v2)



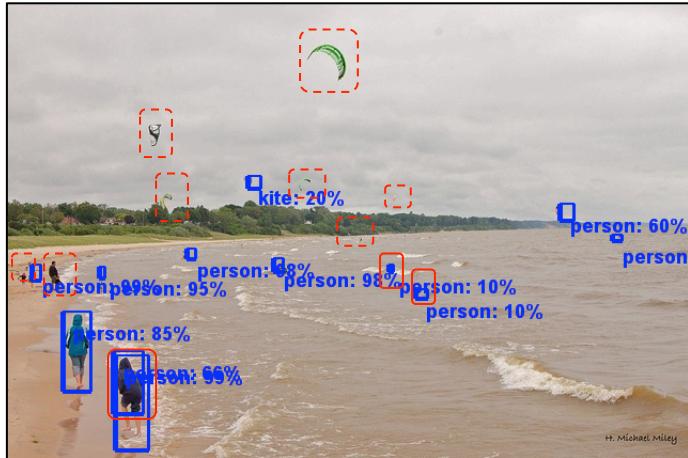
# Race to the Top



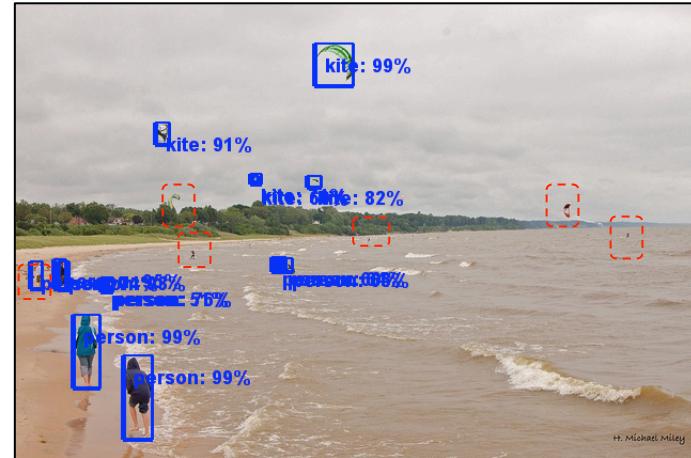
# Race to the Top



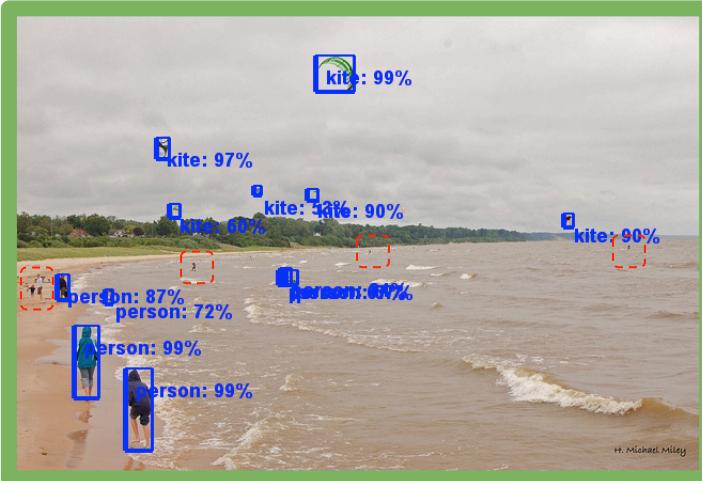
### Inception Resnet SSD



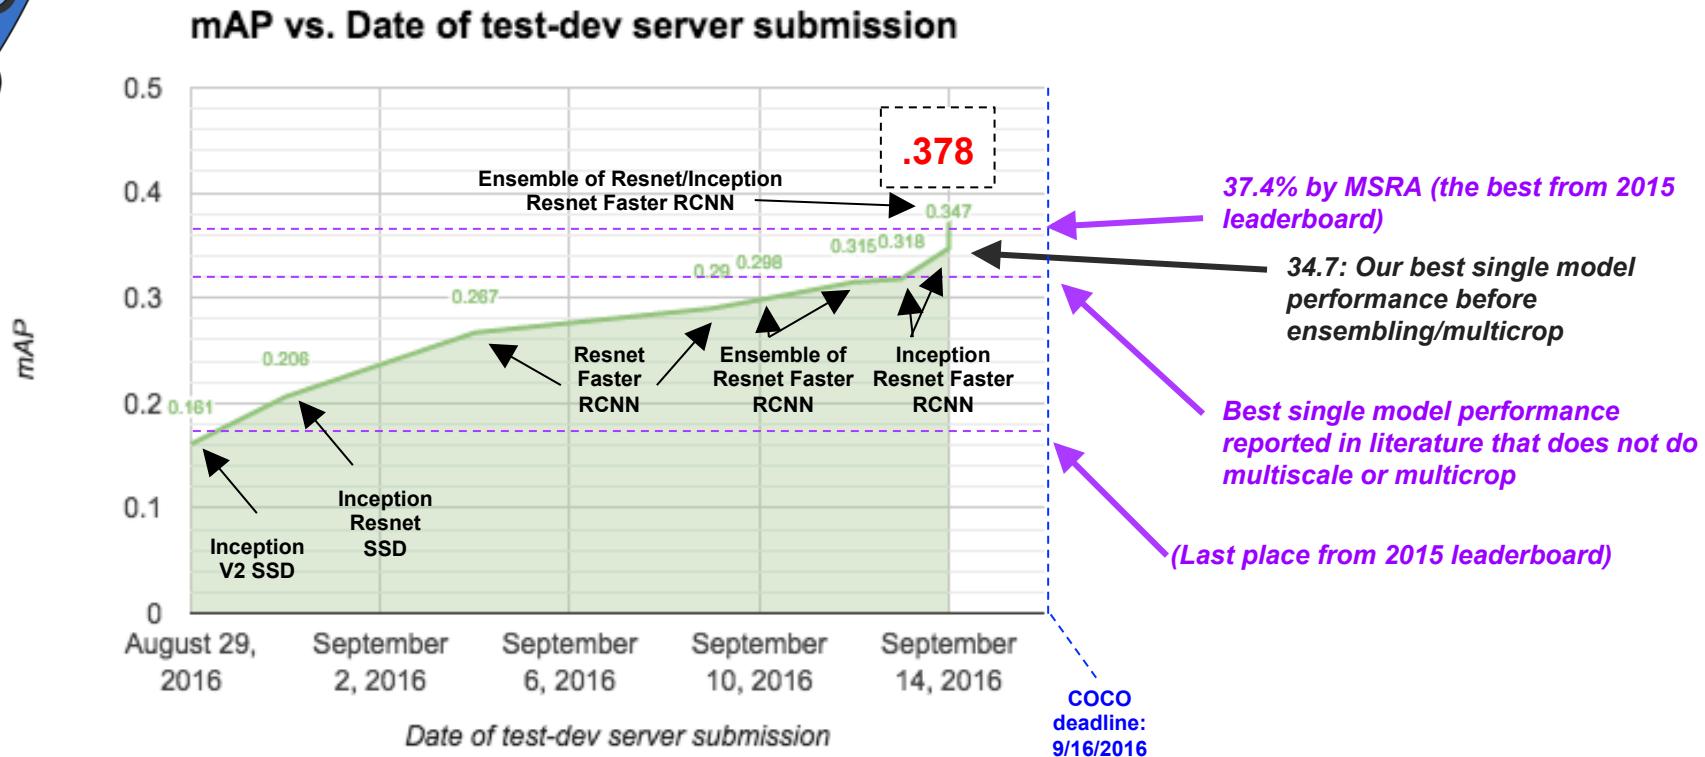
### Resnet Faster RCNN



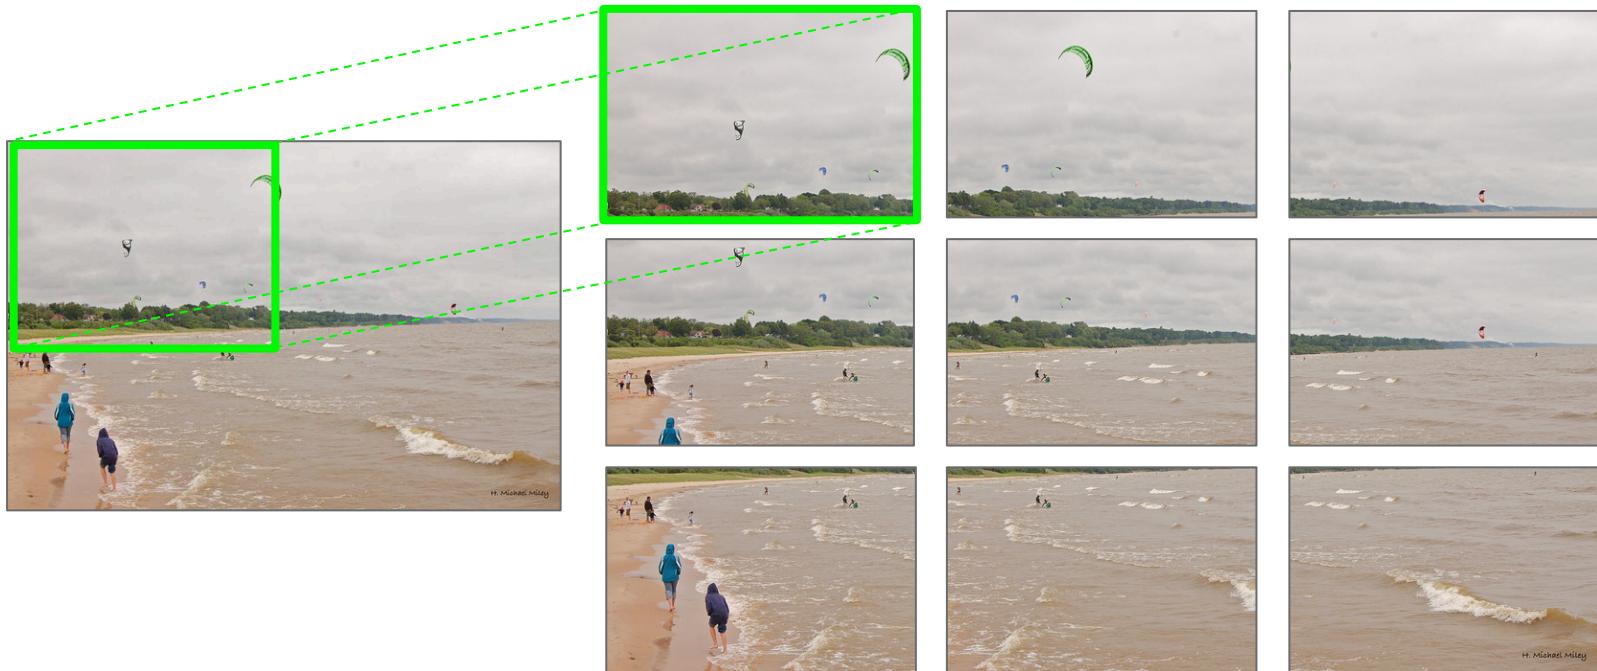
### Inception Resnet Faster RCNN



# Race to the Top

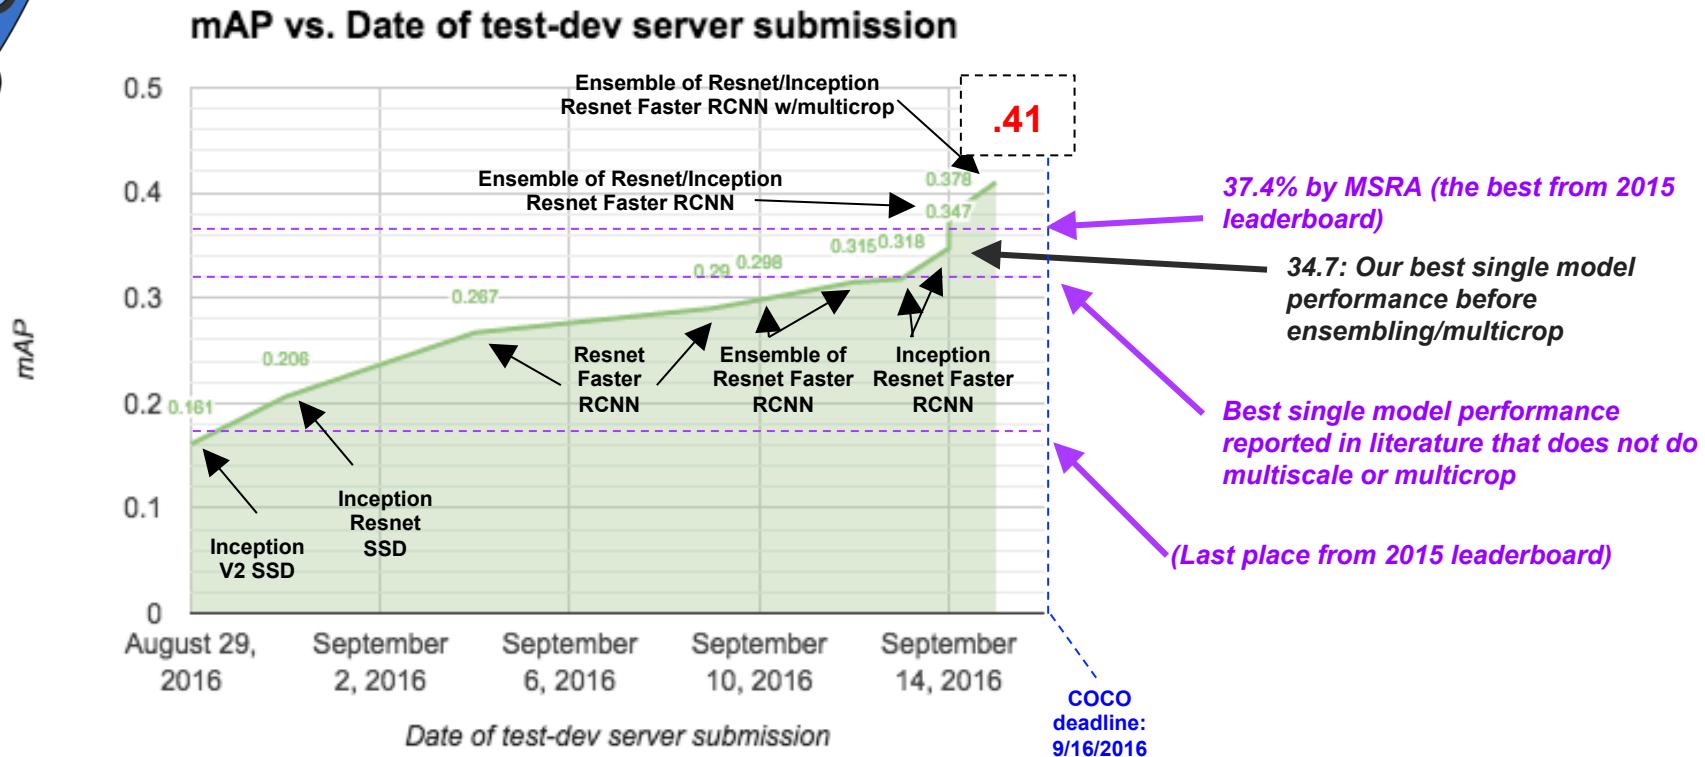


# 10-crop inference



No multiscale training, horizontal flip, box refinement, box voting, global context or ILSVRC detection data

# Race to the Top



	A	B	C	D	E	F	G	H
1	Job ID	Copy I	Feature Extractor	Feature Crop S	Loss Weights	LR	LR Boundary	
2	1	1	inception_resnet	16 34:2	1:1:1:1	.001,.0001,.00001	8,000,001,000,000	
3	1	2	inception_resnet	16 34:2	1:1:1:1	.001,.0001,.00001	8,000,001,000,000	
4	1	3	inception_resnet	16 34:2	1:1:1:1	.001,.0001,.00001	8,000,001,000,000	
5	2	1	inception_resnet	16 17:1	1:1:1:1	.001,.0001,.00001	8,000,001,000,000	
6	2	2	inception_resnet	16 17:1	1:1:1:1	.001,.0001,.00001	8,000,001,000,000	
7	2	3	inception_resnet	16 17:1	1:1:1:1	.001,.0001,.00001	8,000,001,000,000	
8	3	1	inception_resnet	16 34:2	2:1:2:1	.0008,.00008,.000008	8,000,001,000,000	
9	3	2	inception_resnet	16 34:2	2:1:2:1	.0008,.00008,.000008	8,000,001,000,000	
10	3	3	inception_resnet	16 34:2	2:1:2:1	.0008,.00008,.000008	8,000,001,000,000	
11	4	1	inception_resnet	16 17:1	2:1:2:1	.0008,.00008,.000008	8,000,001,000,000	
12	4	2	inception_resnet	16 17:1	2:1:2:1	.0008,.00008,.000008	8,000,001,000,000	
13	4	3	inception_resnet	16 17:1	2:1:2:1	.0008,.00008,.000008	8,000,001,000,000	
14	5	1	inception_resnet	16 34:2	1:1:1:1	.001,.0001,.00001	10,000,001,200,000	
15	5	2	inception_resnet	16 34:2	1:1:1:1	.001,.0001,.00001	10,000,001,200,000	
16	5	3	inception_resnet	16 34:2	1:1:1:1	.001,.0001,.00001	10,000,001,200,000	
17	6	1	inception_resnet	16 17:1	1:1:1:1	.001,.0001,.00001	10,000,001,200,000	
18	6	2	inception_resnet	16 17:1	1:1:1:1	.001,.0001,.00001	10,000,001,200,000	
19	6	3	inception_resnet	16 17:1	1:1:1:1	.001,.0001,.00001	10,000,001,200,000	
20	7	1	inception_resnet	16 34:2	2:1:2:1	.0008,.00008,.000008	10,000,001,200,000	
21	7	2	inception_resnet	16 34:2	2:1:2:1	.0008,.00008,.000008	10,000,001,200,000	
22	7	3	inception_resnet	16 34:2	2:1:2:1	.0008,.00008,.000008	10,000,001,200,000	
23	8	1	inception_resnet	16 17:1	2:1:2:1	.0008,.00008,.000008	10,000,001,200,000	
24	8	2	inception_resnet	16 17:1	2:1:2:1	.0008,.00008,.000008	10,000,001,200,000	
25	8	3	inception_resnet	16 17:1	2:1:2:1	.0008,.00008,.000008	10,000,001,200,000	
26	9	1	resnet101	8 28:4	3:1:3:1	.0002,.00002,.000002	8,000,001,000,000	
27	9	2	resnet101	8 28:4	3:1:3:1	.0002,.00002,.000002	8,000,001,000,000	
28	9	3	resnet101	8 28:4	3:1:3:1	.0002,.00002,.000002	8,000,001,000,000	
29	10	1	resnet101	8 14:2	3:1:3:1	.0002,.00002,.000002	8,000,001,000,000	
30	10	2	resnet101	8 14:2	3:1:3:1	.0002,.00002,.000002	8,000,001,000,000	
31	10	3	resnet101	8 14:2	3:1:3:1	.0002,.00002,.000002	8,000,001,000,000	
32	11	1	resnet101	8 28:4	1:1:1:1	.0003,.00003,.000003	8,000,001,000,000	
33	11	2	resnet101	8 28:4	1:1:1:1	.0003,.00003,.000003	8,000,001,000,000	

# Many Final Experiments

- Best layer of Inception Resnet for RPN?
- Use atrous convolution for dense output?
- Maximize IOU vs Minimize SmoothL1?
- Best learning rate decay schedule?
- Best data augmentation operations (e.g. random cropping, random saturation, random contrast)?

... led to *Model Checkpoint Overload* :(

# Model Selection for Ensembling

checkpoints				
	Name	Size	Type	Date Modified
▼ My Computer	exp_1_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_1_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_1_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_2_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_2_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_2_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_3_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_3_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_3_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_4_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_4_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_4_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_5_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_5_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_5_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_6_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_6_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_6_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_7_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_7_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_7_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_8_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_8_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
	exp_8_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
▼ Network				
	Network			

**Take best K models?  
Or select diverse K-subset of models?**

# Model Selection for Ensembling

The screenshot shows a file explorer window titled "checkpoints". The left sidebar lists "My Computer" and "Network" sections. The main area displays a list of checkpoints, each with a thumbnail icon, name, size (480.8 MB), type (Binary), and date modified (Tue 27 Sep 2016 03:31:01 PM PDT). A black rounded rectangle highlights a table overlaid on the list.

	Model 1 mAP	Model 2 mAP	Model 3 mAP
Car	20%	23%	70%
Dog	81%	80%	15%
Bear	78%	81%	20%
Chair	10%	12%	71%

Take best K models?  
Or select diverse K-subset of models?

# Model Selection for Ensembling

The screenshot shows a file explorer window titled "checkpoints". The left sidebar lists "My Computer" and "Network" sections. Under "My Computer", there are numerous files named "exp\_1\_copy1.ckpt-1500000" through "exp\_8\_copy3.ckpt-1500000", all of which are 480.8 MB binary files modified on Tuesday, September 27, 2016, at 03:31:01 PM PDT. A red rounded rectangle highlights a table titled "Similar Models" in the foreground. The table compares three models (Model 1, Model 2, Model 3) across four categories: Car, Dog, Bear, and Chair. Model 1 has mAP values of 20%, 81%, 78%, and 10% respectively. Model 2 has 23%, 80%, 81%, and 12%. Model 3 has 70%, 15%, 20%, and 71%. Green checkmarks are placed next to the highest mAP values for each category.

	Model 1 mAP	Model 2 mAP	Model 3 mAP
Car	20%	23%	70%
Dog	81%	80%	15%
Bear	78%	81%	20%
Chair	10%	12%	71%

Take best K models?  
Or select diverse K-subset of models?

# Model Selection for Ensembling

checkpoints

File Edit View Go Bookmarks Help

< > ^ | > .. J M E E E

**My Computer**

- Home
- Desktop
- Documents
- Music
- Pictures
- Videos
- Downloads
- Recent
- File System
- Trash

**Network**

- Network

Name	Size	Type	Date Modified
exp_1_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_1_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_1_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_2_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_2_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_2_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_3_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_3_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_3_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_4_copy1.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_4_copy2.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_4_copy3.ckpt-1500000	480.8 MB	Binary	Tue 27 Sep 2016 03:31:01 PM PDT
exp_5_copy1.ckpt-150			
exp_5_copy2.ckpt-150			
exp_5_copy3.ckpt-150			
exp_6_copy1.ckpt-150			
exp_6_copy2.ckpt-150			
exp_6_copy3.ckpt-150			
exp_7_copy1.ckpt-150			
exp_7_copy2.ckpt-150			
exp_7_copy3.ckpt-150			
exp_8_copy1.ckpt-150			
exp_8_copy2.ckpt-150			
exp_8_copy3.ckpt-150			

**Complementary Models**

	Model 1 mAP	Model 2 mAP	Model 3 mAP
Car	20%	23%	70%
Dog	81%	80%	15%
Bear	78%	81%	20%
Chair	10%	12%	71%

Take best K models?  
Or select diverse K-subset of models?

# Diversity Matters

**Model NMS for diverse ensembling:** Greedily select diverse model collection for ensembling, pruning away models too similar to already selected models.

Final ensemble selected for challenge submission				
Individual mean AP (on minival)	Feature Extractor	Output Stride	Location:Classification loss ratio	Location Loss function
32.93	Resnet 101	8	3:1	SmoothL1
33.3	Resnet 101	8	1:1	SmoothL1
34.75	Inception Resnet	16	1:1	SmoothL1
35	Inception Resnet	16	2:1	SmoothL1+IOU
35.64	Inception Resnet	8	1:1	SmoothL1

# Diversity Matters

**Model NMS for diverse ensembling:** Greedily select diverse model collection for ensembling, pruning away models too similar to already selected models.

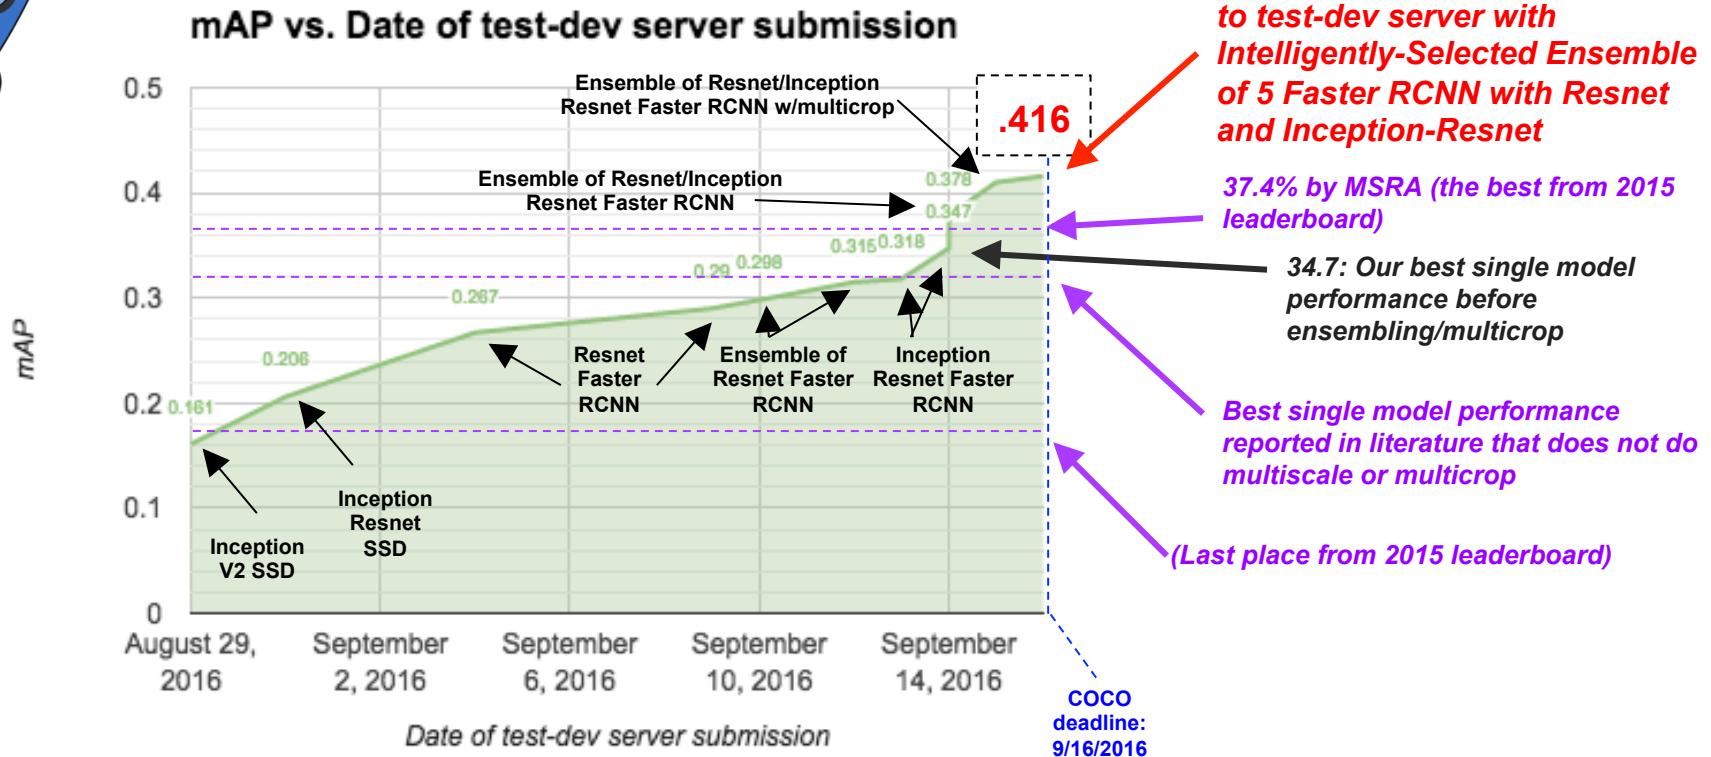
Final ensemble selected for challenge submission				
Individual mean AP (on minival)	Feature Extractor	Output Stride	Location:Classification loss ratio	Location Loss function
32.93	Resnet 101	8	3:1	SmoothL1
33.3	Resnet 101	8	1:1	SmoothL1
34.75	Inception Resnet	16	1:1	SmoothL1
35	Inception Resnet	16	2:1	SmoothL1+IOU
35.64	Inception Resnet	8	1:1	SmoothL1

# Diversity Matters

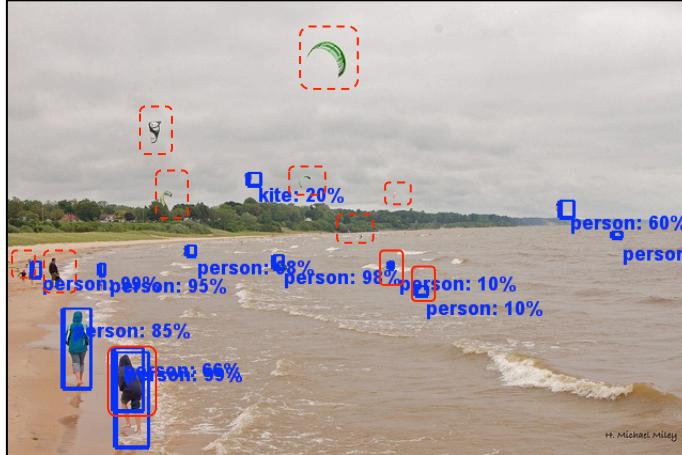
**Model NMS for diverse ensembling:** Greedily select diverse model collection for ensembling, pruning away models too similar to already selected models.

Final ensemble selected for challenge submission				
Individual mean AP (on minival)	Feature Extractor	Output Stride	Location:Classification loss ratio	Location Loss function
32.93	Resnet 101	8	3:1	SmoothL1
33.3	Resnet 101	8	1:1	SmoothL1
34.75	Inception Resnet	16	1:1	SmoothL1
35	Inception Resnet	16	2:1	SmoothL1+IOU
35.64	Inception Resnet	8	1:1	SmoothL1

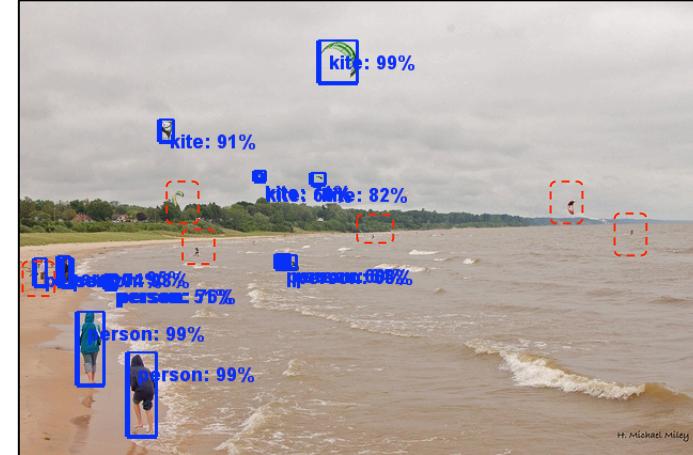
# Race to the Top



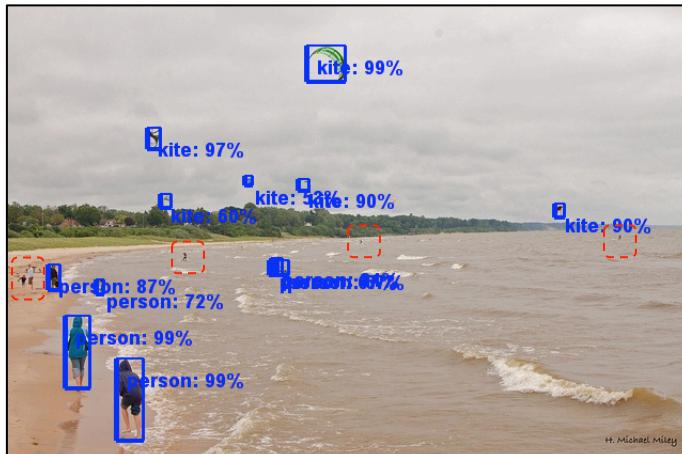
### Inception Resnet SSD



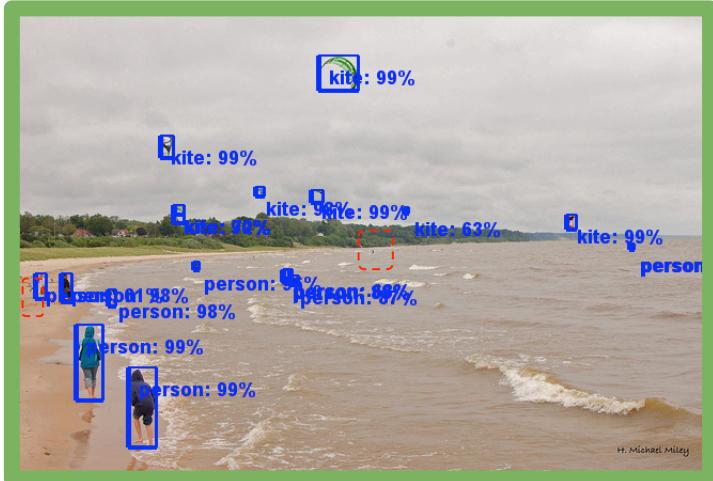
### Resnet Faster RCNN



### Inception Resnet Faster RCNN



### Final ensemble with multicrop inference



# This Talk

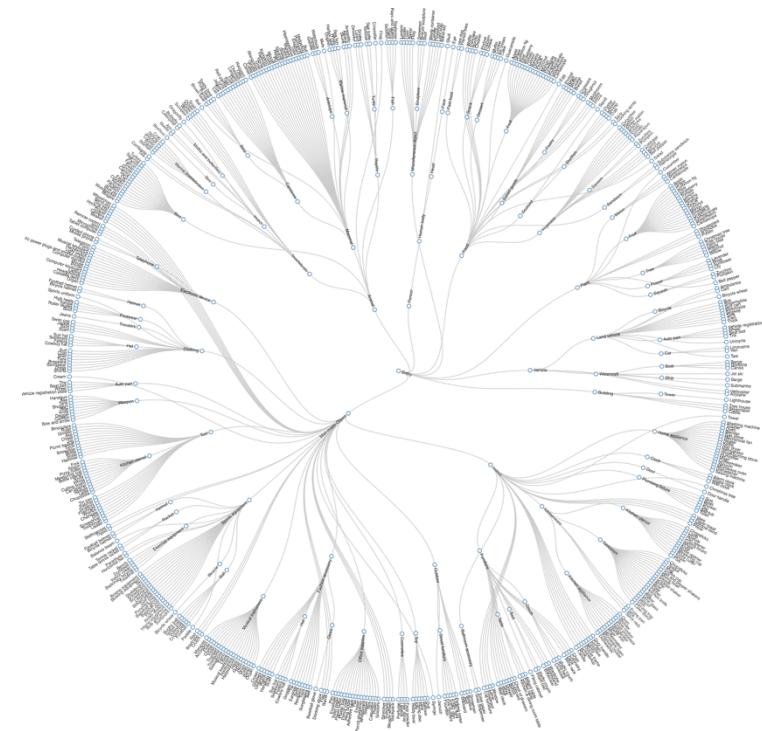
- The Tensorflow Object Detection API
- **Pushing the envelope on Accuracy**
  - Better Models
  - **Better Data**
- Pushing the envelope on Speed
  - Lightweight Models
  - Adaptive Inference

# JFT-300M: a massive dataset from ImageSearch data

- 300M web images
- 375M image label pairs

# JFT-300M: a massive dataset from ImageSearch data

- 300M web images
- 375M image label pairs
- ~ 19K categories



# JFT-300M: a massive dataset from ImageSearch data

- 300M web images
- 375M image label pairs
- ~ 18K categories
- ~ 20% label noise
- Unknown recall
- Long-tail distribution

Tortoise:

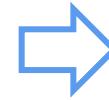
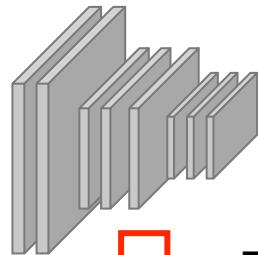
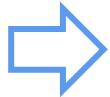


vs.



# Revisiting Unreasonable Effectiveness of Data in Deep Learning Era [Sun et al, '17]

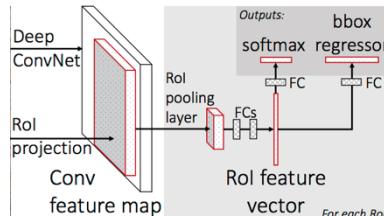
JFT 300M



18K labels

 **PASCAL2**  
Pattern Analysis, Statistical Modelling and  
Computational Learning

 **coco**  
Common Objects in Context

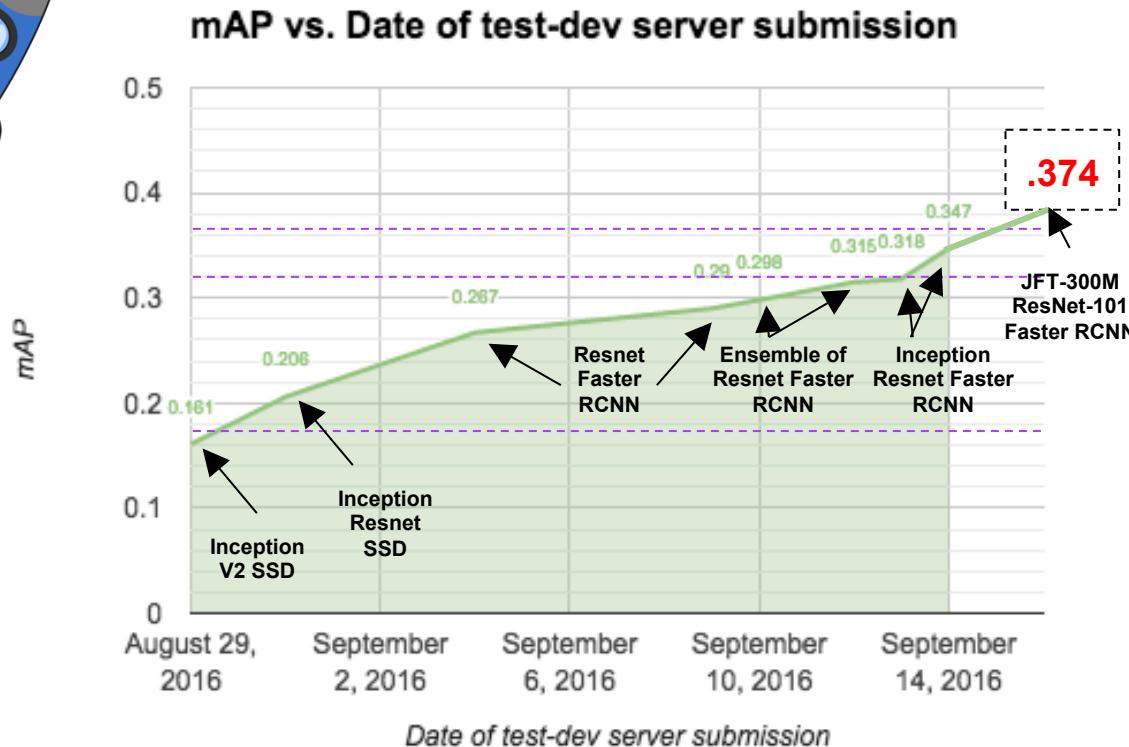


Transfer weights



Detections

# Better Representation Learning Helps!

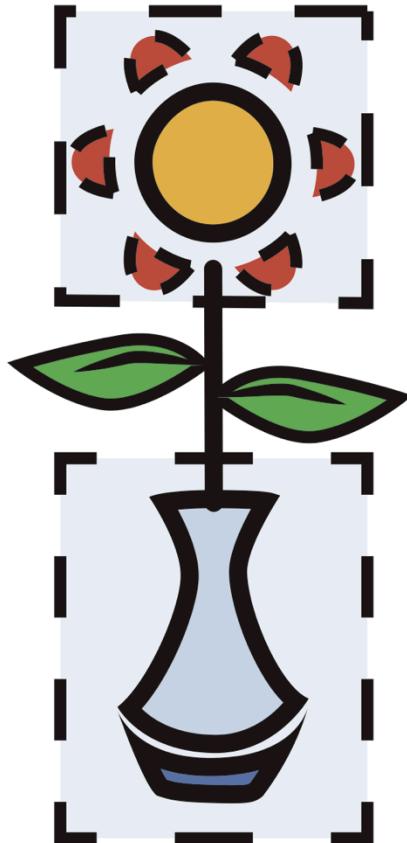


Using a JFT-300M pre-trained checkpoint to replace ImageNet ones:

- 2.7% gain over best single model
- 3.1% gain over comparable ResNet model

# Open Images

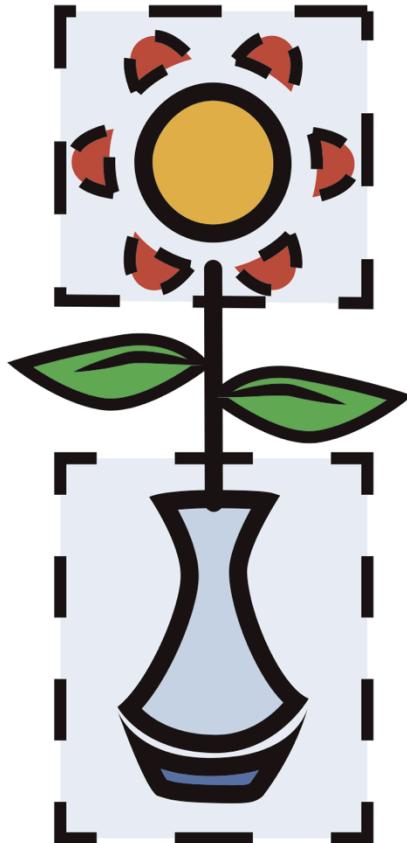
[GITHUB.COM/OPENIMAGES](https://github.com/openimages)



- **5000 classes**
- Validation and Test sets:
  - **41k images**
- Training set:
  - **9M images**

# Open Images

GITHUB.COM/OPENIMAGES



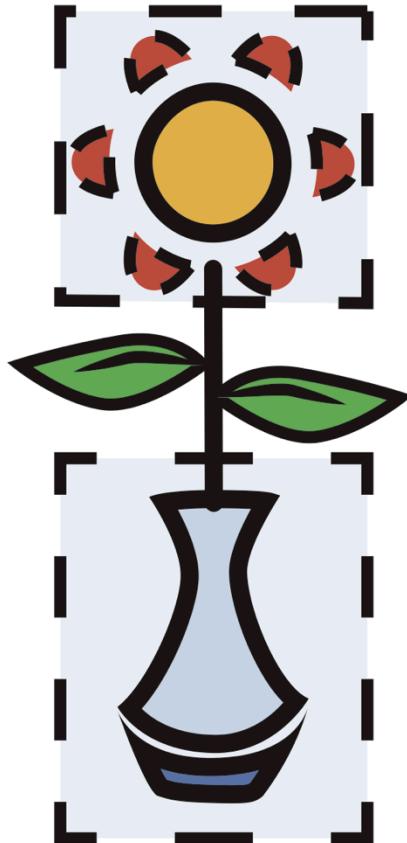
balcony, stairs, facade, iron, door, interior design, gate, architecture, handrail, baluster, window, arch



cutlery, tableware, metal, tool, spoon, fork

# Open Images

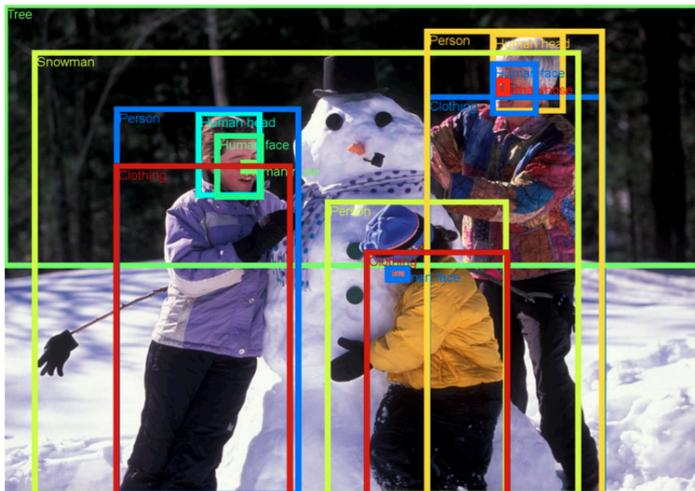
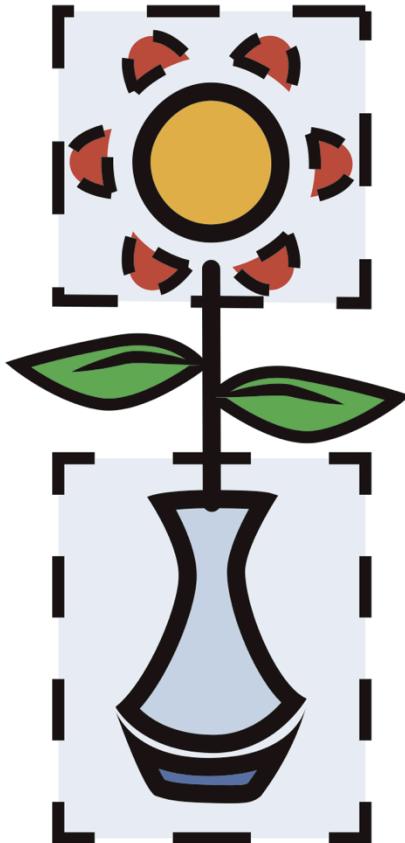
[GITHUB.COM/OPENIMAGES](https://github.com/openimages)



- **600 boxable object classes**
- Validation and Test sets:
  - **830k bounding-boxes on 167k images**
  - all object instances manually boxed
- Training set:
  - **1.2M bounding-boxes**
  - 1 box per class per image
  - produced by human verification technique

# Open Images

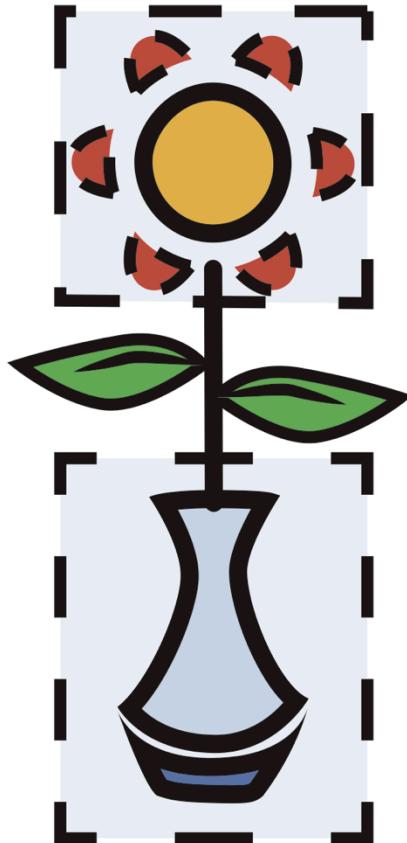
GITHUB.COM/OPENIMAGES



Annotated images from the Open Images dataset. Left: [FAMILY MAKING A SNOWMAN](#) by [mwvchamber](#). Right: [STANZA STUDENTI.S.S. ANNUNZIATA](#) by [ersupalermo](#). Both images used under [CC BY 2.0](#) license.

# Open Images

GITHUB.COM/OPENIMAGES



*Stay tuned for pre-trained  
Tensorflow Object Detection  
models on the Open Images  
dataset...*



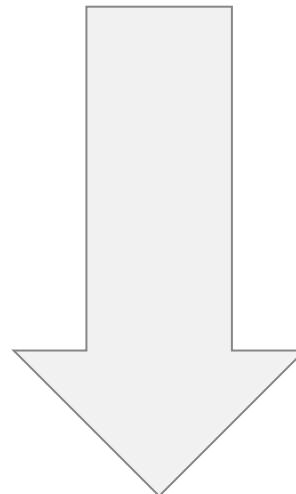
"Krasin I., Duerig T., Alldrin N., Ferrari V., Abu-El-Haija S., Kuznetsova A., Rom H., Uijlings J., Popov S., Veit A., Belongie S., Gomes V., Gupta A., Sun C., Chechik G., Cai D., Feng Z., Narayanan D., Murphy K. OpenImages: A public dataset for large-scale multi-label and multi-class image classification, 2017. Available from <https://github.com/openimages>".

# This Talk

- The Tensorflow Object Detection API
- Pushing the envelope on Accuracy
  - Better Models
  - Better Data
- **Pushing the envelope on Speed**
  - **Lightweight Models**
  - Adaptive Inference

# Vision Models are Getting Bigger and Better

- AlexNet
- GoogleNet, VGG
- Inception V2, V3
- ResNet
- Inception ResNet



Bigger and Better

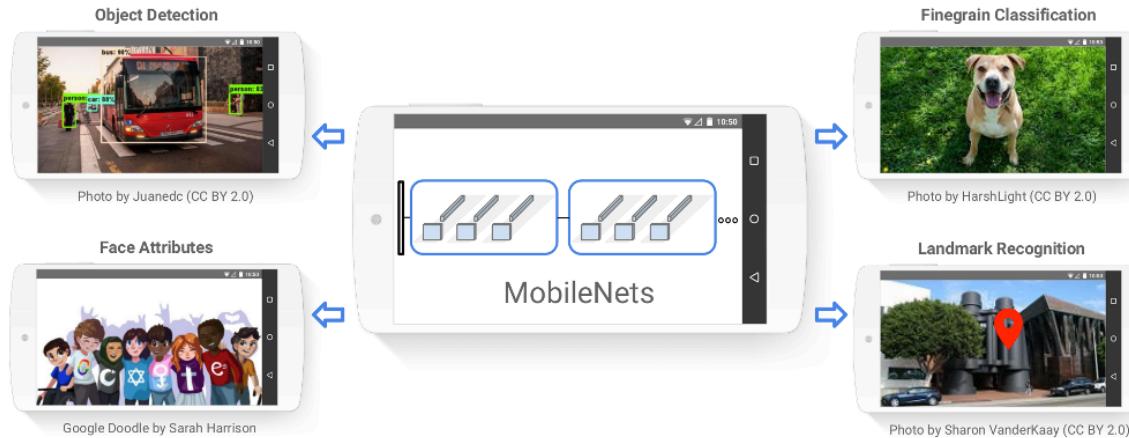


Example:  
GoogleNet

[slide credit: Andrew Howard]

# Mobile Vision at Google

**Long-Term Goal: Enable Computer Vision on Mobile Devices.**



This means optimizing for:

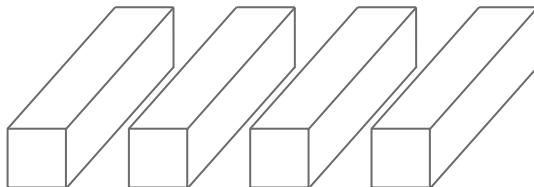
- Latency
- Power Consumption
- Memory Footprint
- Disk Footprint - Binary and Model Size

[slide credit: Andrew Howard]

# Convolution vs Depthwise Separable Convolutions

**Convolution:**

**Filter and Combine in 1 Step:**



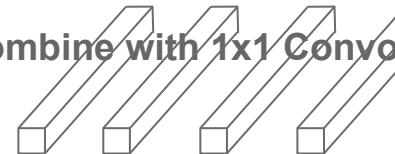
**Depthwise Separable Convolution:**

**Filter and Combine in 2 Steps:**

**Filter with Depthwise Convolutions:**

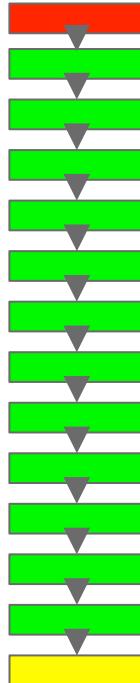


**Combine with 1x1 Convolutions:**



- Breaks the computational dependence between the size of filters and the number of output channels.
- In practice for 3x3 Convolution, saves ~8x computation at the cost of ~1% accuracy.

# MobileNet Architecture



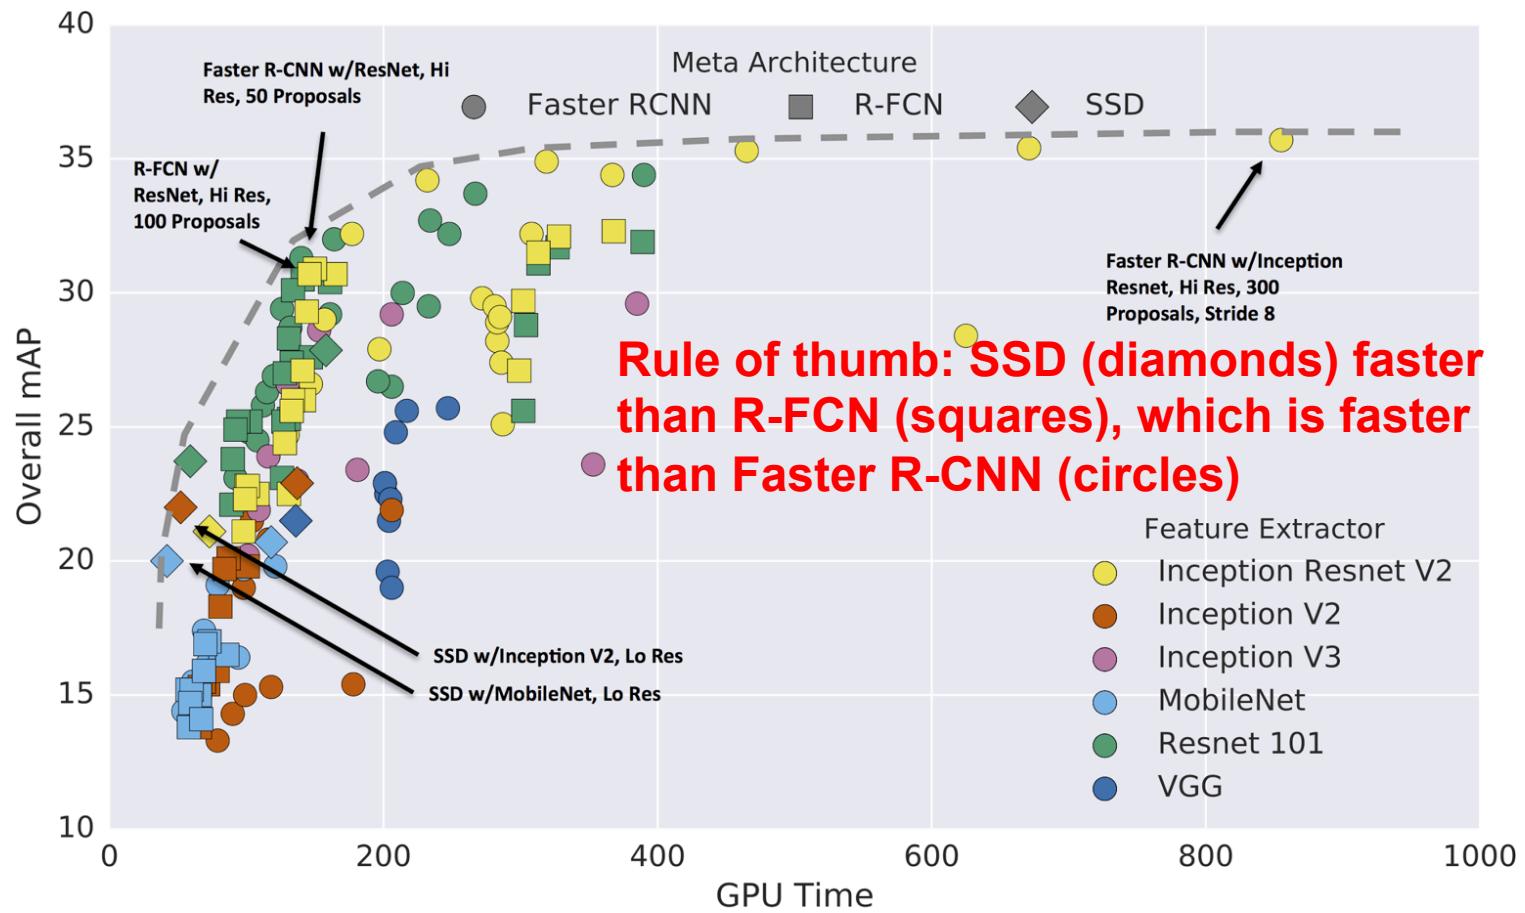
Type	Output Depth	Output Resolution
Convolution	32	112
Separable Convolution	64	112
Separable Convolution	128	56
Separable Convolution	128	56
Separable Convolution	256	28
Separable Convolution	256	28
Separable Convolution	512	14
Avg Pool + FC	1000	1

- 95% of computation is 1x1 convolutions efficiently implemented with GEMMs.

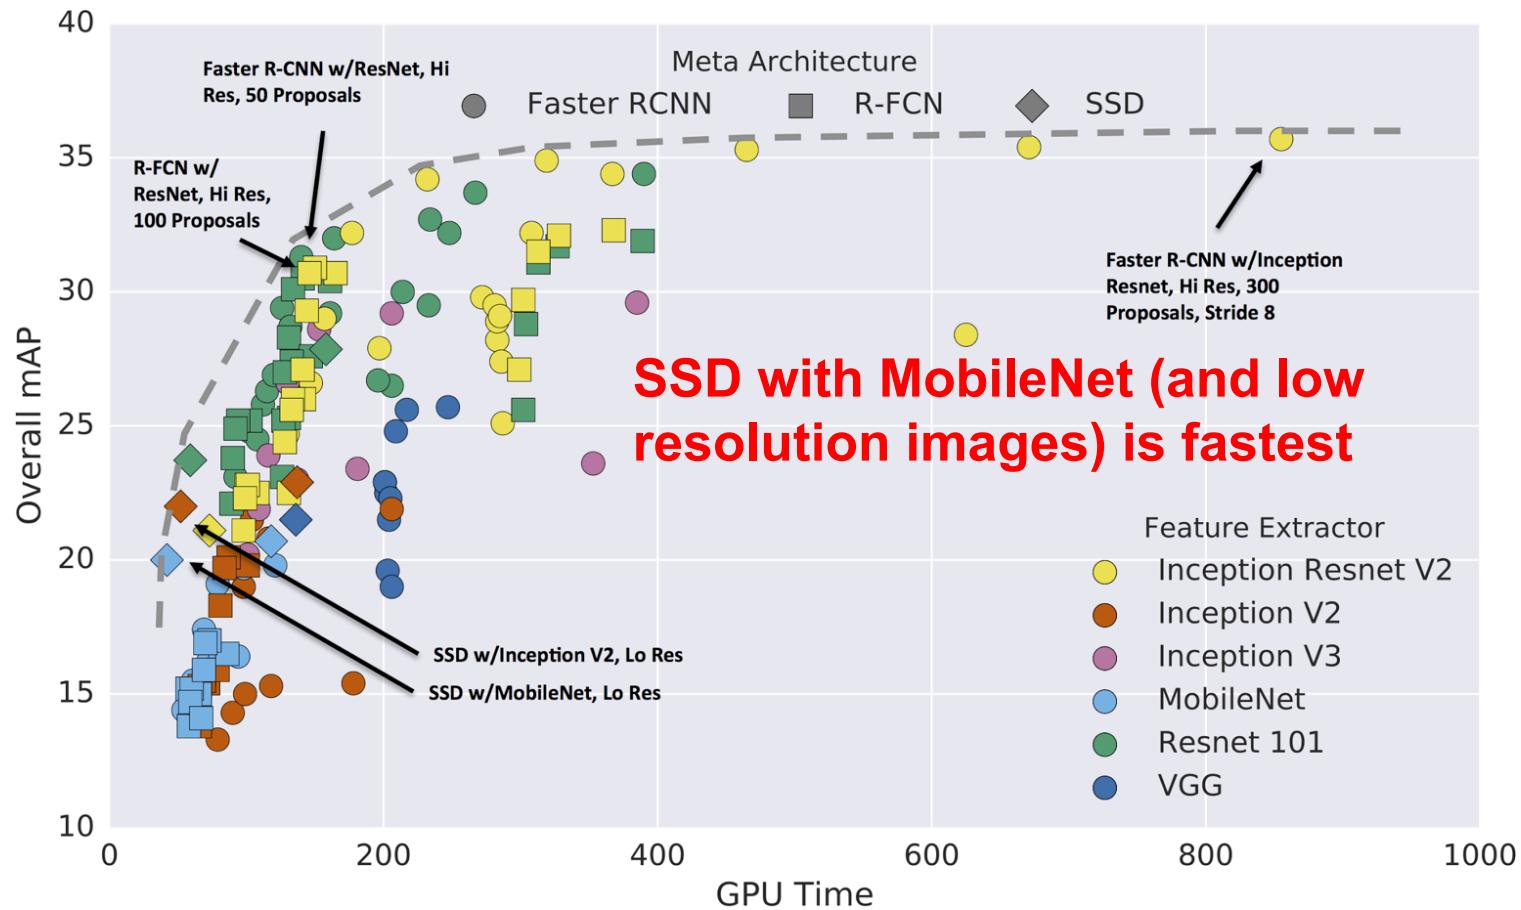
# MobileNet on Imagenet

Model	ImageNet Accuracy	Million Mult-Adds	Million Parameters
Conv MobileNet	71.7%	4866	29.3
MobileNet	70.6%	569	4.2

Model	ImageNet Accuracy	Million Mult-Adds	Million Parameters
1.0 MobileNet-224	70.6%	569	4.2
GoogleNet	69.8%	1550	6.8
VGG 16	71.5%	15300	138



Huang, J., Rathod, V., Sun, C., Zhu, M., Korattikara, A., Fathi, A., ... & Murphy, K. Speed/accuracy trade-offs for modern convolutional object detectors. CVPR 2017



Huang, J., Rathod, V., Sun, C., Zhu, M., Korattikara, A., Fathi, A., ... & Murphy, K. Speed/accuracy trade-offs for modern convolutional object detectors. CVPR 2017

# Detection Model Zoo

elephant: 99%



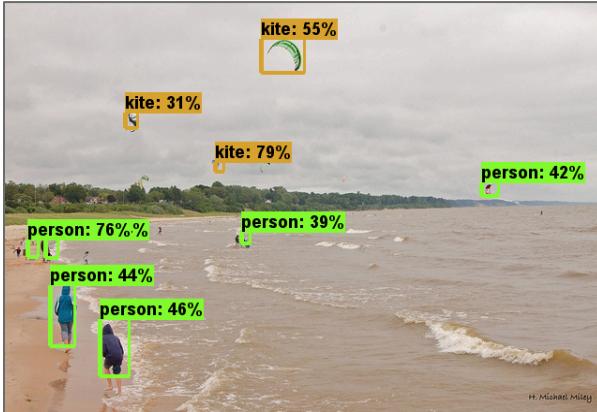
elephant: 99%



## Model summary

- |                                                                   | minival mAP |
|-------------------------------------------------------------------|-------------|
| (Fastest) SSD w/MobileNet (Low Resolution)                        | 19.3        |
| (Fastest) SSD w/Inception V2 (Low Resolution)                     | 22          |
| (Sweet Spot) Faster R-CNN w/Resnet 101, 100 Proposals             | 32          |
| (Sweet Spot) R-FCN w/Resnet 101, 300 Proposals                    | 30.4        |
| (Most Accurate) Faster R-CNN w/Inception Resnet V2, 300 Proposals | 35.7        |

**SSD w/MobileNet  
(Low Resolution)**



**SSD w/Inception V2  
(Low Resolution)**



**Faster R-CNN w/Resnet101,  
100 proposals**



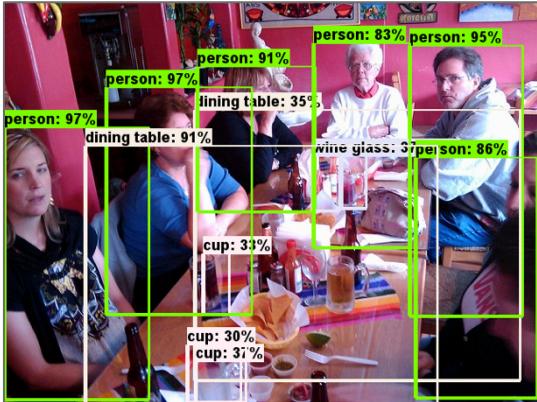
**RFCN w/Resnet101, 300**



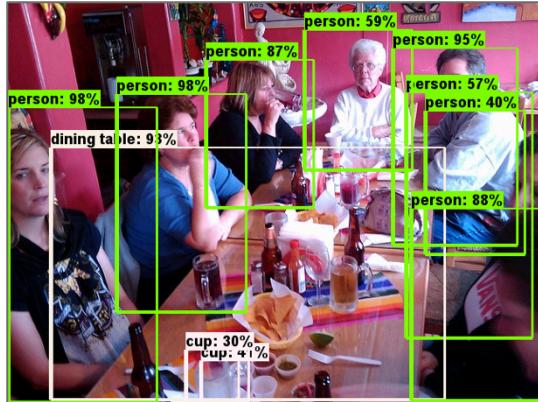
**Faster R-CNN w/Inception  
Resnet V2, 300 proposals**



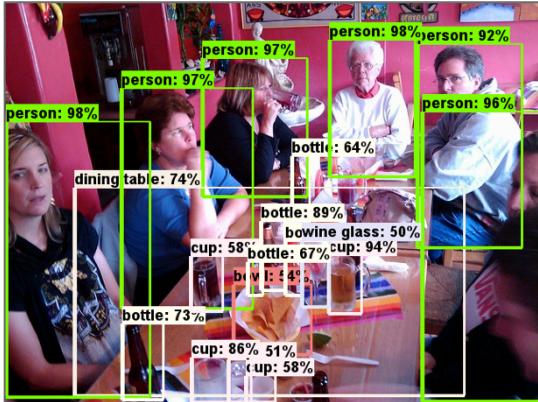
## SSD w/MobileNet (Low Resolution)



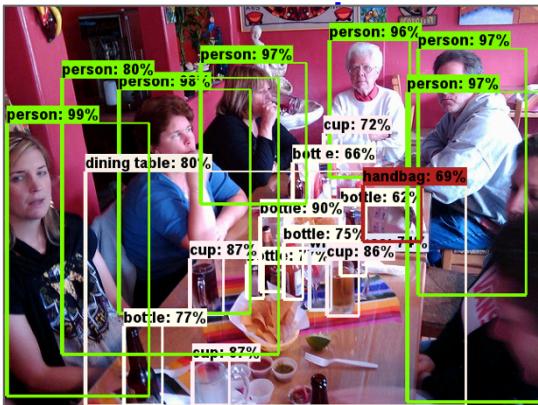
## SSD w/Inception V2 (Low Resolution)



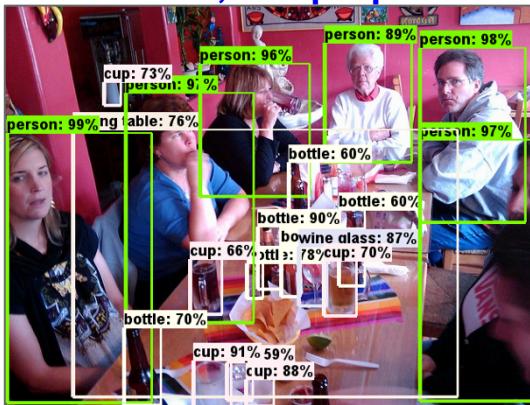
## Faster R-CNN w/Resnet101, 100 proposals



## RFCN w/Resnet101, 300



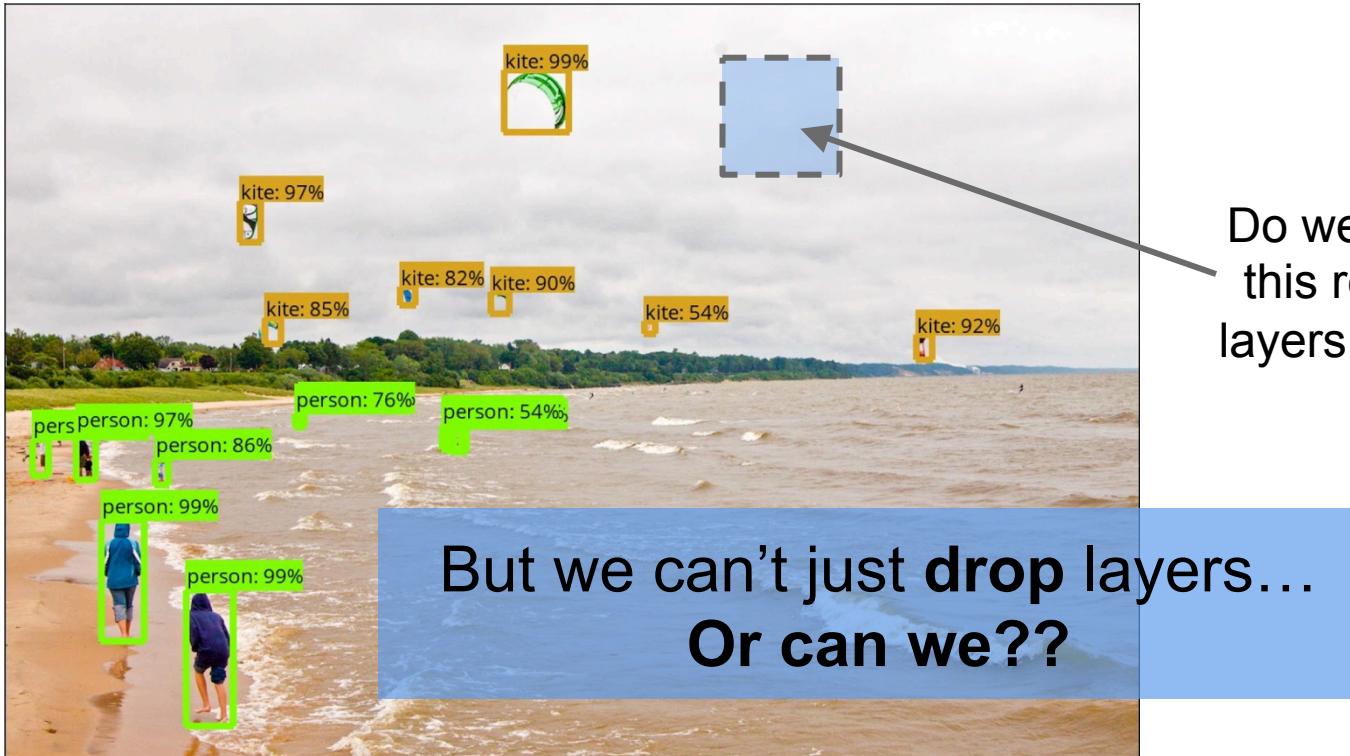
## Faster R-CNN w/Inception Resnet V2, 300 proposals



# This Talk

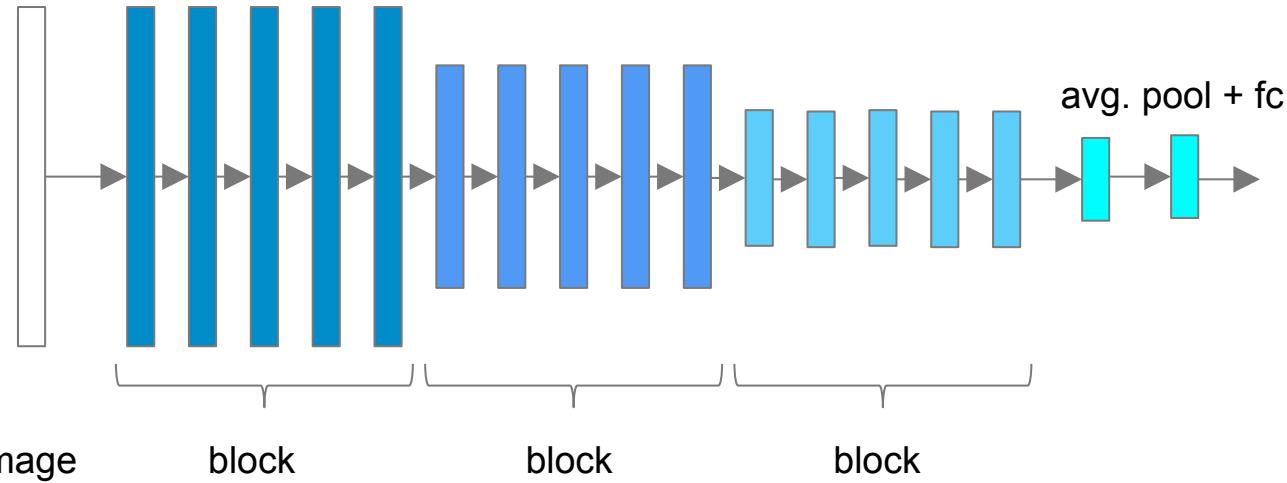
- The Tensorflow Object Detection API
- Pushing the envelope on Accuracy
  - Better Models
  - Better Data
- **Pushing the envelope on Speed**
  - Lightweight Models
  - **Adaptive Inference**

# Idea: Not all parts of an image are equally interesting!



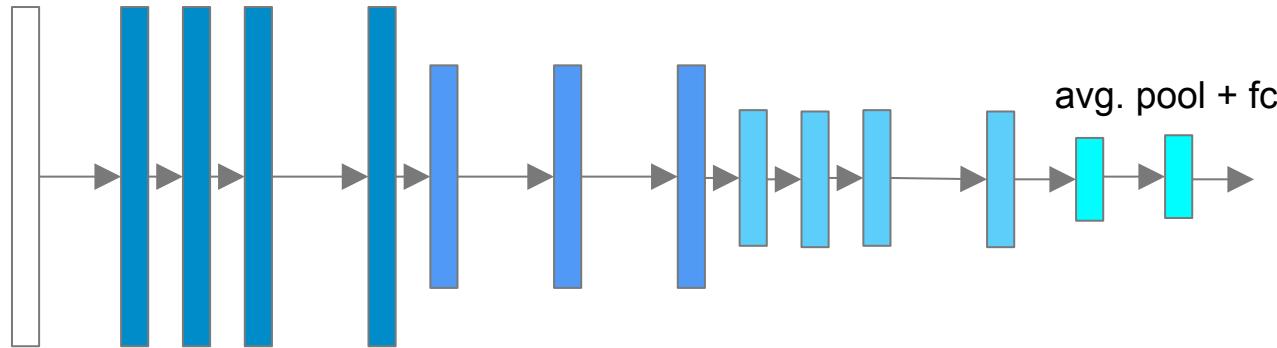
Do we really need to pass this region through ~100 layers of our trained CNN?

# Residual Networks (ResNet)



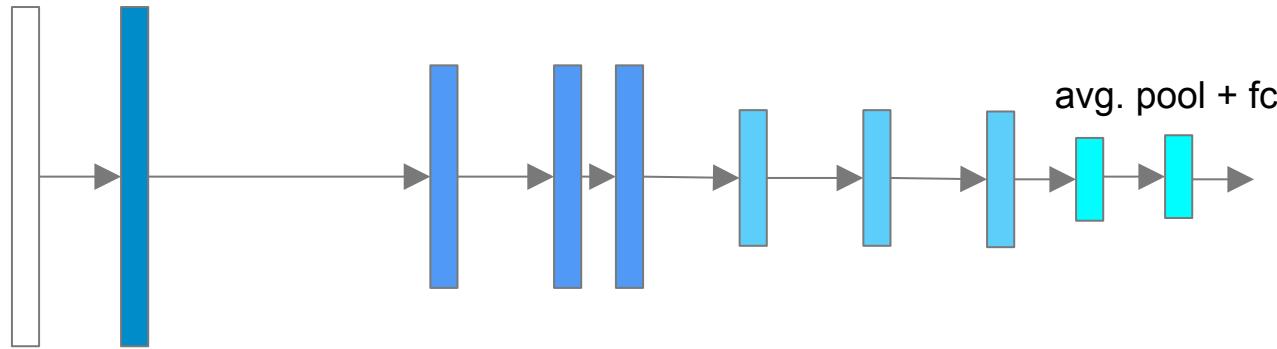
He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 770-778).

# ResNet with Stochastic Depth



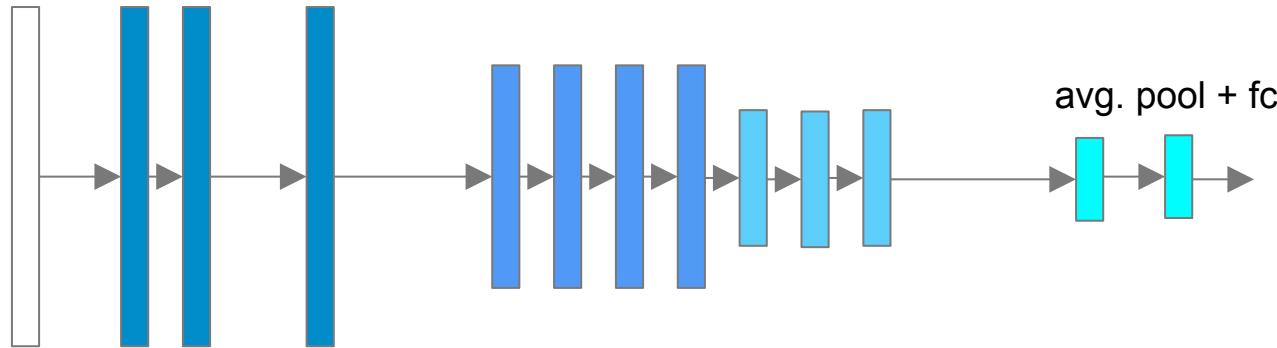
Huang, G., Sun, Y., Liu, Z., Sedra, D., & Weinberger, K. Q. (2016, October). Deep networks with stochastic depth. In *European Conference on Computer Vision* (pp. 646-661). Springer International Publishing.

# ResNet with Stochastic Depth



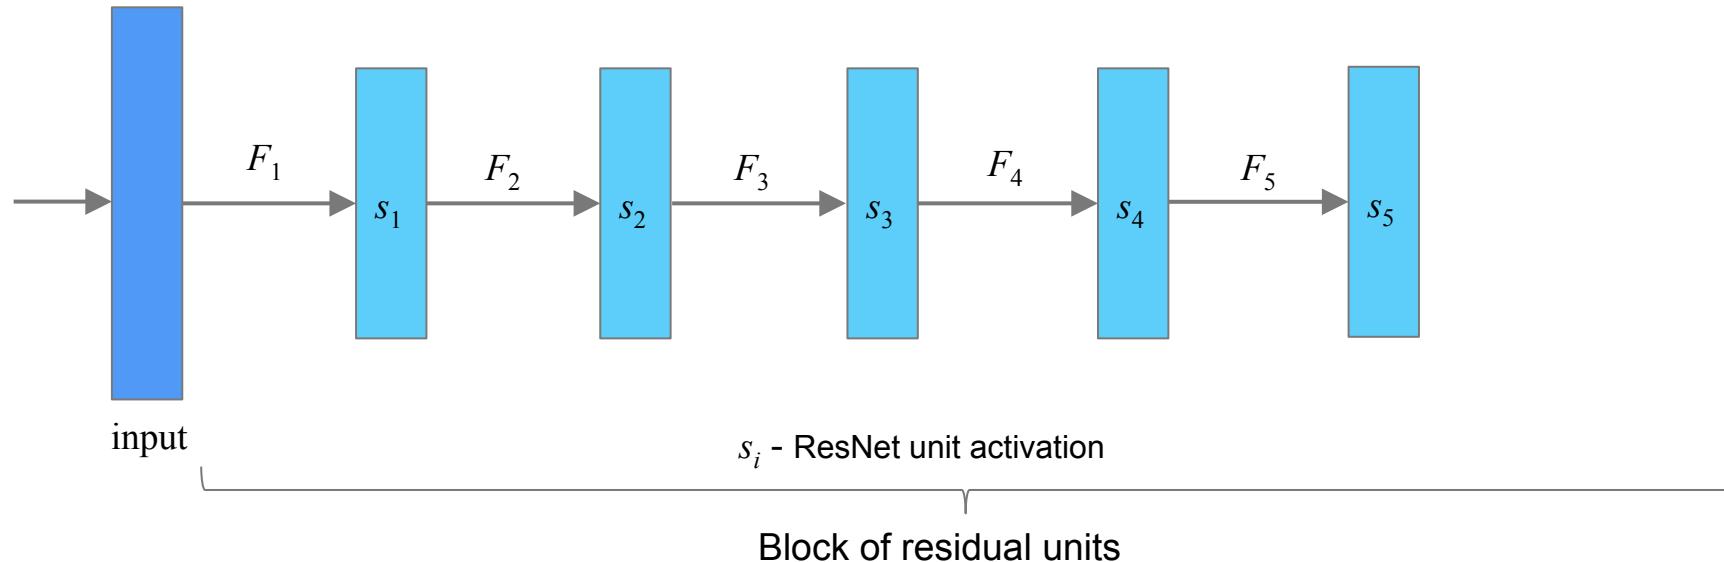
Huang, G., Sun, Y., Liu, Z., Sedra, D., & Weinberger, K. Q. (2016, October). Deep networks with stochastic depth. In *European Conference on Computer Vision* (pp. 646-661). Springer International Publishing.

# ResNet with Stochastic Depth



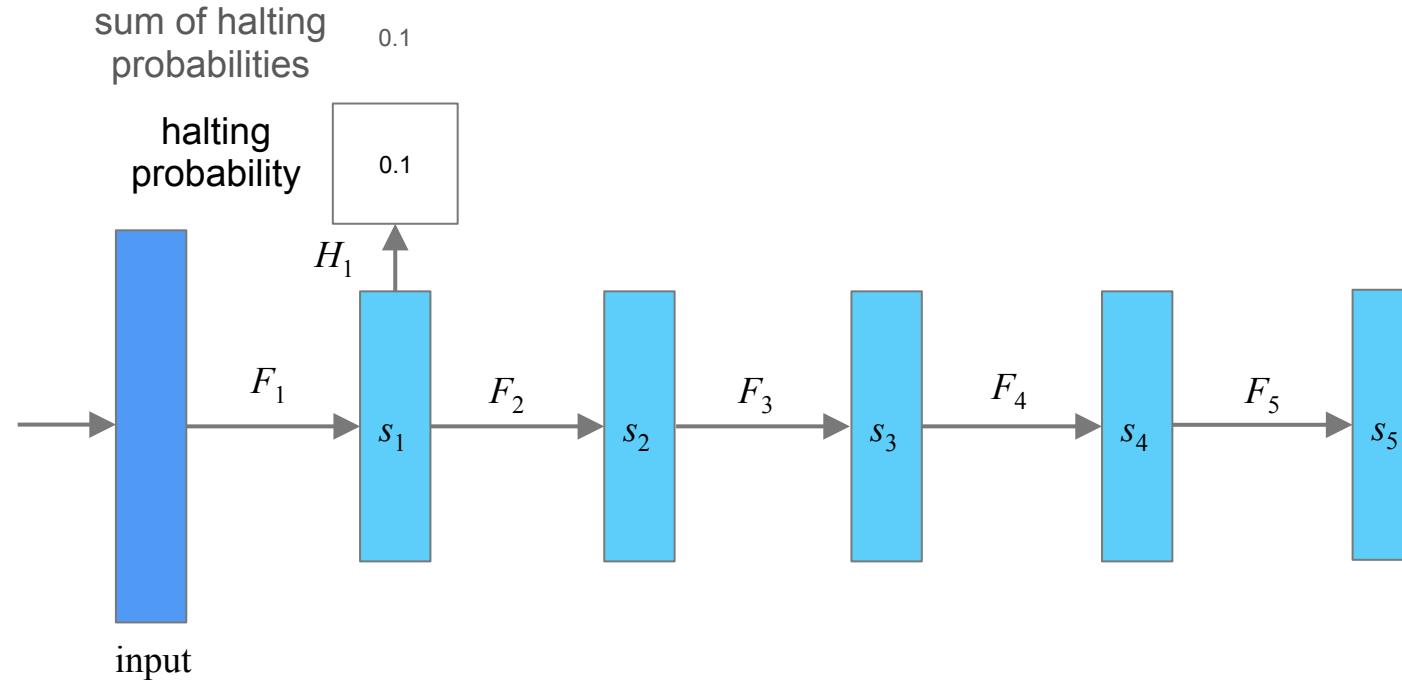
Huang, G., Sun, Y., Liu, Z., Sedra, D., & Weinberger, K. Q. (2016, October). Deep networks with stochastic depth. In *European Conference on Computer Vision* (pp. 646-661). Springer International Publishing.

# Adaptive Computation Time for Residual Networks

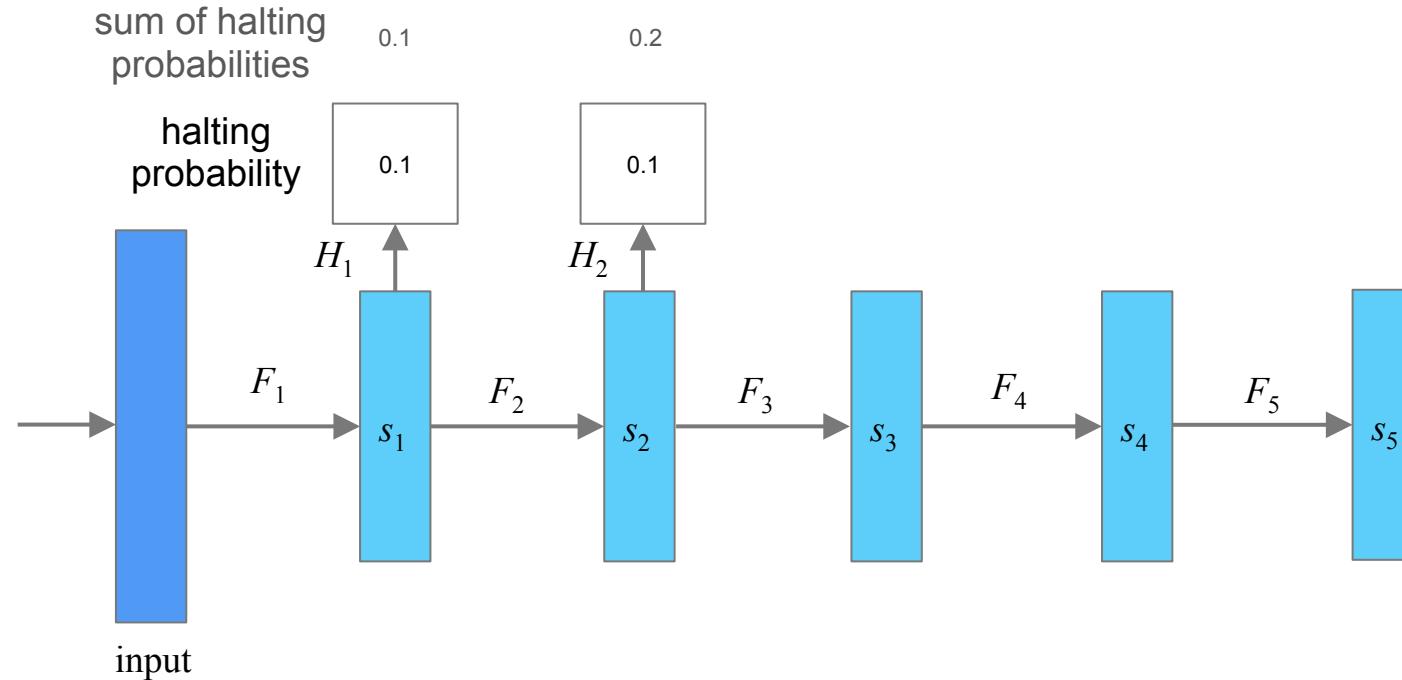


Figurnov, M., Collins, M. D., Zhu, Y., Zhang, L., Huang, J., Vetrov, D., & Salakhutdinov, R. Spatially adaptive computation time for residual networks. CVPR 2017

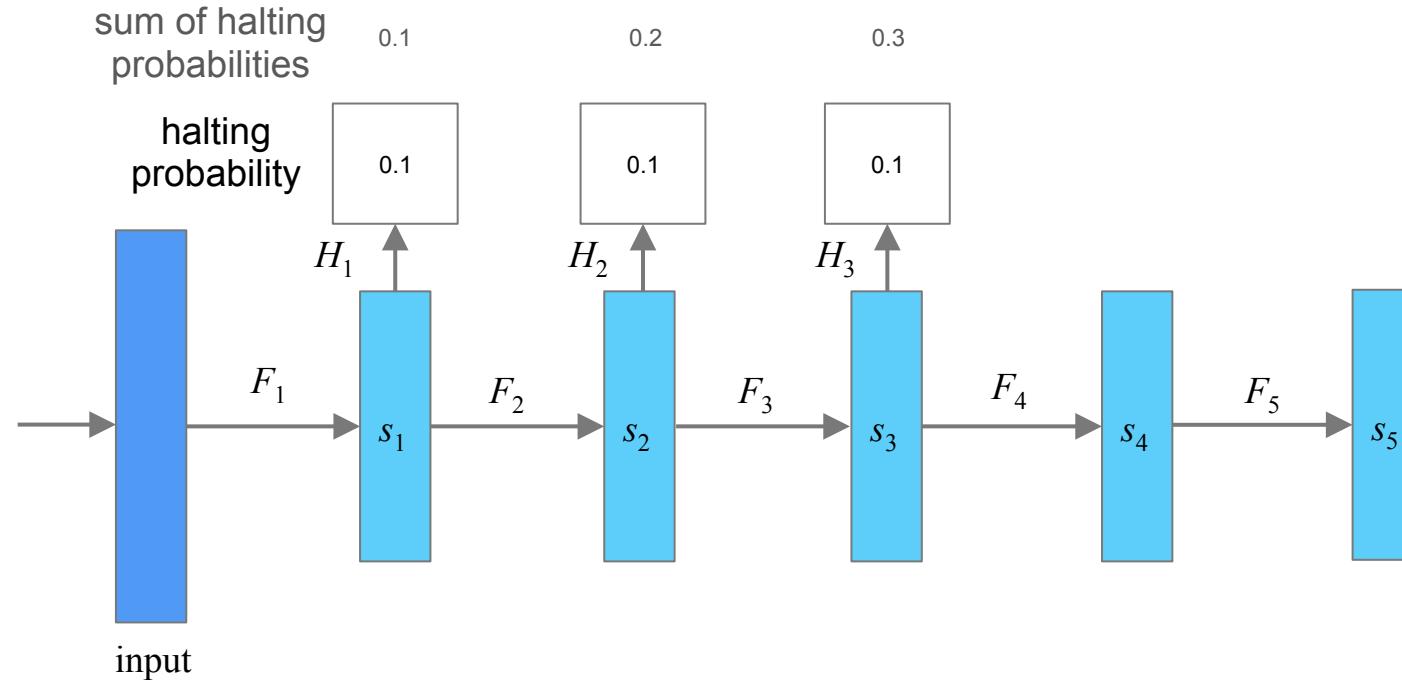
# Adaptive Computation Time for Residual Networks



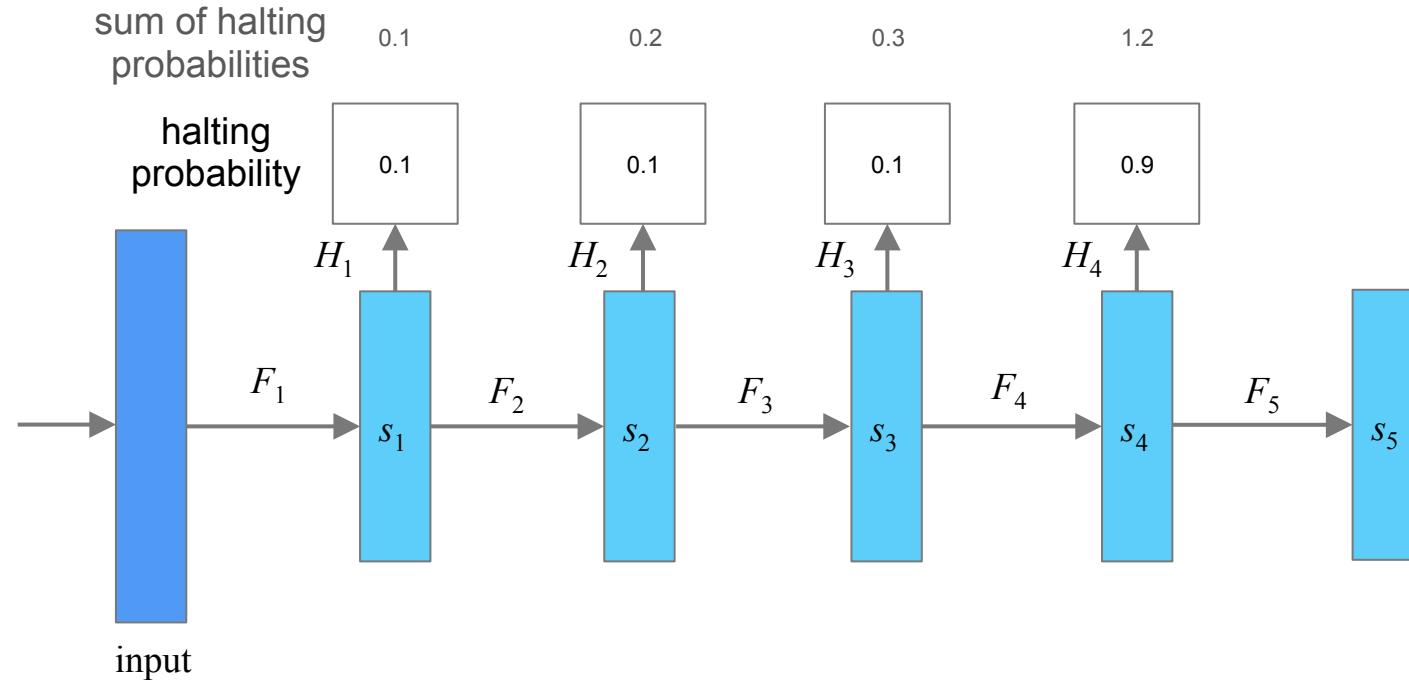
# Adaptive Computation Time for Residual Networks



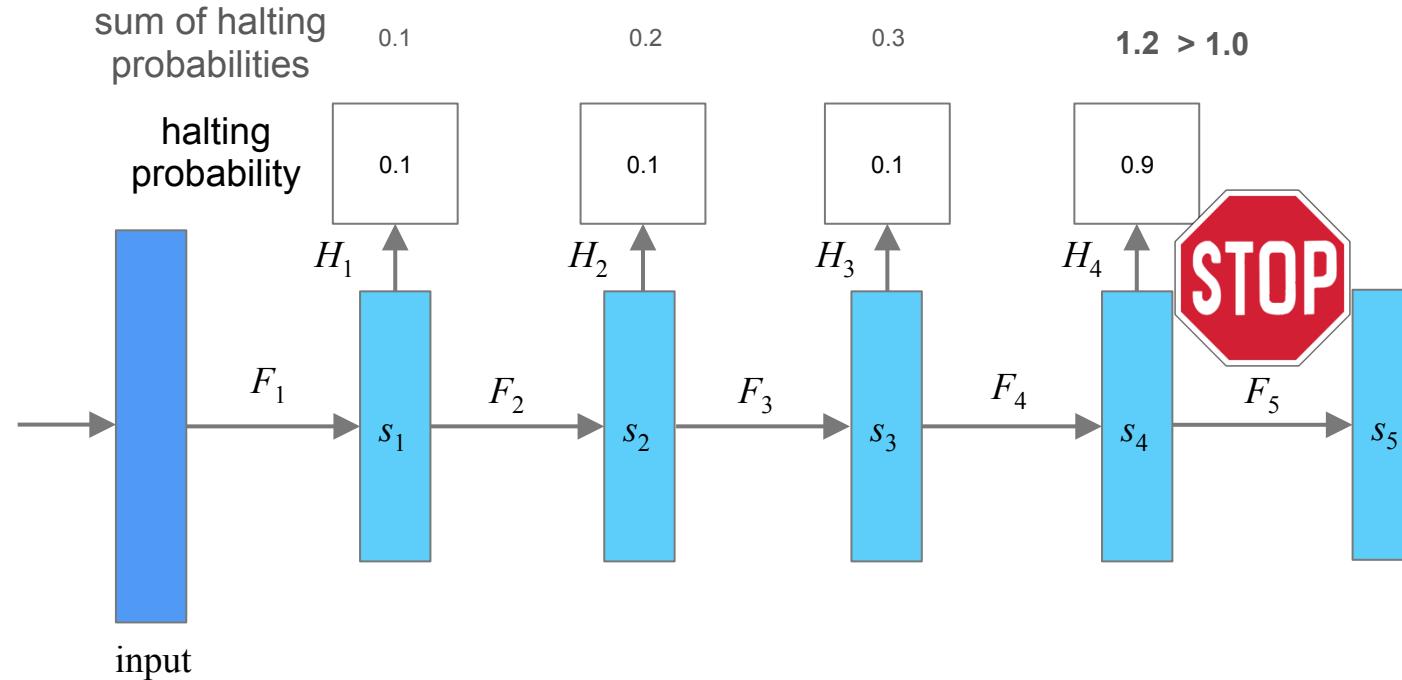
# Adaptive Computation Time for Residual Networks



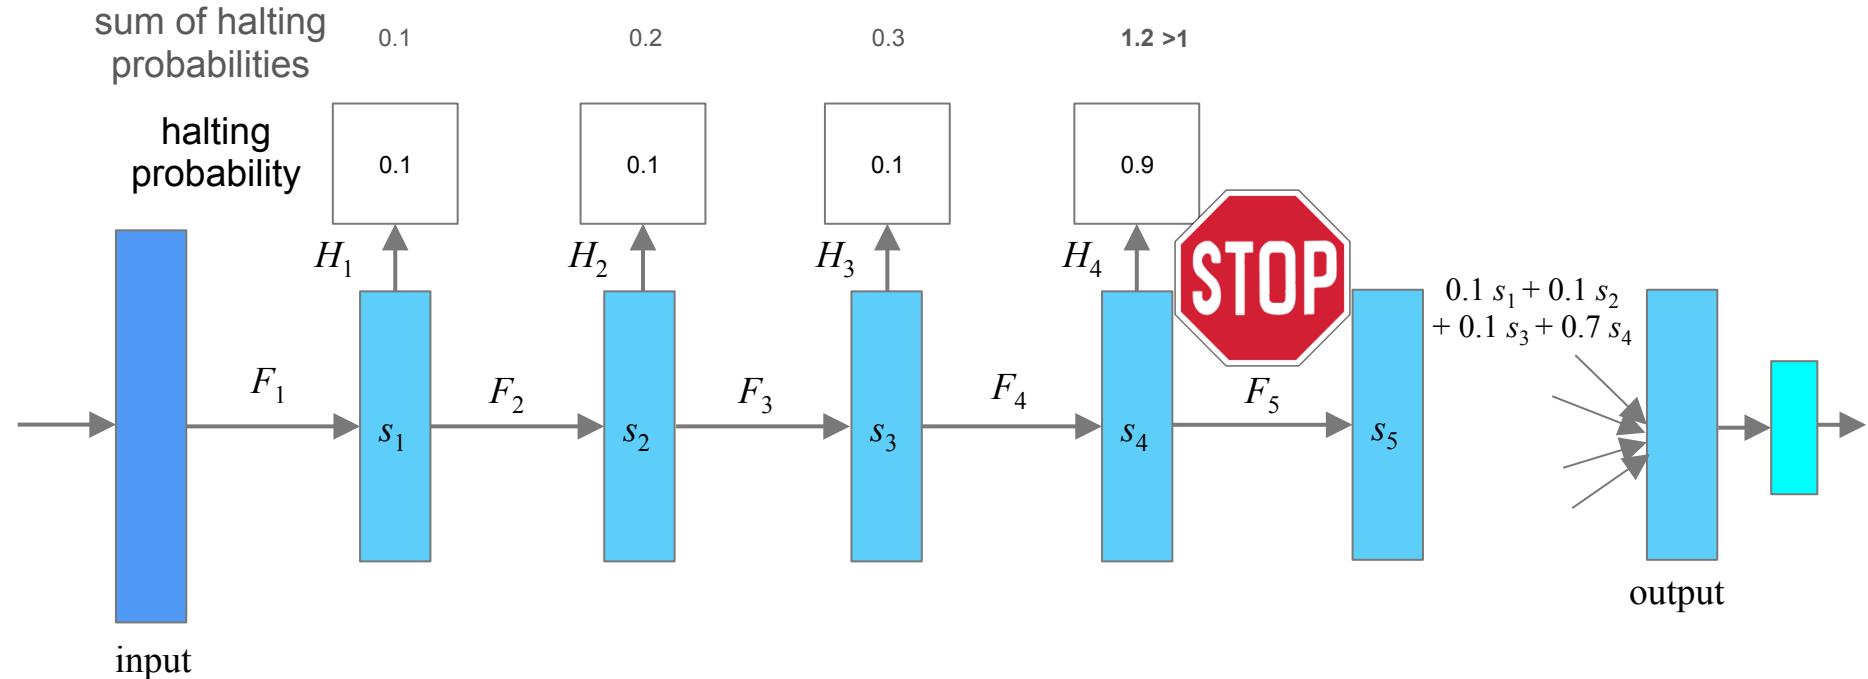
# Adaptive Computation Time for Residual Networks



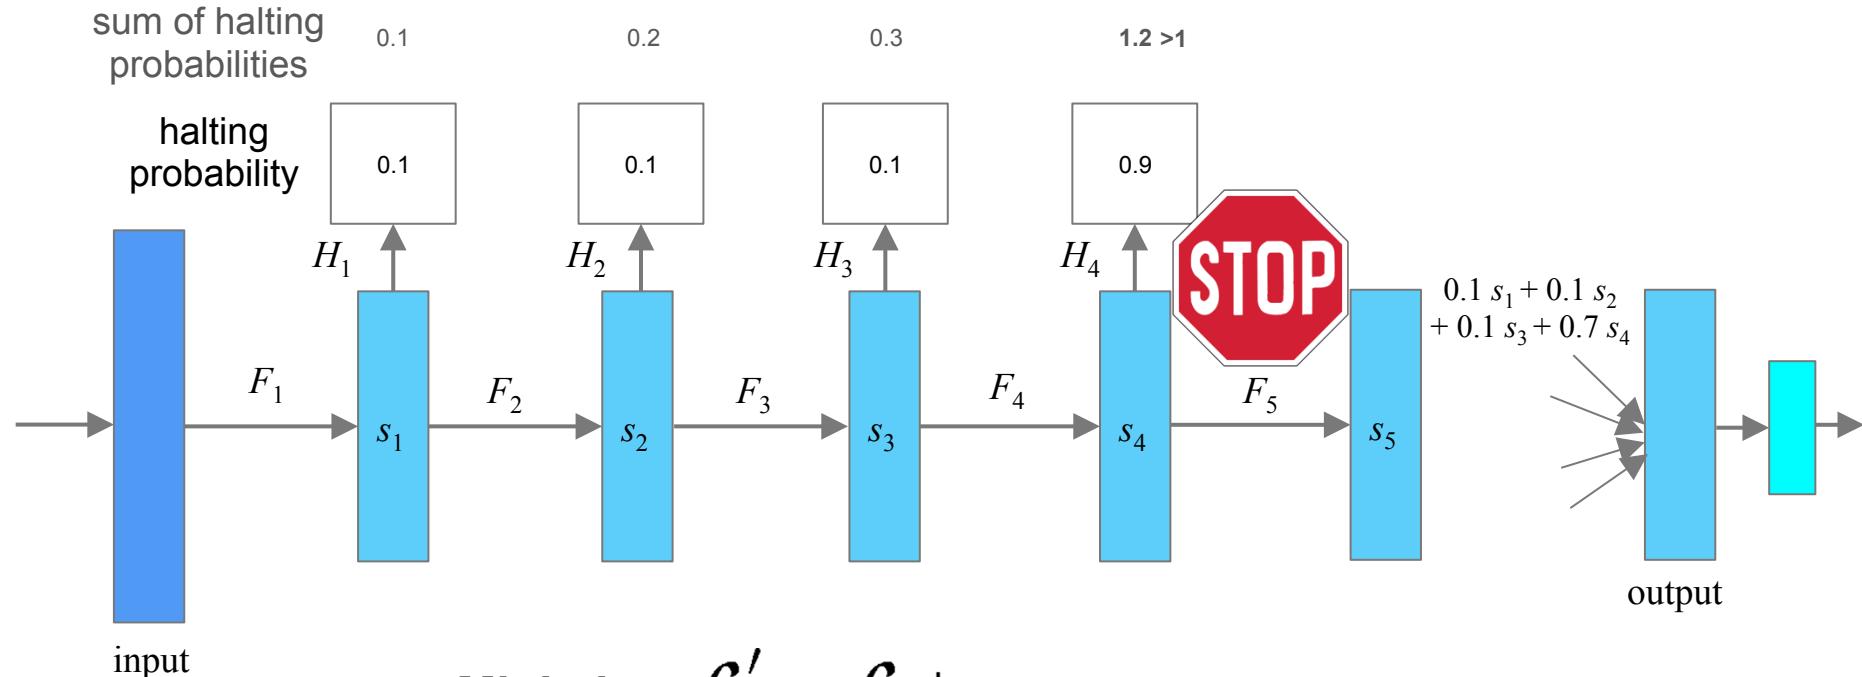
# Adaptive Computation Time for Residual Networks



# Adaptive Computation Time for Residual Networks



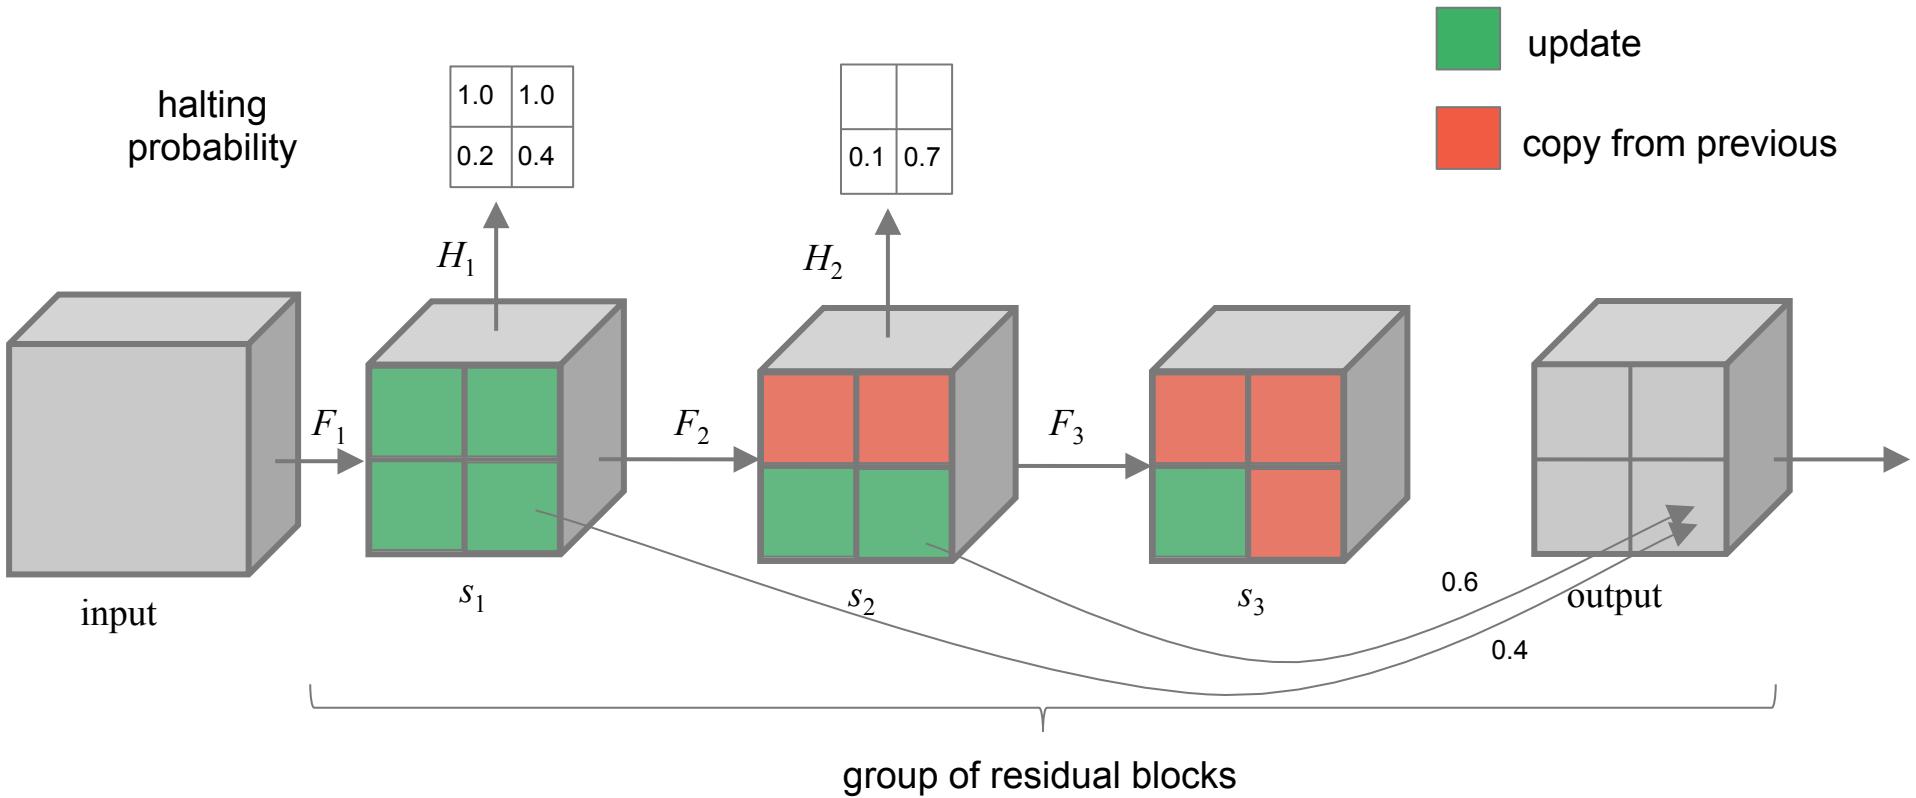
# Adaptive Computation Time for Residual Networks



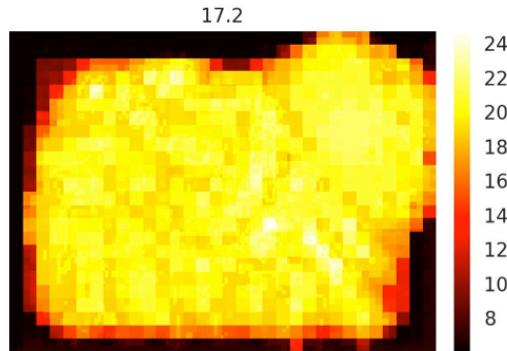
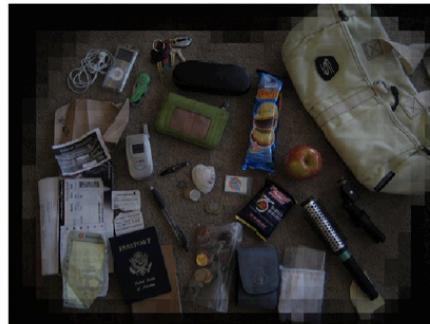
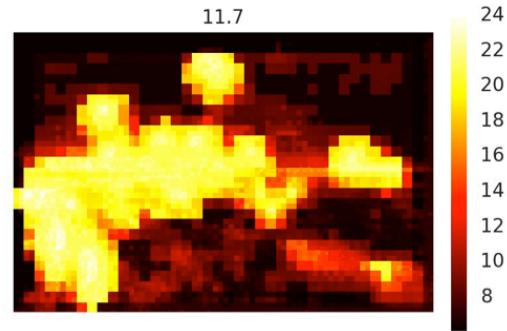
$$\text{Minimize } \mathcal{L}' = \mathcal{L} + \tau \cdot \rho$$

**Ponder cost:** (differentiable) upper bound on number of layers processed

# Spatially Adaptive Computation Time for Residual Networks



# Spatially Adaptive Computation Time for Residual Networks [Figurnov et al '17]



*When used with Faster R-CNN, similar accuracy as full ResNet, only 50% of the FLOPs!*

# Thanks!



Speech Processing

Algorithms  
and Theory

Natural Language  
Processing

Machine Intelligence

Mobile Systems

Security,  
Privacy,  
and Abuse Prevention

General  
Science

Information Retrieval  
and the Web

## SPEAKER

Jonathan Huang ([jonathanhuang@google.com](mailto:jonathanhuang@google.com))

## Special thanks to

Derek Chow, Alireza Fathi, Ian Fischer, Sergio Guadarrama, Jonathan Huang, Anoop Korattikara, Kevin Murphy, Vivek Rathod, Yang Song, Chen Sun, Matt Tang, Zbigniew Wojna, Menglong Zhu, Tom Duerig, Dumitru Erhan, Jitendra Malik, George Papandreou, Dominik Roblek, Chuck Rosenberg, Nathan Silberman, Abhinav Srivastava, Rahul Sukthankar, Christian Szegedy, Jasper Uijlings, Jay Yagnik, Xiangxin Zhu, Andrew Howard, Michael Figurnov, Vittorio Ferrari, Victor Gomes, Michael Hays